

A Not that Comprehensive Introduction to Neural Programming

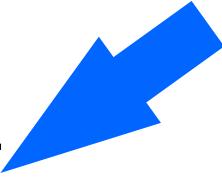
Lili Mou
doublepower.mou@gmail.com

Outline

- Introduction
- Symbolic execution
 - (Fully supervised) Neural programmer interpreter
 - (Weakly supervised) Neural symbolic machine
 - “Spurious programs” and inductive programming
 - Learning semantic parsers from denotations
 - DeepCoder
 - More thoughts on “spurious programs”
- Distributed execution
 - Learning to execute
 - Neural enquirer
 - Differentiable neural computer
- Hybrid execution
 - Neural programmer
 - Coupling approach

Outline

- Introduction
- Symbolic execution
 - (Fully supervised) Neural programmer interpreter
 - (Weakly supervised) Neural symbolic machine
 - “Spurious programs” and inductive programming
 - Learning semantic parsers from denotations
 - DeepCoder
 - More thoughts on “spurious programs”
- Distributed execution
 - Learning to execute
 - Neural enquirer
 - Differentiable neural computer
- Hybrid execution
 - Neural programmer
 - Coupling approach



NEURAL PROGRAMMER-INTERPRETERS

Scott Reed & Nando de Freitas

Google DeepMind

London, UK

scott.ellison.reed@gmail.com

nandodefreitas@google.com

MAKING NEURAL PROGRAMMING ARCHITECTURES GENERALIZE VIA RECURSION

Jonathon Cai, Richard Shin, Dawn Song

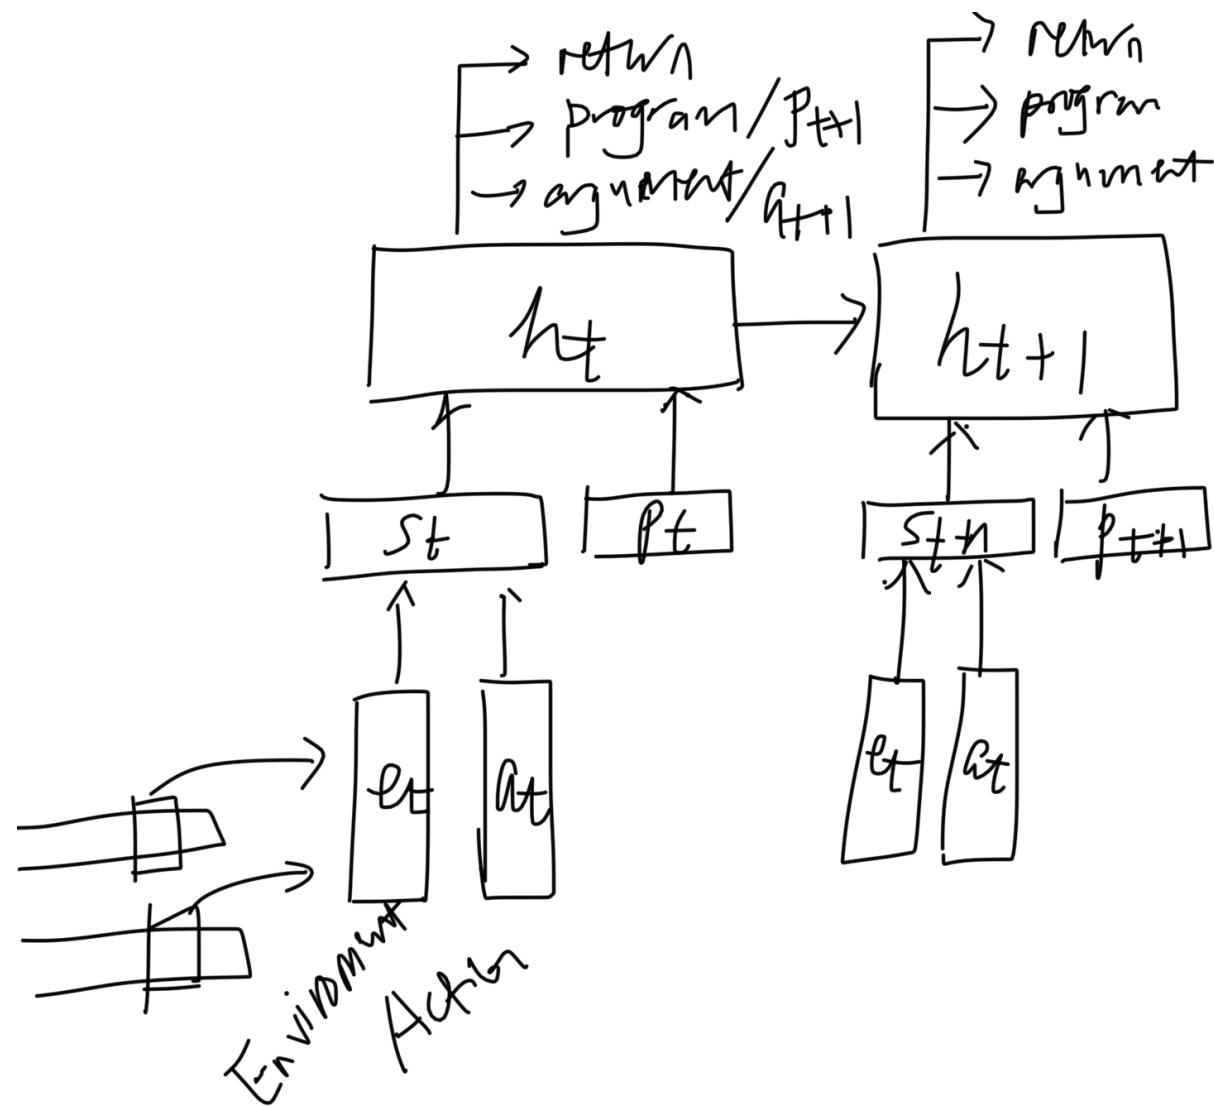
Department of Computer Science
University of California, Berkeley
Berkeley, CA 94720, USA

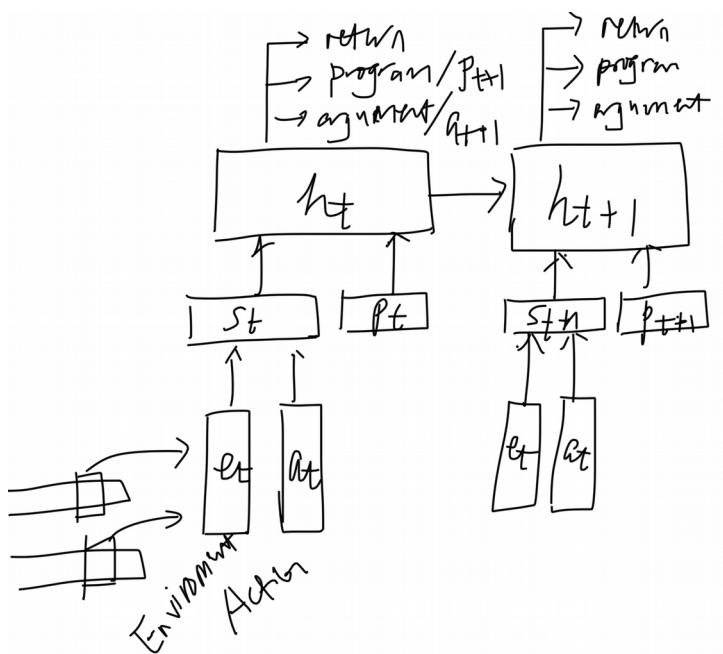
{jonathon, ricshin, dawnsong}@cs.berkeley.edu

$$s_t = f_{enc}(e_t, a_t)$$

$$h_t = f_{lstm}(s_t, p_t, h_{t-1})$$

$$r_t = f_{end}(h_t), p_{t+1} = f_{prog}(h_t), a_{t+1} = f_{arg}(h_t)$$





When the program is called, the current context including the caller's memory state is stored on a stack; when the program returns, the stored context is popped off the stack to resume execution in the previous caller's context.

Algorithm 1 Neural programming inference

```

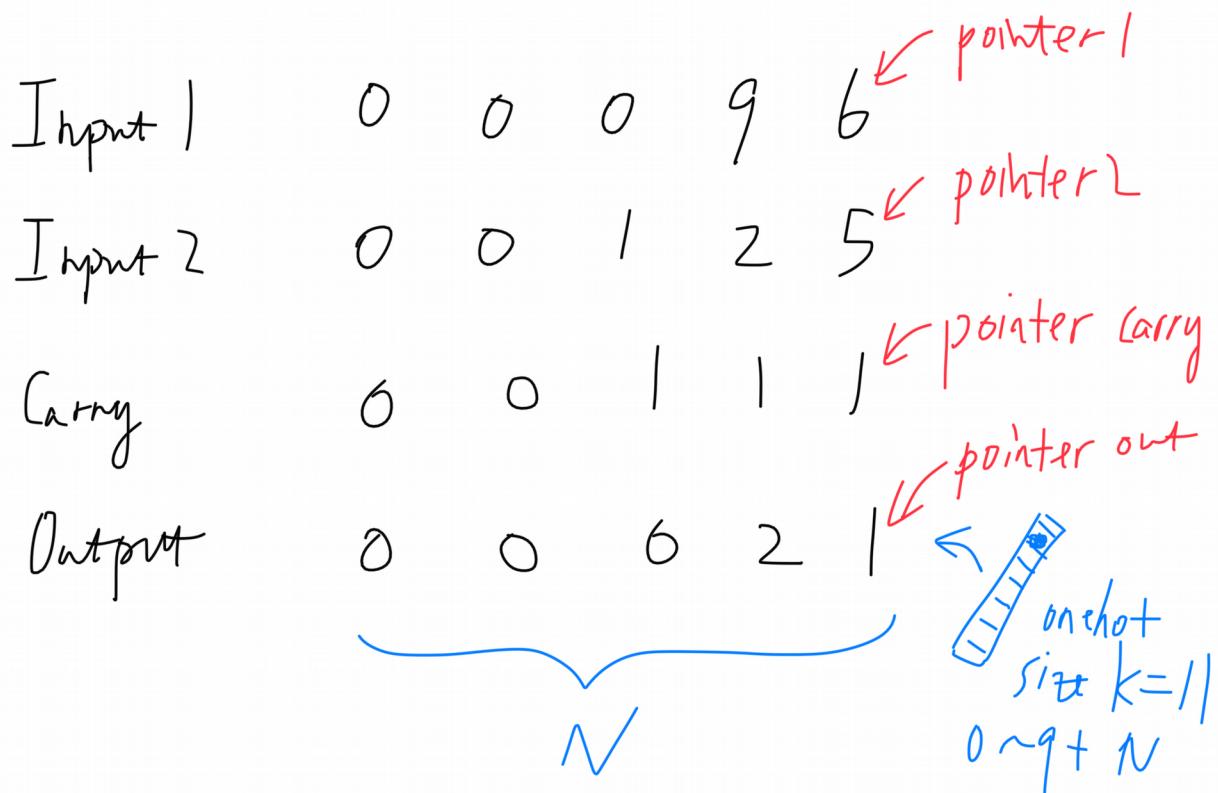
1: Inputs: Environment observation  $e$ , program  $p$ , arguments  $a$ , stop threshold  $\alpha$ 
2: function RUN( $e, p, a$ )
3:    $h \leftarrow \mathbf{0}, r \leftarrow 0$ 
4:   while  $r < \alpha$  do
5:      $s \leftarrow f_{enc}(e, a), h \leftarrow f_{lstm}(s, p, h)$ 
6:      $r \leftarrow f_{end}(h), p_2 \leftarrow f_{prog}(h), a_2 \leftarrow f_{arg}(h)$ 
7:     if  $p$  is a primitive function then
8:        $e \leftarrow f_{env}(e, p, a).$ 
9:     else
10:      function RUN( $e, p_2, a_2$ )

```

Example

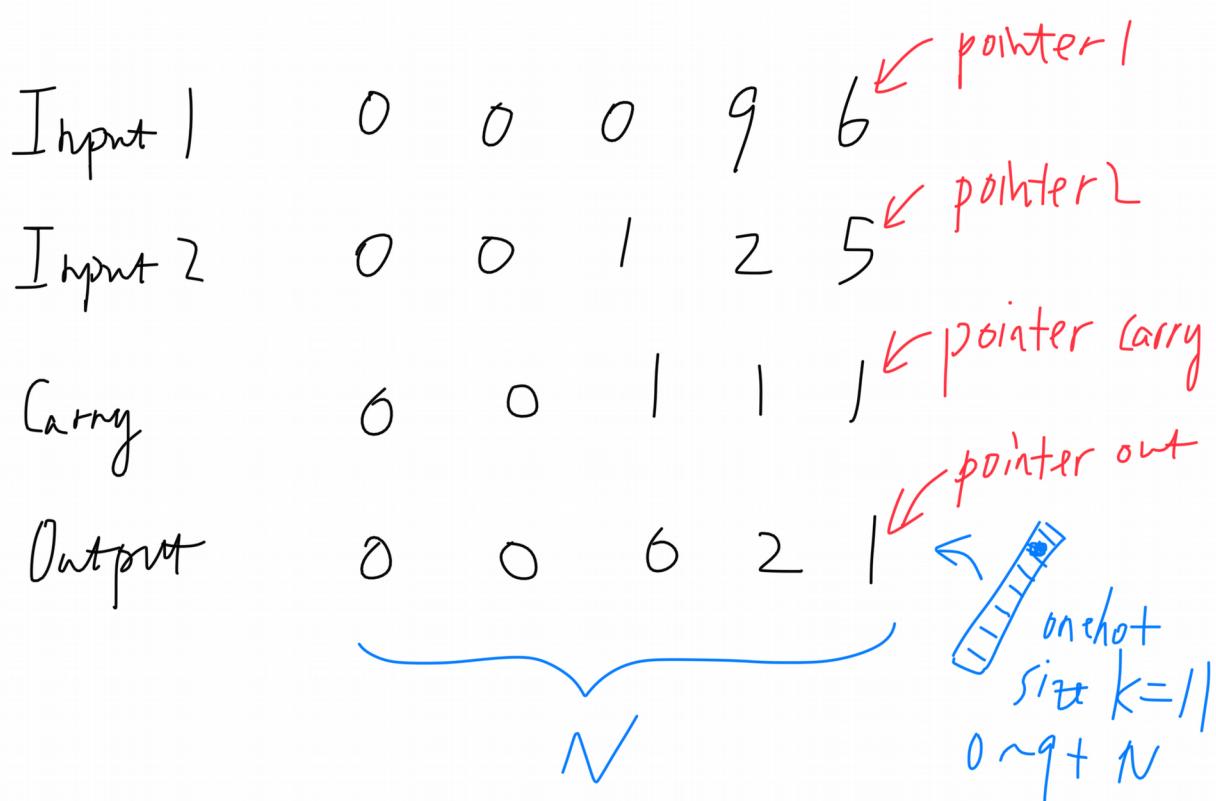
- Grade-school addition

$$f_{enc}(Q, i_1, i_2, i_3, i_4, a_t) = \text{MLP}([Q(1, i_1), Q(2, i_2), Q(3, i_3), Q(4, i_4), a_t(1), a_t(2), a_t(3)])$$



Example

- Grade-school addition



ADD Add for one digit (a wrapper)

ADD1

WRITE OUT 1

CARRY

PTR CARRY LEFT

WRITE CARRY 1

PTR CARRY RIGHT

LSHIFT

PTR INP1 LEFT

PTR INP2 LEFT

PTR CARRY LEFT

PTR OUT LEFT

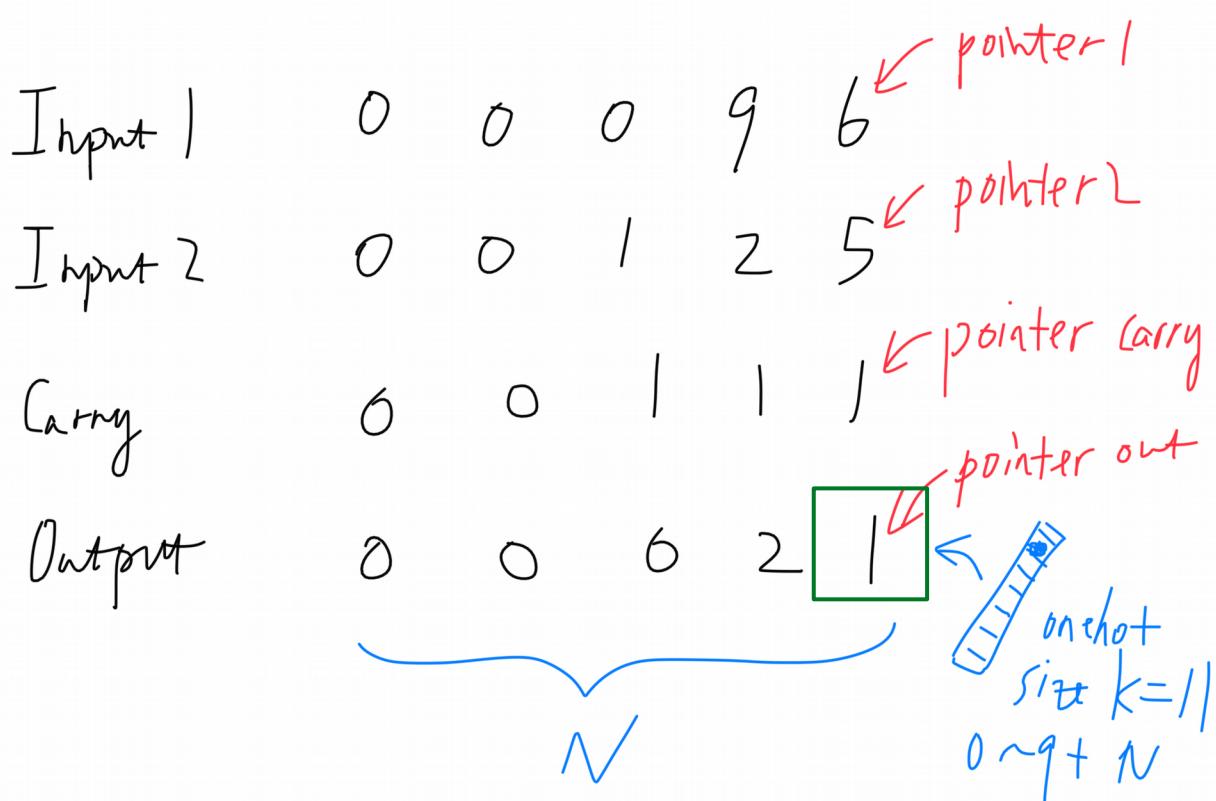
ADD1

...

Example

- Grade-school addition

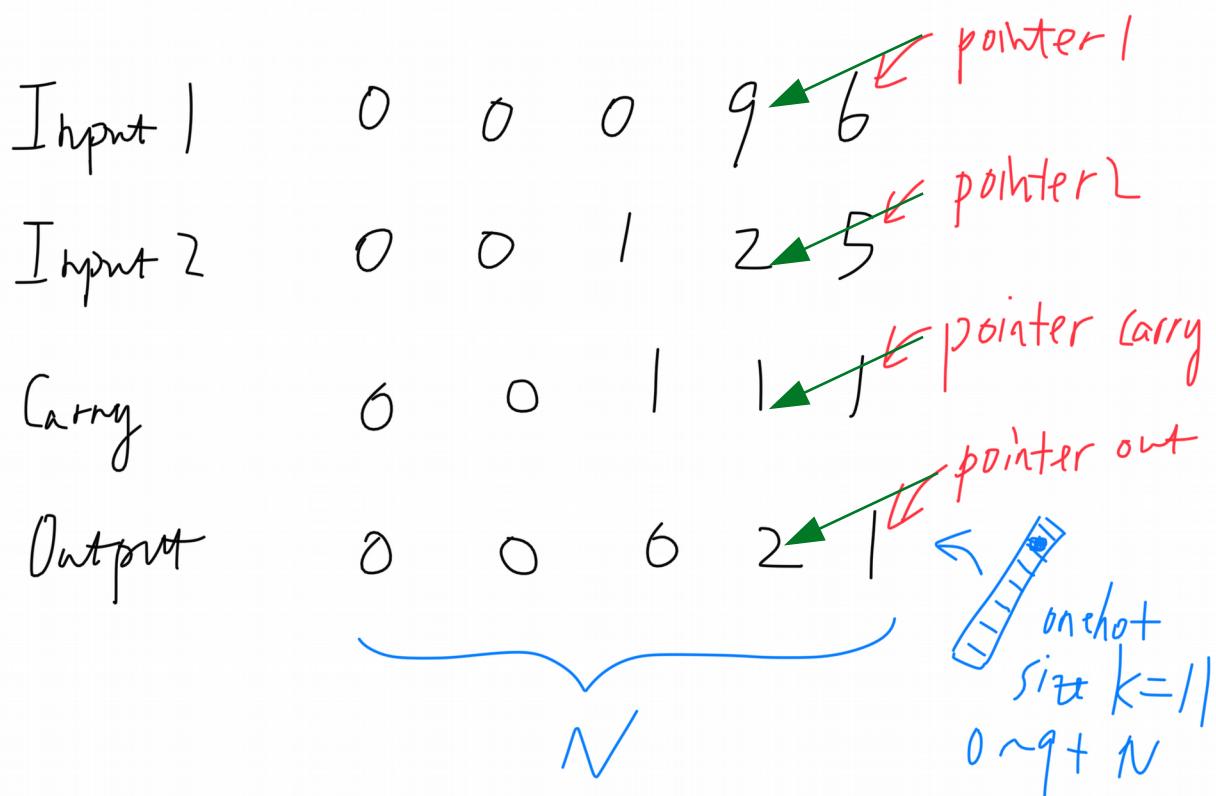
$$f_{enc}(Q, i_1, i_2, i_3, i_4, a_t) = \text{MLP}([Q(1, i_1), Q(2, i_2), Q(3, i_3), Q(4, i_4), a_t(1), a_t(2), a_t(3)]).$$



ADD	Write "1" at the output slip
ADD1	
	WRITE OUT 1
CARRY	
	PTR CARRY LEFT
	WRITE CARRY 1
	PTR CARRY RIGHT
LSHIFT	
	PTR INP1 LEFT
	PTR INP2 LEFT
	PTR CARRY LEFT
	PTR OUT LEFT
ADD1	
	...

Example

- Grade-school addition



ADD

ADD1

WRITE OUT 1

CARRY

PTR CARRY LEFT

WRITE CARRY 1

PTR CARRY RIGHT

LSHIFT

PTR INP1 LEFT

PTR INP2 LEFT

PTR CARRY LEFT

PTR OUT LEFT

ADD1 Move printers left

...

Recursion?

Non-Recursive

```
1 ADD
2 ADD1
3 WRITE OUT 1
4 CARRY
5 PTR CARRY LEFT
6 WRITE CARRY 1
7 PTR CARRY RIGHT
8 LSHIFT
9 PTR INP1 LEFT
10 PTR INP2 LEFT
11 PTR CARRY LEFT
12 PTR OUT LEFT
13 ADD1
14 ...
```

Recursive

```
1 ADD
2 ADD1
3 WRITE OUT 1
4 CARRY
5 PTR CARRY LEFT
6 WRITE CARRY 1
7 PTR CARRY RIGHT
8 LSHIFT
9 PTR INP1 LEFT
10 PTR INP2 LEFT
11 PTR CARRY LEFT
12 PTR OUT LEFT
13 ADD
14 ...
```

The Magic Here

- No magic
- Fully supervised
 - What “ADD” sees is
 - ADD1, LSHIFT, ADD, RETURN
 - What “ADD1” sees is
 - WRITE (parameters: OUT, 1)
 - CARRY, RETURN
 - Etc.

Recursive

```
ADD  
ADD1  
WRITE OUT 1  
CARRY  
PTR CARRY LEFT  
WRITE CARRY 1  
PTR CARRY RIGHT  
LSHIFT  
PTR INP1 LEFT  
PTR INP2 LEFT  
PTR CARRY LEFT  
PTR OUT LEFT  
ADD  
...
```

Design Philosophy

- Define a “local” context, which the current step of execution fully depends on.
- Manually program in terms of the above “language.”
- Generate execution sequences (a.k.a. program traces) with some (not many) examples.

Outline

- Introduction
- Symbolic execution
 - (Fully supervised) Neural programmer interpreter
 - (Weakly supervised) Neural symbolic machine
 - “Spurious programs” and inductive programming
 - Learning semantic parsers from denotations
 - DeepCoder
 - More thoughts on “spurious programs”
- Distributed execution
 - Learning to execute
 - Neural enquirer
 - Differentiable neural computer
- Hybrid execution
 - Neural programmer
 - Coupling approach



Neural Symbolic Machines: Learning Semantic Parsers on Freebase with Weak Supervision

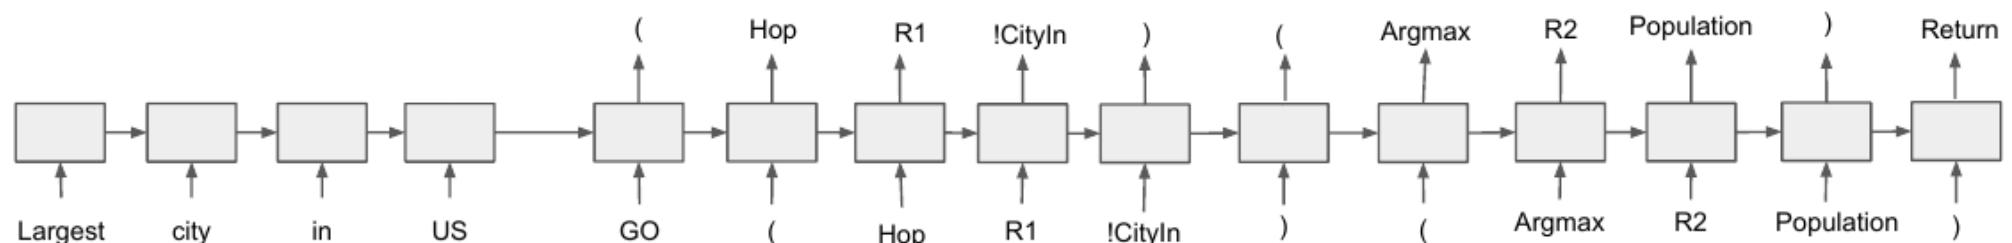
Chen Liang*, Jonathan Berant[†], Quoc Le, Kenneth D. Forbus, Ni Lao

Northwestern University, Evanston, IL
`{chenliang2013,forbus}@u.northwestern.edu`

Tel-Aviv University, Tel Aviv-Yafo, Israel
`joberant@cs.tau.ac.il`

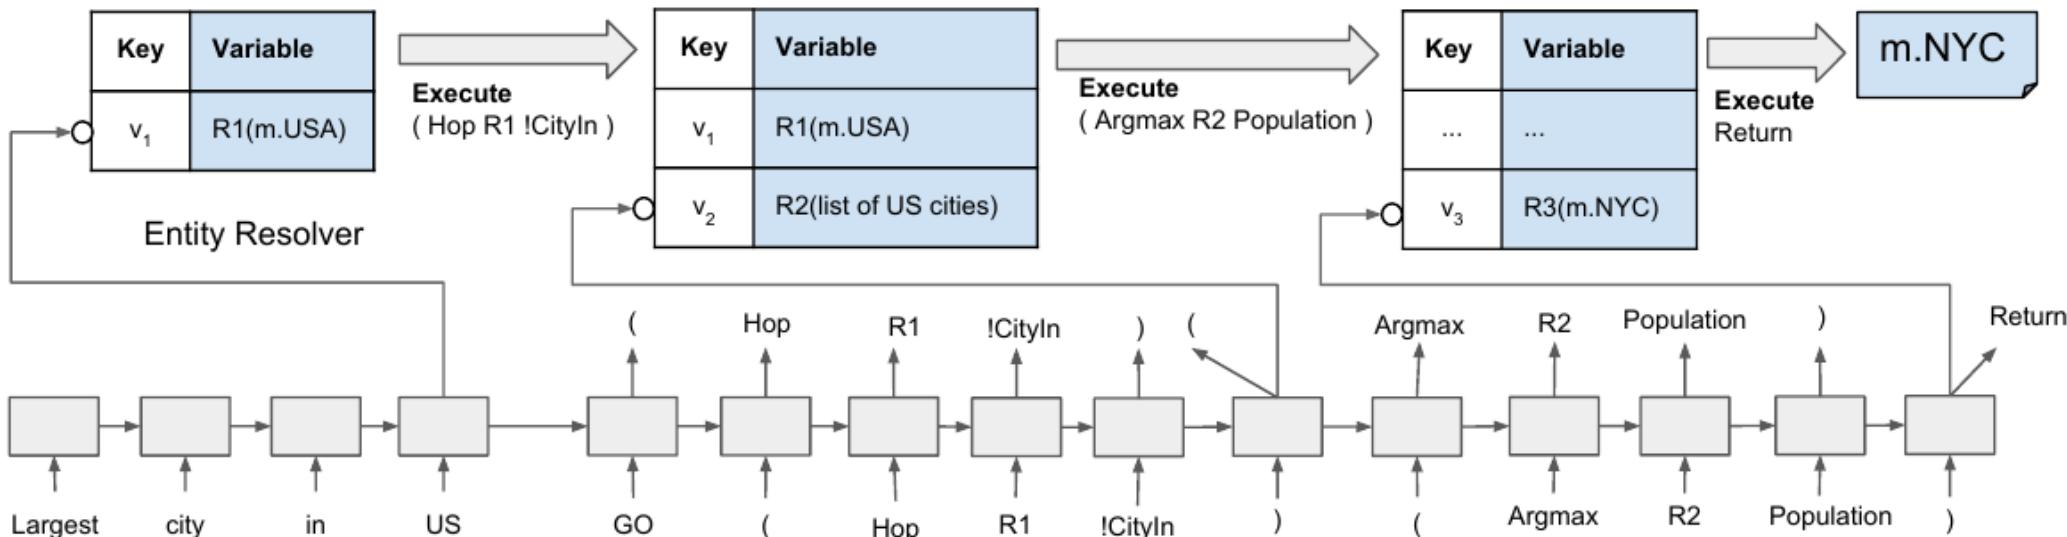
Google Inc., Mountain View, CA
`{qvl,nlao}@google.com`

Architecture



- Generally seq2seq

Architecture



- But with tricks
 - Entities marked with GRU output as embeddings
 - Searched with attention

Training

- Generally REINFORCE

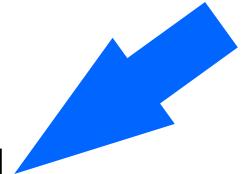
$$\nabla_{\theta} J^{RL}(\theta) = \sum_q \sum_{a_{0:T}} P(a_{0:T}|q, \theta) [R(q, a_{0:T}) - B(q)] \nabla_{\theta} \log P(a_{0:T}|q, \theta)$$

Training

- But with more tricks
 - Beam search instead of sampling
 - Keep the “best” obtained programs so far and perform supervised training in addition to REINFORCE

Outline

- Introduction
- Symbolic execution
 - (Fully supervised) Neural programmer interpreter
 - (Weakly supervised) Neural symbolic machine
 - “Spurious programs” and inductive programming
 - Learning semantic parsers from denotations
 - DeepCoder
 - More thoughts on “spurious programs”
- Distributed execution
 - Learning to execute
 - Neural enquirer
 - Differentiable neural computer
- Hybrid execution
 - Neural programmer
 - Coupling approach



Inferring Logical Forms From Denotations

Panupong Pasupat

Computer Science Department
Stanford University

ppasupat@cs.stanford.edu

Percy Liang

Computer Science Department
Stanford University

pliang@cs.stanford.edu

(ACL-16)

- Spurious programs
- Fictitious worlds

The Problem of “Spurious Programs”

Year	Venue	Position	Event	Time
2001	Hungary	2nd	400m	47.12
2003	Finland	1st	400m	46.69
2005	Germany	11th	400m	46.62
2007	Thailand	1st	relay	182.05
2008	China	7th	relay	180.32

x : “Where did the last 1st place finish occur?”

y : Thailand

Consistent
Correct
$z_1: \mathbf{R}[\text{Venue}].\text{argmax}(\text{Position.1st}, \text{Index})$ Among rows with Position = 1st, pick the one with maximum index, then return the Venue of that row.
$z_2: \mathbf{R}[\text{Venue}].\text{Index.max}(\mathbf{R}[\text{Index}].\text{Position.1st})$ Find the maximum index of rows with Position = 1st, then return the Venue of the row with that index.
$z_3: \mathbf{R}[\text{Venue}].\text{argmax}(\text{Position.Number.1},$ $\mathbf{R}[\lambda x. \mathbf{R}[\text{Date}].\mathbf{R}[\text{Year}].x])$ Among rows with Position number 1, pick one with latest date in the Year column and return the Venue.
Spurious
$z_4: \mathbf{R}[\text{Venue}].\text{argmax}(\text{Position.Number.1},$ $\mathbf{R}[\lambda x. \mathbf{R}[\text{Number}].\mathbf{R}[\text{Time}].x])$ Among rows with Position number 1, pick the one with maximum Time number. Return the Venue.
$z_5: \mathbf{R}[\text{Venue}].\text{Year.Number.(}$ $\mathbf{R}[\text{Number}].\mathbf{R}[\text{Year}].\text{argmax}(\text{Type.Row}, \text{Index}) - 1)$ Subtract 1 from the Year in the last row, then return the Venue of the row with that Year.
Inconsistent
$\tilde{z}: \mathbf{R}[\text{Venue}].\text{argmin}(\text{Position.1st}, \text{Index})$ Among rows with Position = 1st, pick the one with minimum index, then return the Venue. (= Finland)

Solution: Generating fictitious data

Year	Venue	Position	Event	Time
2001	Hungary	2nd	400m	47.12
2003	Finland	1st	400m	46.69
2005	Germany	11th	400m	46.62
2007	Thailand	1st	relay	182.05
2008	China	7th	relay	180.32

x : “Where did the last 1st place finish occur?”

y : Thailand

Year	Venue	Position	Event	Time
2001	Finland	7th	relay	46.62
2003	Germany	1st	400m	180.32
2005	China	1st	relay	47.12
2007	Hungary	7th	relay	182.05

	w	w_1	w_2	\dots
z_1	Thailand	China	Finland	\dots
z_2	Thailand	China	Finland	$\dots \} q_1$
z_3	Thailand	China	Finland	$\dots \} q_2$
z_4	Thailand	Germany	China	$\dots \} q_3$
z_5	Thailand	China	China	
z_6	Thailand	China	China	
:	:	:	:	

Figure 4: From the example in Figure 1, we generate a table for the fictitious world w_1 .

- Equivalent classes
 - $q_1 = \{z_1, z_2, z_3\}$
 - $q_2 = \{z_4\}$
 - $q_3 = \{z_5, z_6\}$
- Choose worlds with most information gain

Outline

- Introduction
- Symbolic execution
 - (Fully supervised) Neural programmer interpreter
 - (Weakly supervised) Neural symbolic machine
 - “Spurious programs” and inductive programming
 - Learning semantic parsers from denotations
 - DeepCoder
 - More thoughts on “spurious programs”
- Distributed execution
 - Learning to execute
 - Neural enquirer
 - Differentiable neural computer
- Hybrid execution
 - Neural programmer
 - Coupling approach

DEEPCODER: LEARNING TO WRITE PROGRAMS

Matej Balog*

Department of Engineering
University of Cambridge

Alexander L. Gaunt, Marc Brockschmidt,

Sebastian Nowozin, Daniel Tarlow

Microsoft Research

```

a ← [int]
b ← FILTER (<0) a
c ← MAP (*4) b
d ← SORT c
e ← REVERSE d

```

An input-output example:

Input:

[-17, -3, 4, 11, 0, -5, -9, 13, 6, 6, -8, 11]

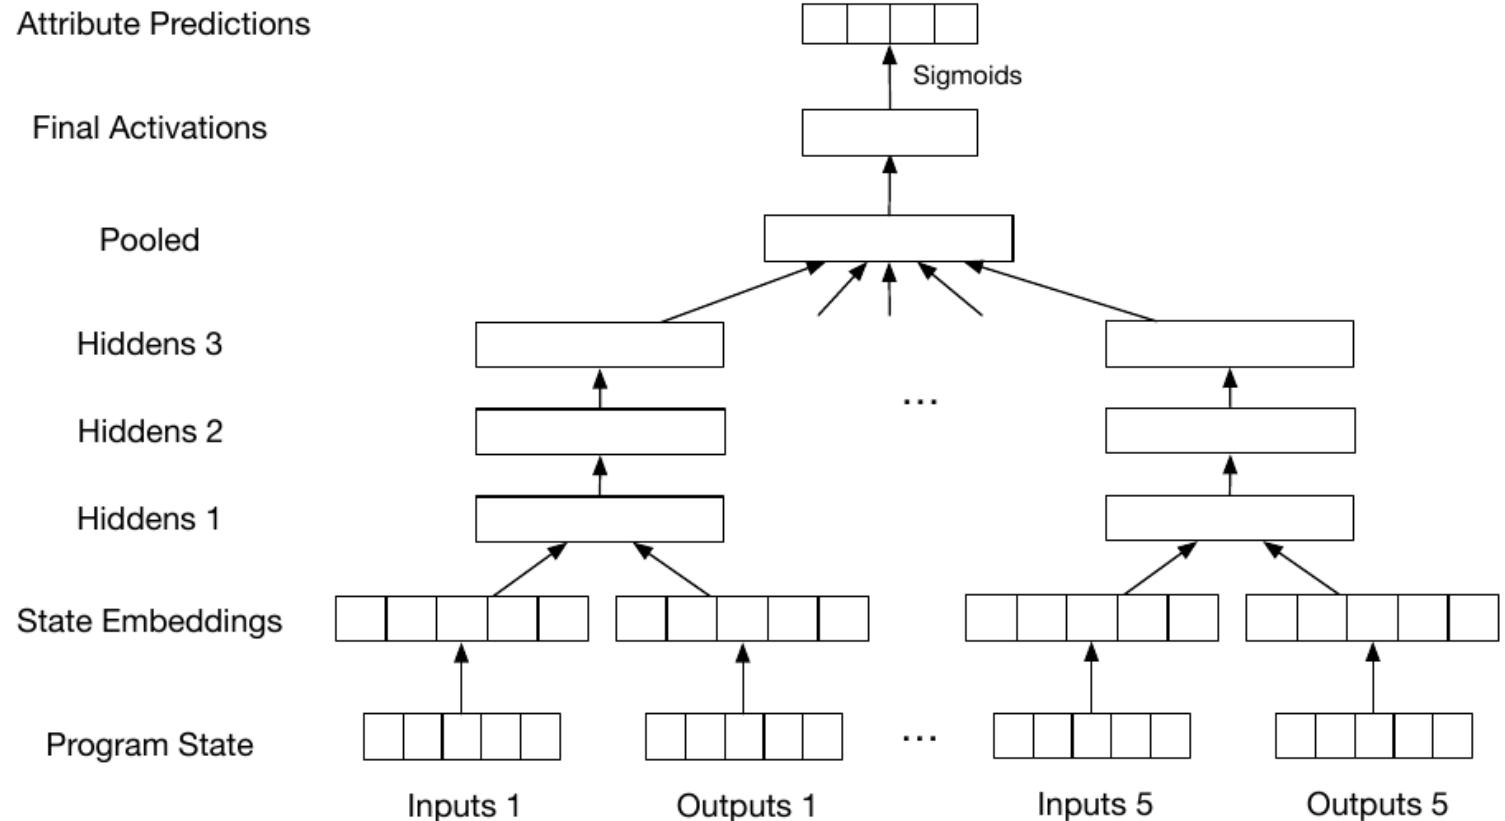
Output:

[-12, -20, -32, -36, -68]

Figure 1: An example program in our DSL that takes a single integer array as its input.

(+1)	(-1)	(*2)	(/2)	(*-1)	(**2)	(*3)	(/3)	(*4)	(/4)	(>0)	(>0)	(%2==1)	(%2==0)	HEAD	LAST	MAP	FILTER	SORT	REVERSE	TAKE	DROP	ACCESS	ZIPWITH	SCANL1	+ . * . MIN MAX COUNT MINIMUM MAXIMUM SUM									
.0	.0	.1	.0	.0	.0	.0	.0	1.0	.0	.0	1.0	.0	.2	.0	.0	1.0	1.0	1.0	.7	.0	.1	.0	.4	.0	.0	.1	.0	.2	.1	.0	.0	.0	.0	.0

Figure 2: Neural network predicts the probability of each function appearing in the source code.



Search

- DFS
 - DFS can opt to consider the functions as ordered by their predicted probabilities from the neural network
- “Sort and add” enumeration
- Sketch
- Not really “structured” prediction

Outline

- Introduction
- Symbolic execution
 - (Fully supervised) Neural programmer interpreter
 - (Weakly supervised) Neural symbolic machine
 - “Spurious programs” and inductive programming
 - Learning semantic parsers from denotations
 - DeepCoder
 - More thoughts on “spurious programs”
- Distributed execution
 - Learning to execute
 - Neural enquirer
 - Differentiable neural computer
- Hybrid execution
 - Neural programmer
 - Coupling approach

Outline

- Introduction
- Symbolic execution
 - (Fully supervised) Neural programmer interpreter
 - (Weakly supervised) Neural symbolic machine
 - “Spurious programs” and inductive programming
 - Learning semantic parsers from denotations
 - DeepCoder
 - More thoughts on “spurious programs”
- Distributed execution
- Hybrid execution

LEARNING TO EXECUTE

Wojciech Zaremba*

New York University

woj.zaremba@gmail.com

Essentially a
seq2seq model

Ilya Sutskever

Google

ilyasu@google.com

Input:

```
j=8584
for x in range(8):
    j+=920
b=(1500+j)
print((b+7567))
```

Target: 25011.

Input:

```
i=8827
c=(i-5347)
print((c+8704) if 2641<8500 else 5308)
```

Target: 12184.

Outline

- Introduction
- Symbolic execution
 - (Fully supervised) Neural programmer interpreter
 - (Weakly supervised) Neural symbolic machine
 - “Spurious programs” and inductive programming
 - Learning semantic parsers from denotations
 - DeepCoder
 - More thoughts on “spurious programs”
- Distributed execution
 - Learning to execute
 - Neural enquirer
 - Differentiable neural computer
- Hybrid execution
 - Neural programmer
 - Coupling approach

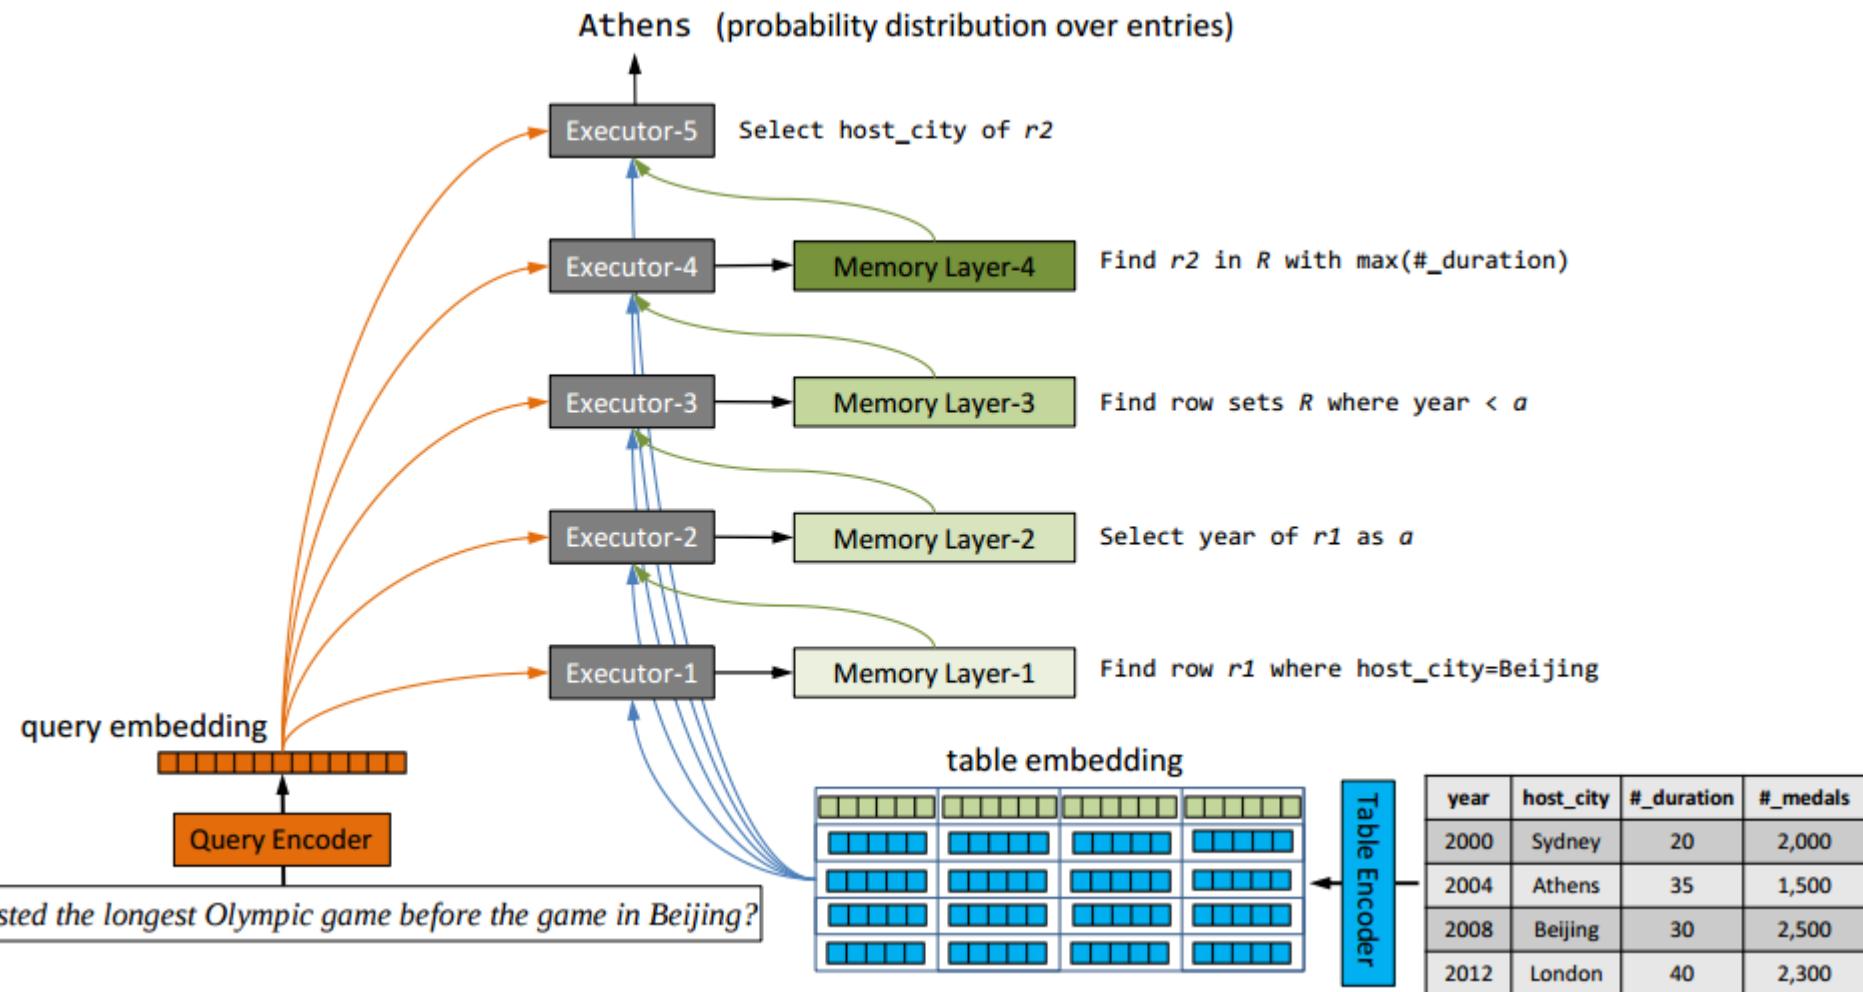
Neural Enquirer: Learning to Query Tables in Natural Language

Pengcheng Yin¹ Zhengdong Lu² Hang Li² Ben Kao¹

¹ Department of Computer Science, The University of Hong Kong

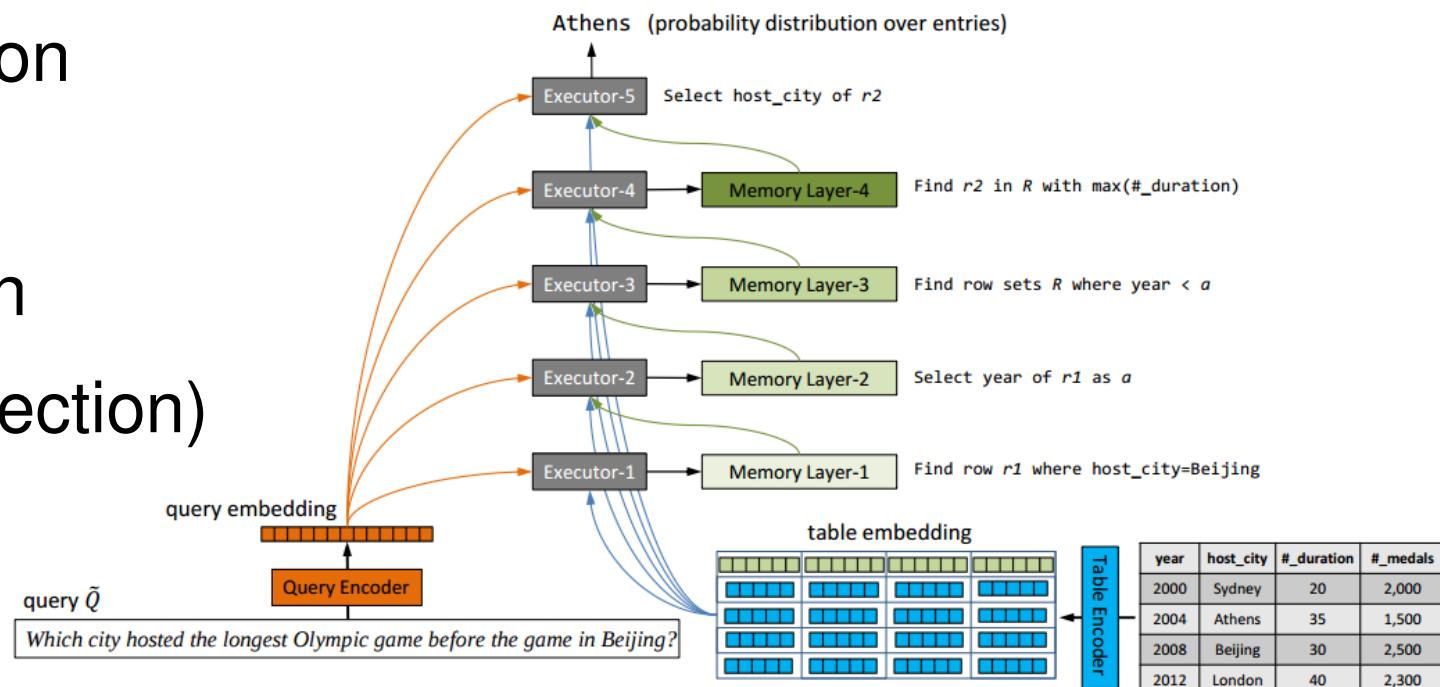
² Noah's Ark Lab, Huawei Technologies

¹{pcyin,kao}@cs.hku.hk ²{Lu.Zhengdong,HangLi.HL}@huawei.com



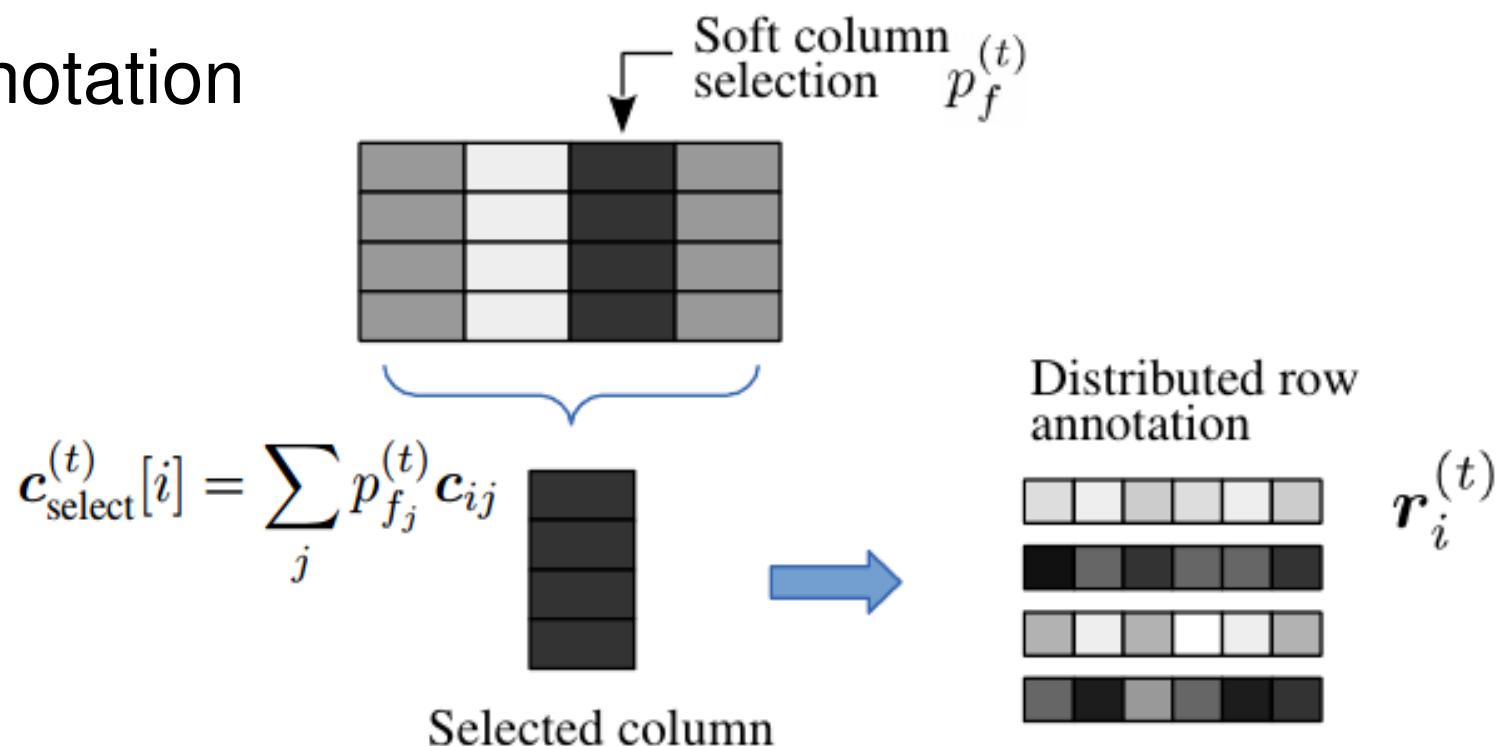
Three Modules

- Query encoder: BiRNN
- Table encoder
 - Concatenation of cell and field embeddings
 - Further processed by a multi-layer perceptron
- Executor
 - Column attention
(soft selection)
 - Row annotation
(distributed selection)



Executor

- The result of one-step execution softmax attention over columns and a distributed
 - Column attention
 - Row annotation



Details

- Let $\mathbf{r}_i^{(t-1)}$ be the previous step's row annotation results, where the subscript i indexes a particular row.

- Last step's execution information

$$\mathbf{g}^{(t-1)} = \text{MaxPool}_i \left\{ \mathbf{r}_i^{(t-1)} \right\}$$

- Current step

- Column attention $p_{f_j}^{(t)} = \text{softmax} \left(\text{MLP} \left([\mathbf{q}; \mathbf{f}_j; \mathbf{g}^{(t-1)}] \right) \right)$

- Row annotation

$$\mathbf{c}_{\text{select}}^{(t)}[i] = \sum_j p_{f_j}^{(t)} \mathbf{c}_{ij}$$

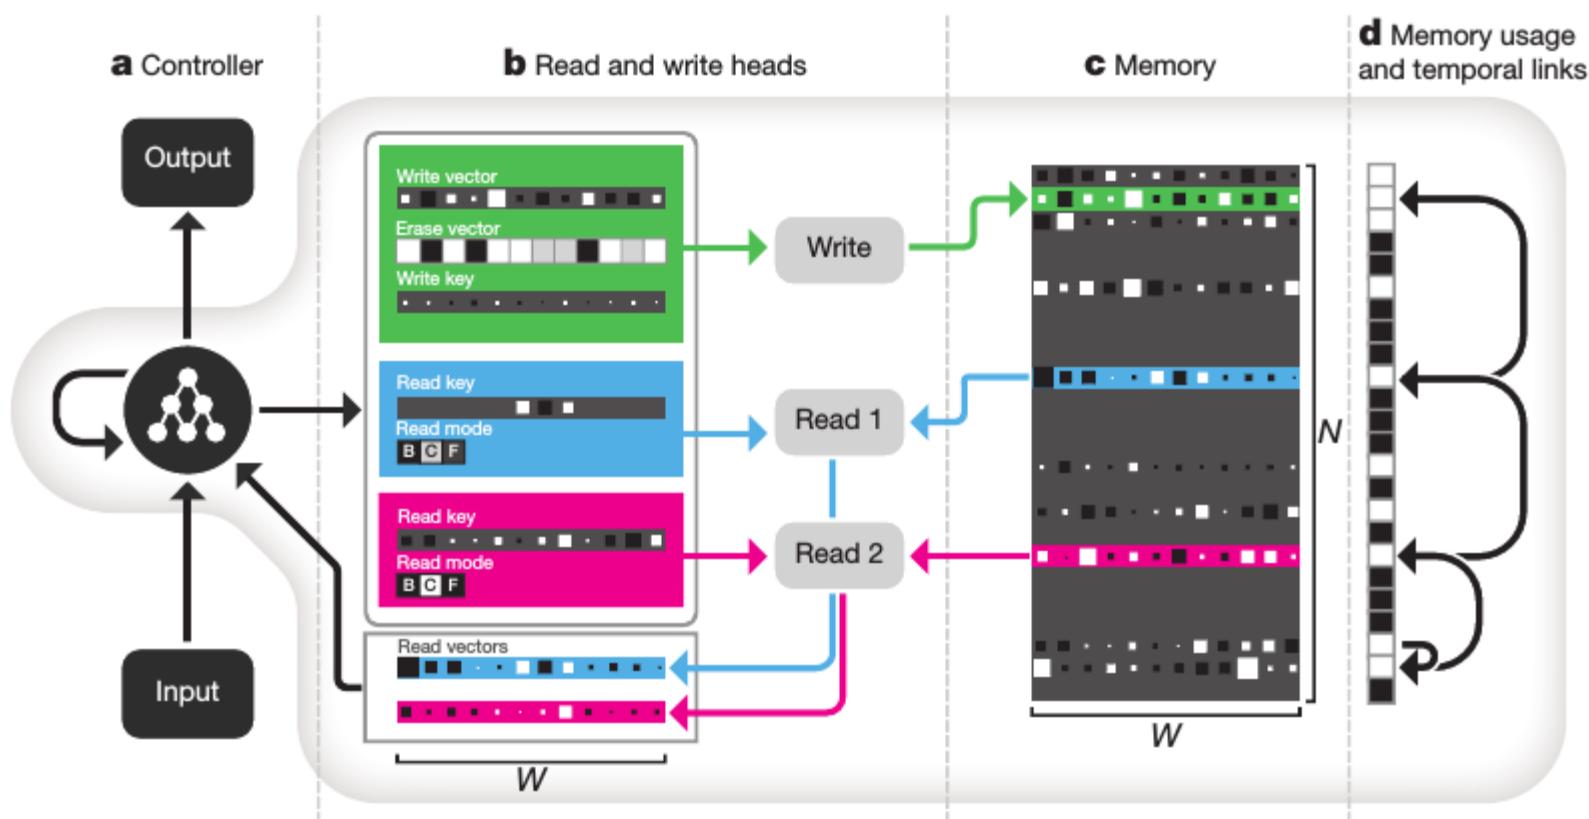
$$\mathbf{r}_i^{(t)} = \text{MLP} \left(\left[\mathbf{q}, \mathbf{g}^{(t-1)}, \mathbf{r}^{(t-1)}, \mathbf{c}_{\text{select}}^{(t)}[i] \right] \right)$$

Outline

- Introduction
- Symbolic execution
 - (Fully supervised) Neural programmer interpreter
 - (Weakly supervised) Neural symbolic machine
 - “Spurious programs” and inductive programming
 - Learning semantic parsers from denotations
 - DeepCoder
 - More thoughts on “spurious programs”
- Distributed execution
 - Learning to execute
 - Neural enquirer
 - Differentiable neural computer
- Hybrid execution
 - Neural programmer
 - Coupling approach

Hybrid computing using a neural network with dynamic external memory

Alex Graves^{1*}, Greg Wayne^{1*}, Malcolm Reynolds¹, Tim Harley¹, Ivo Danihelka¹, Agnieszka Grabska-Barwińska¹, Sergio Gómez Colmenarejo¹, Edward Grefenstette¹, Tiago Ramalho¹, John Agapiou¹, Adrià Puigdomènech Badia¹, Karl Moritz Hermann¹, Yori Zwols¹, Georg Ostrovski¹, Adam Cain¹, Helen King¹, Christopher Summerfield¹, Phil Blunsom¹, Koray Kavukcuoglu¹ & Demis Hassabis¹



Memory and its Access

- Memory: $M \in \mathbb{R}^{N \times W}$

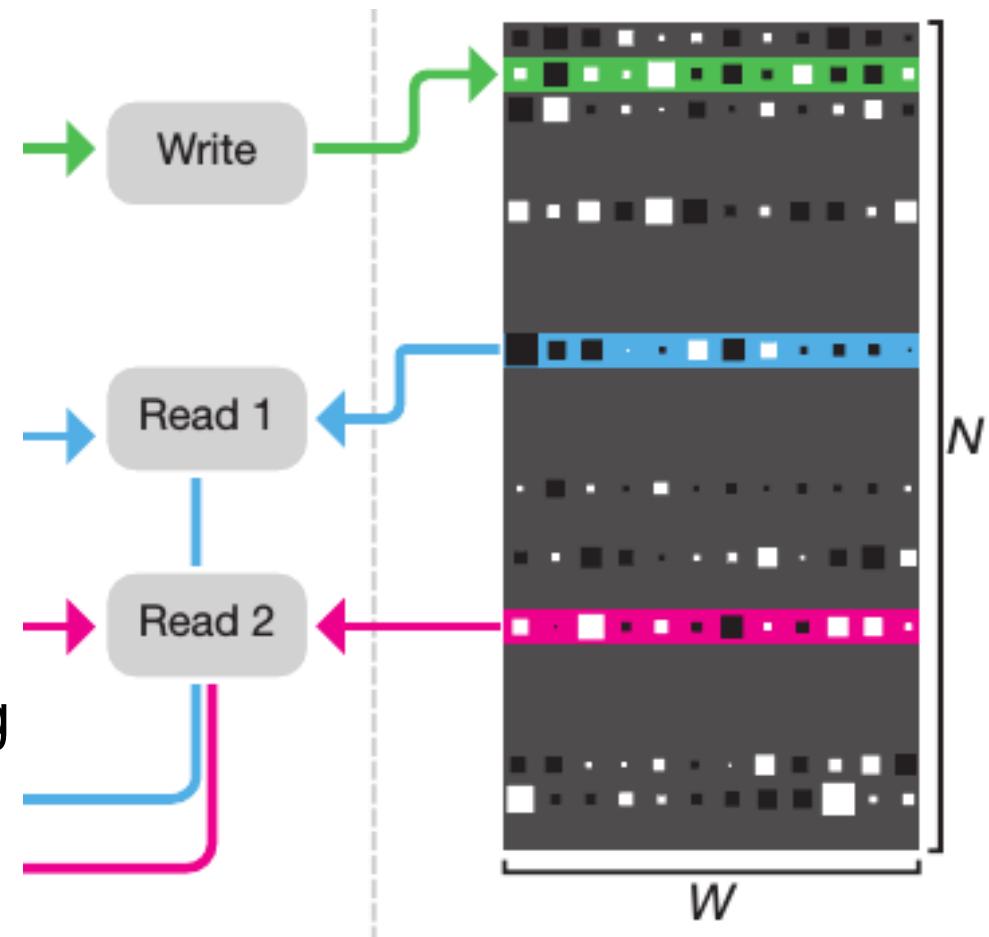
- N : # of memory slots
 - W : Width/dim of a slot

- Reader:

- R readers
 - Content addressing
 - + Temporal linking addressing

- Writer:

- Exact 1 writer
 - Content addressing + dynamic memory allocation



Architecture

- Input
 - Current signal
 - Previous step's memory read
- Output

- Value

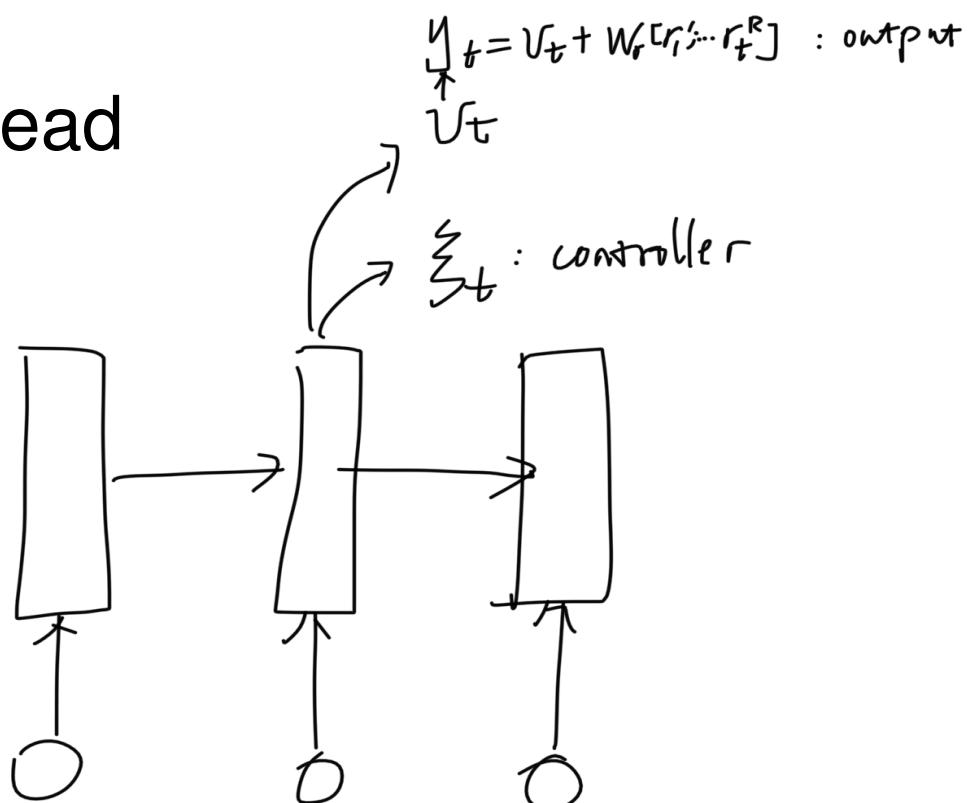
$$\mathbf{v}_t = W_v[\mathbf{h}_t^1; \dots; \mathbf{h}_t^L]$$

$$\mathbf{y}_t = \mathbf{v}_t + W_r[\mathbf{r}_t^1; \dots; \mathbf{r}_t^R]$$

(so that output is dependent on read memory)

- Control $\xi_t = W_\xi[\mathbf{h}_t^1; \dots; \mathbf{h}_t^L]$

$$\xi_t = [k_t^{r,1}; \dots; k_t^{r,R}; \hat{\beta}_t^{r,1}; \dots; \hat{\beta}_t^{r,R}; k_t^w; \hat{\beta}_t^w; \hat{\epsilon}_t; v_t; \hat{f}_t^1; \dots; \hat{f}_t^R; \hat{g}_t^a; \hat{g}_t^w; \hat{\pi}_t^1; \dots; \hat{\pi}_t^R]$$



$$x = [x_t; r_{t-1}^1; \dots; r_{t-1}^R]$$

Controller: Preliminary

- To obtain a gate: sigmoid
- To obtain a distribution: softmax
- To constrain in [1, infinity)

$$\text{oneplus}(x) = 1 + \log(1 + e^x)$$

- A distribution

$$\mathcal{S}_N = \left\{ \boldsymbol{\alpha} \in \mathbb{R}^N : \alpha_i \in [0,1], \sum_{i=1}^N \alpha_i = 1 \right\}$$

- No more than a distribution (detracted by a default null prob.)

$$\Delta_N = \left\{ \boldsymbol{\alpha} \in \mathbb{R}^N : \alpha_i \in [0,1], \sum_{i=1}^N \alpha_i \leq 1 \right\}$$

A Glance at the Control Signals

- R read keys $\{\mathbf{k}_t^{r,i} \in \mathbb{R}^W; 1 \leq i \leq R\}$;
- R read strengths $\{\beta_t^{r,i} = \text{oneplus}(\hat{\beta}_t^{r,i}) \in [1, \infty); 1 \leq i \leq R\}$;
- the write key $\mathbf{k}_t^w \in \mathbb{R}^W$;
- the write strength $\beta_t^w = \text{oneplus}(\hat{\beta}_t^w) \in [1, \infty)$;
- the erase vector $\mathbf{e}_t = \sigma(\hat{\mathbf{e}}_t) \in [0,1]^W$;
- the write vector $\mathbf{v}_t \in \mathbb{R}^W$;
- R free gates $\{f_t^i = \sigma(\hat{f}_t^i) \in [0,1]; 1 \leq i \leq R\}$;
- the allocation gate $g_t^a = \sigma(\hat{g}_t^a) \in [0,1]$;
- the write gate $g_t^w = \sigma(\hat{g}_t^w) \in [0,1]$; and
- R read modes $\{\boldsymbol{\pi}_t^i = \text{softmax}(\hat{\boldsymbol{\pi}}_t^i) \in \mathcal{S}_3; 1 \leq i \leq R\}$.

Memory Reading and Writing

- Reader
 - Reading weights $\mathbf{w}_t^{r,1}, \dots, \mathbf{w}_t^{r,R} \in \Delta_N$
 - Read vectors $\mathbf{r}_t^i = M_t^\top \mathbf{w}_t^{r,i}$
where $M: < N \times W >$, each $\mathbf{w}: < N \times 1 >$
(no more than attention)
- Writer
 - Writing weight regarding a memory slot $\mathbf{w}_t^w \in \Delta_N$
 - Erasing vector regarding a memory dimension
(given by the controller) $\mathbf{e}_t = \sigma(\hat{\mathbf{e}}_t) \in [0,1]^W$
 - Writing operation $M_t = M_{t-1} \circ (E - \mathbf{w}_t^w \mathbf{e}_t^\top) + \mathbf{w}_t^w \mathbf{v}_t^\top$
E: 1's | w → 1 & e → 1 ==> M is determined by v

Memory Addressing

- Content-based addressing
 - e.g., select 5
- Dynamic memory allocation
 - Free(.), Malloc(.) and write to it
- Temporal memory linkage
 - Move to the next element in a linked list
- C.f. CopyNet, Latent Predictor Network

Content-Based Addressing

$$\mathcal{C}(M, \mathbf{k}, \beta)[i] = \frac{\exp\{\mathcal{D}(\mathbf{k}, M[i, \cdot])\beta\}}{\sum_j \exp\{\mathcal{D}(\mathbf{k}, M[j, \cdot])\beta\}} \quad \text{Size: } \langle N \rangle$$

\mathbf{k} and beta given by controller

$$\mathcal{D}(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u} \cdot \mathbf{v}}{|\mathbf{u}| |\mathbf{v}|}$$

- R read keys $\{\mathbf{k}_t^{\text{r},i} \in \mathbb{R}^W; 1 \leq i \leq R\}$;
- R read strengths $\{\beta_t^{\text{r},i} = \text{oneplus}(\hat{\beta}_t^{\text{r},i}) \in [1, \infty); 1 \leq i \leq R\}$;
- the write key $\mathbf{k}_t^{\text{w}} \in \mathbb{R}^W$;
- the write strength $\beta_t^{\text{w}} = \text{oneplus}(\hat{\beta}_t^{\text{w}}) \in [1, \infty)$;

Dynamic Memory Allocation

- Memory usage vector $\mathbf{u}_t \in [0, 1]^N$

1=used, 0=unused $\mathbf{u}_0 = \mathbf{0}$

- Memory retention vector $\psi_t \in [0, 1]^N$

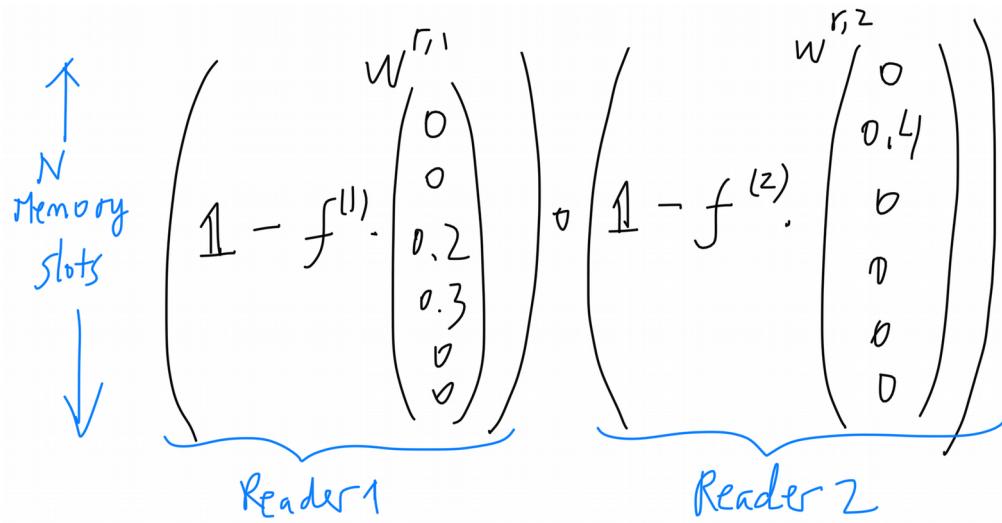
how much each location will not be freed by the free gates

$$\psi_t = \prod_{i=1}^R \left(1 - f_t^i \mathbf{w}_{t-1}^{r,i} \right)$$

where free gates are given by the controller

$$\{f_t^i = \sigma(\hat{f}_t^i) \in [0,1]; 1 \leq i \leq R\}$$

Free is in the sense of “malloc and free,” rather than “free-style.”



Read more => retain less
 f larger => retrain less

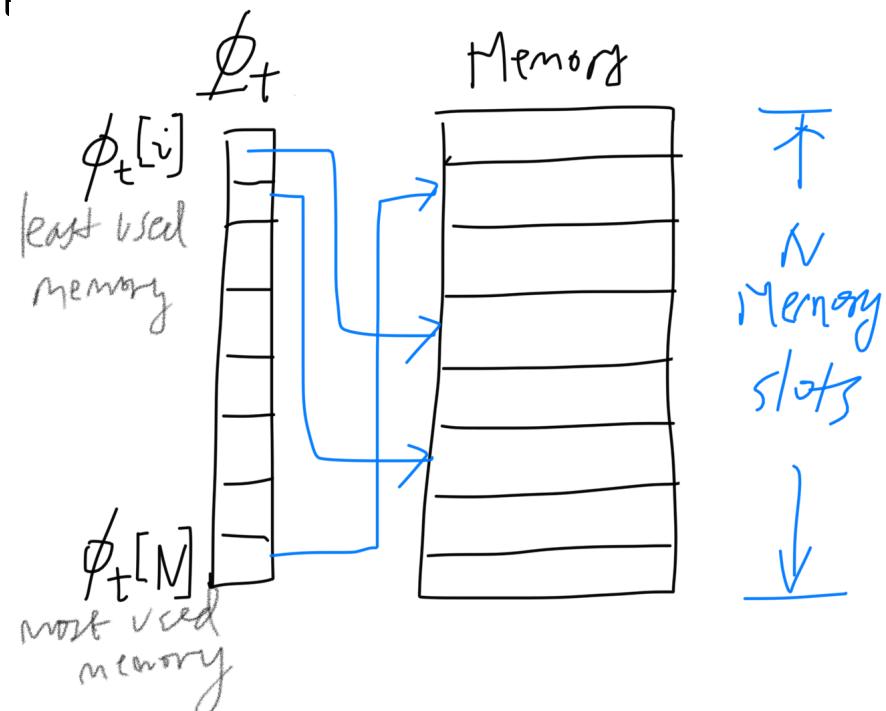
Dynamic Memory Allocation

- Memory usage vector recursion

- 1=used, 0=unused $\mathbf{u}_0 = \mathbf{0}$
- $\mathbf{u}_t = (\mathbf{u}_{t-1} + \mathbf{w}_{t-1}^W - \mathbf{u}_{t-1} \circ \mathbf{w}_{t-1}^W) \circ \psi_t$
- $\mathbf{u} = [\mathbf{u} * (1 - \mathbf{w}) + 1 * \mathbf{w}] * \text{psi}$
- \mathbf{w}_{t-1}^W : last step's writing weight
- write more => use more

- Free list $\phi_t \in \mathbb{Z}^N$

- Sorted list of memory usage
- $\phi_t[1]$: least used memory index



Dynamic Memory Allocation

- Allocation weighting

$$\alpha_t[\phi_t[j]] = (1 - \mathbf{u}_t[\phi_t[j]]) \prod_{i=1}^{j-1} \mathbf{u}_t[\phi_t[i]]$$

- $1 - \mathbf{u}$: used more => allocate less
- $\phi_t \in \mathbb{Z}^N$: sorted list of memory usage
- $\prod_{i=1}^{j-1} \mathbf{u}_t[\phi_t[i]]$ least used=> allocate more

Write Weighting

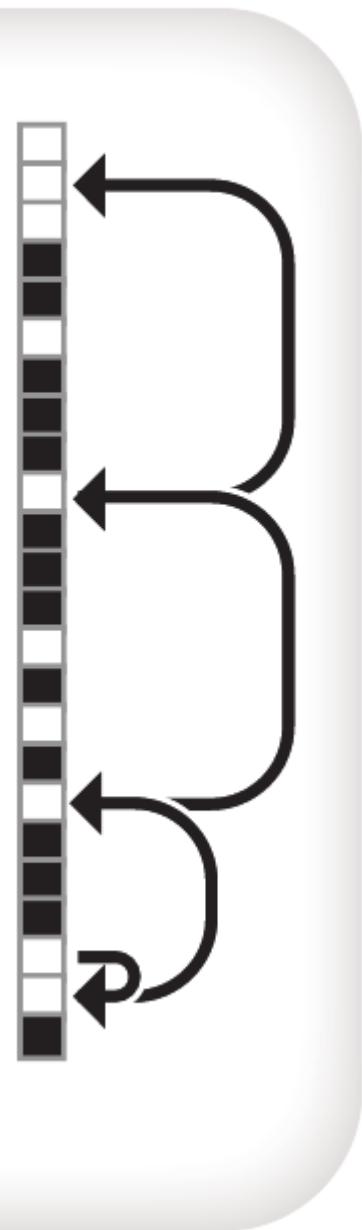
$$\mathbf{w}_t^w = g_t^w \left[g_t^a \mathbf{a}_t + (1 - g_t^a) \mathbf{c}_t^w \right]$$

- g^w : writing gate, deciding to write or not
- g^a : memory allocation gate for write weighting
- Both given by the controller

the allocation gate $g_t^a = \sigma(\hat{g}_t^a) \in [0,1]$;

the write gate $g_t^w = \sigma(\hat{g}_t^w) \in [0,1]$;

d Memory usage
and temporal links



Read Weighting

$$\mathbf{w}_t^{\text{r},i} = \boldsymbol{\pi}_t^i[1]\mathbf{b}_t^i + \boldsymbol{\pi}_t^i[2]\mathbf{c}_t^{\text{r},i} + \boldsymbol{\pi}_t^i[3]\mathbf{f}_t^i$$

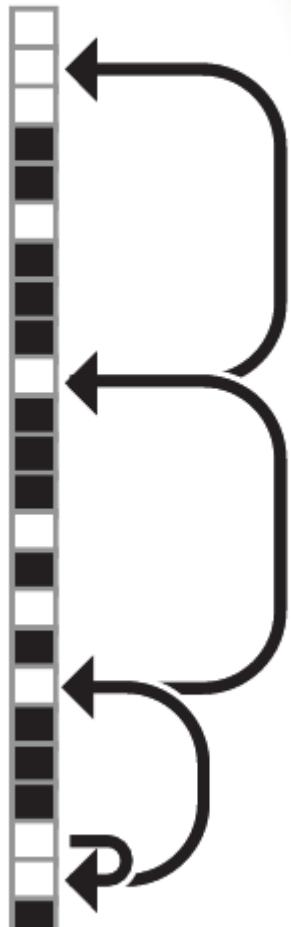
Temporal memory linkage



$$\mathbf{w}_t^{r,i} = \pi_t^i[1]\mathbf{b}_t^i + \pi_t^i[2]\mathbf{c}_t^{r,i} + \pi_t^i[3]\mathbf{f}_t^i$$

- Precedence weighting $\mathbf{p}_t \in \Delta_N$
$$\mathbf{p}_0 = \mathbf{0}$$
$$\mathbf{p}_t = \left(1 - \sum_i \mathbf{w}_t^w[i]\right) \mathbf{p}_{t-1} + \mathbf{w}_t^w$$
 - Degree to which a slot is lastly written
 - Current writing weight plus \mathbf{p}_{t-1}
 - \mathbf{p}_{t-1} discounted by summed writing weight

Temporal memory linkage



$$\mathbf{w}_t^{r,i} = \boldsymbol{\pi}_t^i[1]\mathbf{b}_t^i + \boldsymbol{\pi}_t^i[2]\mathbf{c}_t^{r,i} + \boldsymbol{\pi}_t^i[3]\mathbf{f}_t^i$$

- Linking matrix $L_t \in [0, 1]^{N \times N}$

$$L_0[i, j] = 0 \quad \forall i, j$$

$$L_t[i, i] = 0 \quad \forall i$$

$$L_t[i, j] = (1 - \mathbf{w}_t^w[i] - \mathbf{w}_t^w[j])L_{t-1}[i, j] + \mathbf{w}_t^w[i]\mathbf{p}_{t-1}[j]$$

- Degree to which location i is written after j
- slot i written at the current step
- slot j written at the last step
==> $L[i, j] \rightarrow 1$
- Slot i not written ==> $L[i, j]$ tends to remain
- Slot j written ==> $L[i, j] \rightarrow 0$, i.e., refresh

Temporal memory linkage



$$\mathbf{w}_t^{r,i} = \pi_t^i[1]\mathbf{b}_t^i + \pi_t^i[2]\mathbf{c}_t^{r,i} + \pi_t^i[3]\mathbf{f}_t^i$$

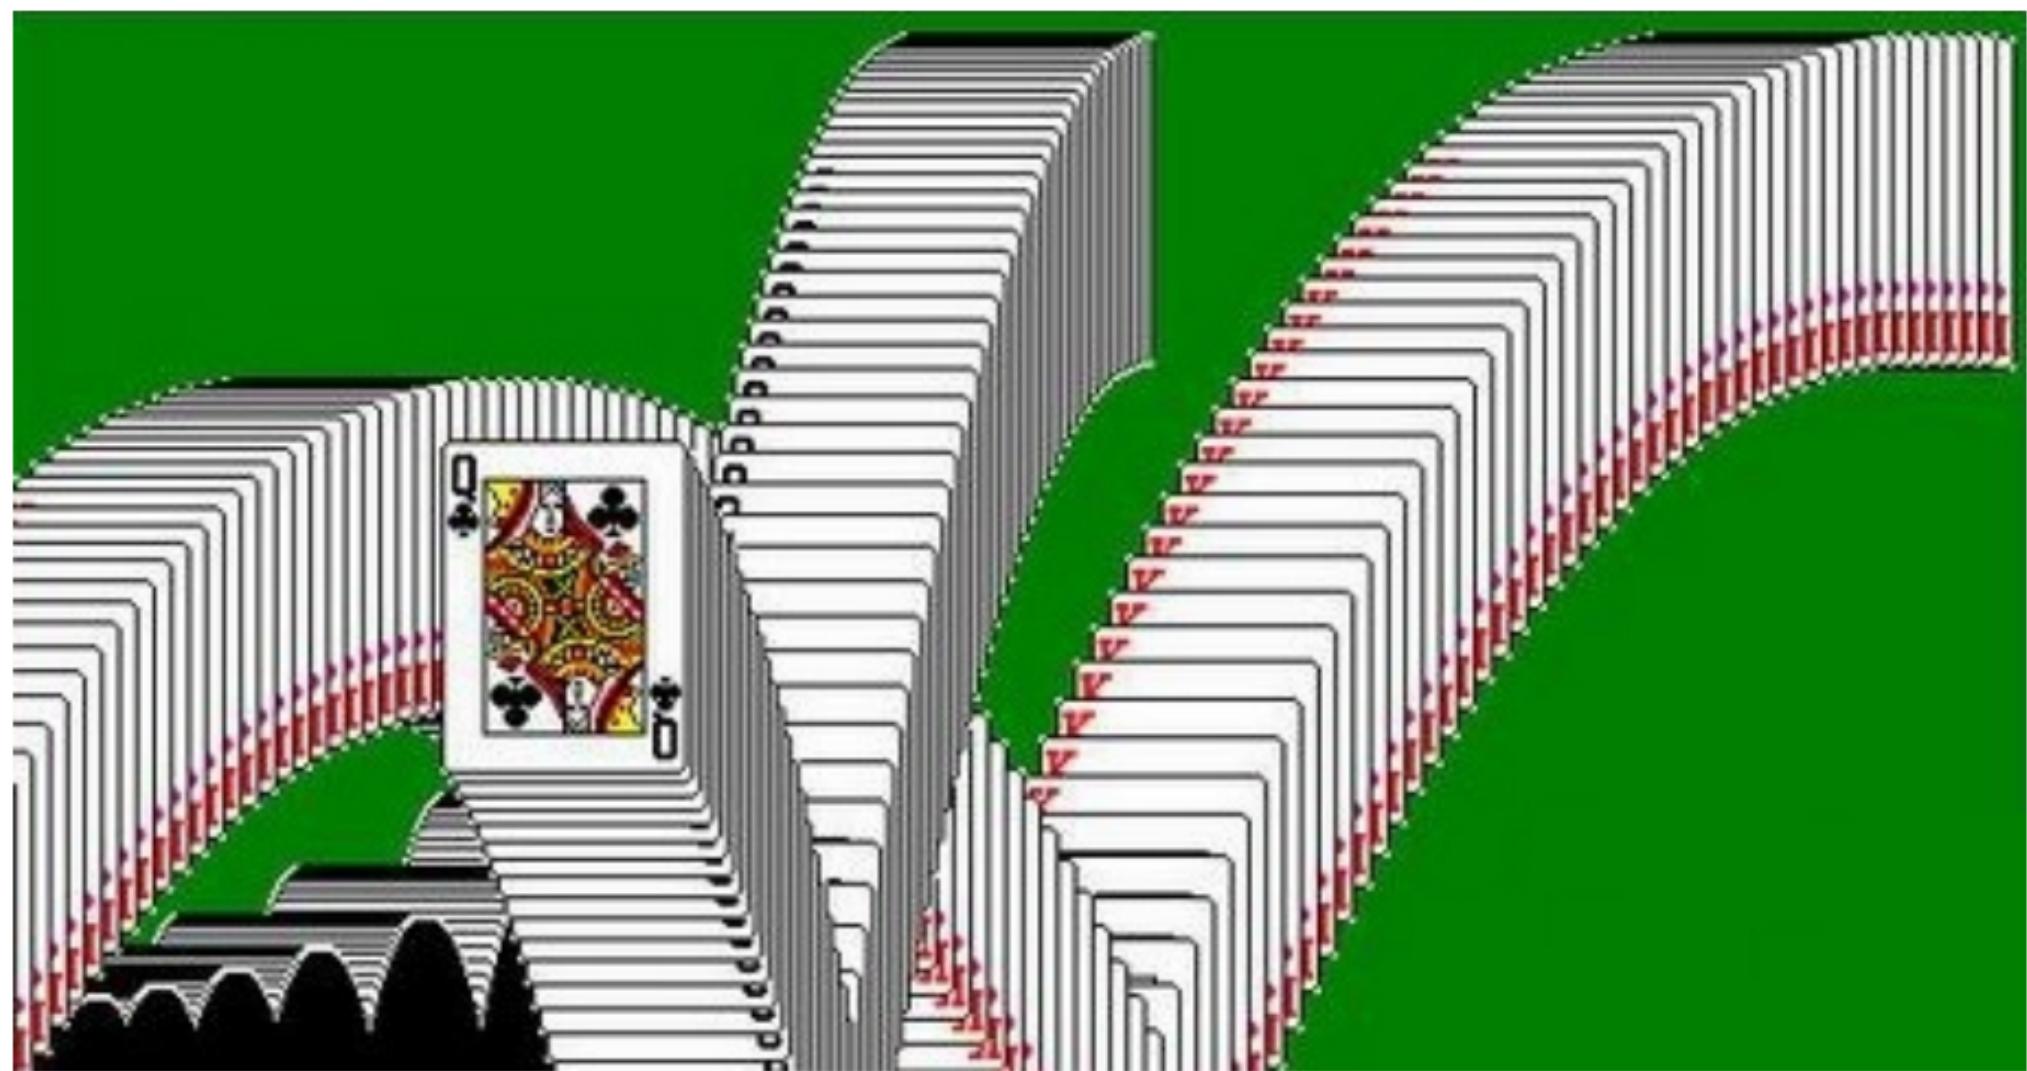
- Backward/forward weighting

$$\mathbf{b}_t^i \in \Delta_N \quad \mathbf{f}_t^i \in \Delta_N$$

$$\mathbf{f}_t^i = L_t \mathbf{w}_{t-1}^{r,i}$$

$$\mathbf{b}_t^i = L_t^\top \mathbf{w}_{t-1}^{r,i}$$

Linking weight * last step's reading weight



Outline

- Introduction
- Symbolic execution
 - (Fully supervised) Neural programmer interpreter
 - (Weakly supervised) Neural symbolic machine
 - “Spurious programs” and inductive programming
 - Learning semantic parsers from denotations
 - DeepCoder
 - More thoughts on “spurious programs”
- Distributed execution
 - Learning to execute
 - Neural enquirer
 - Differentiable neural computer
- Hybrid execution
 - Neural programmer
 - Coupling approach

NEURAL PROGRAMMER: INDUCING LATENT PROGRAMS WITH GRADIENT DESCENT

Arvind Neelakantan*

University of Massachusetts Amherst

arvind@cs.umass.edu

Quoc V. Le

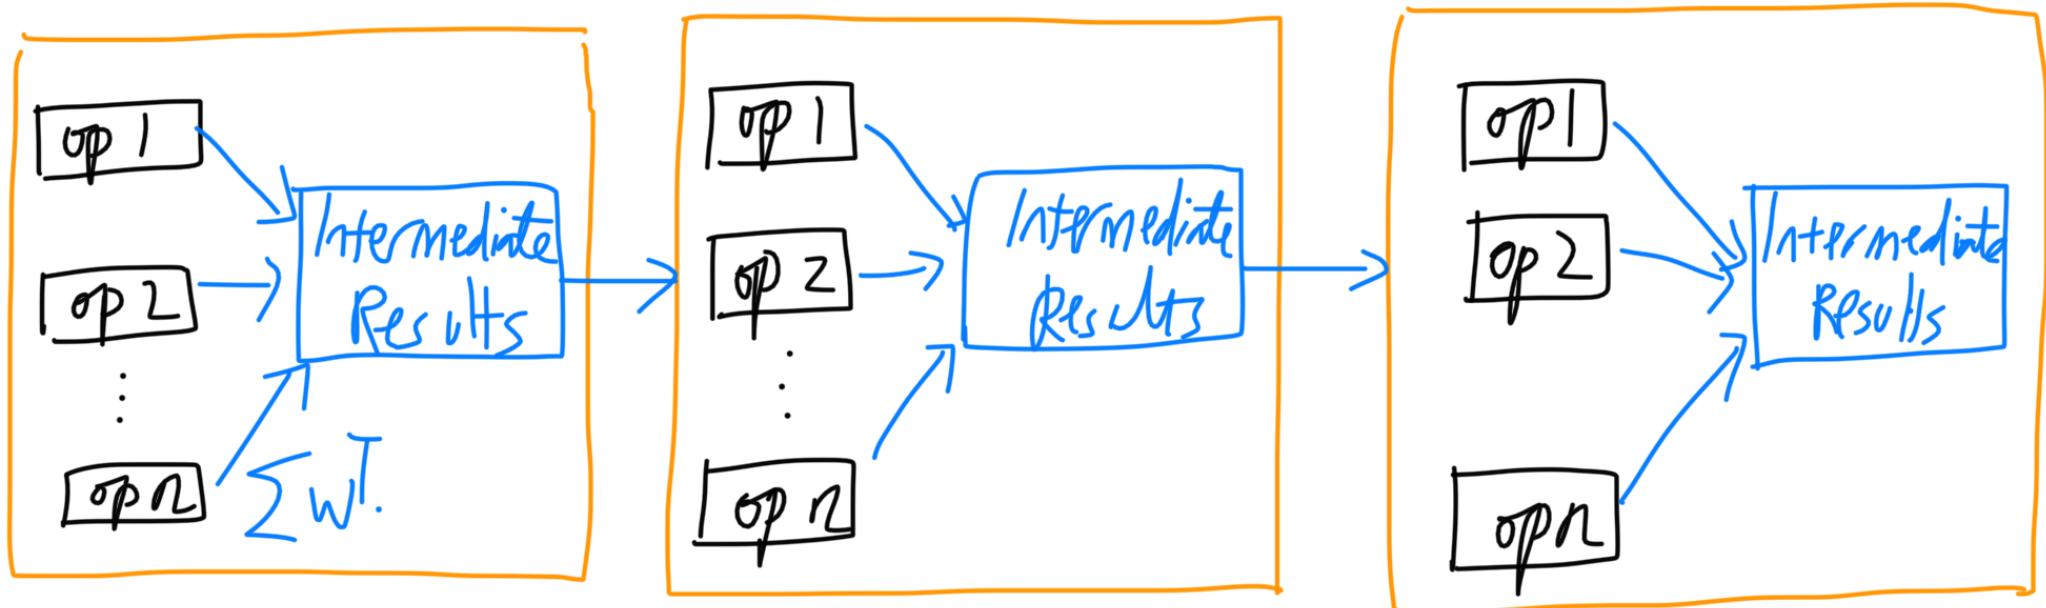
Google Brain

qvl@google.com

Ilya Sutskever

Google Brain

ilyasu@google.com

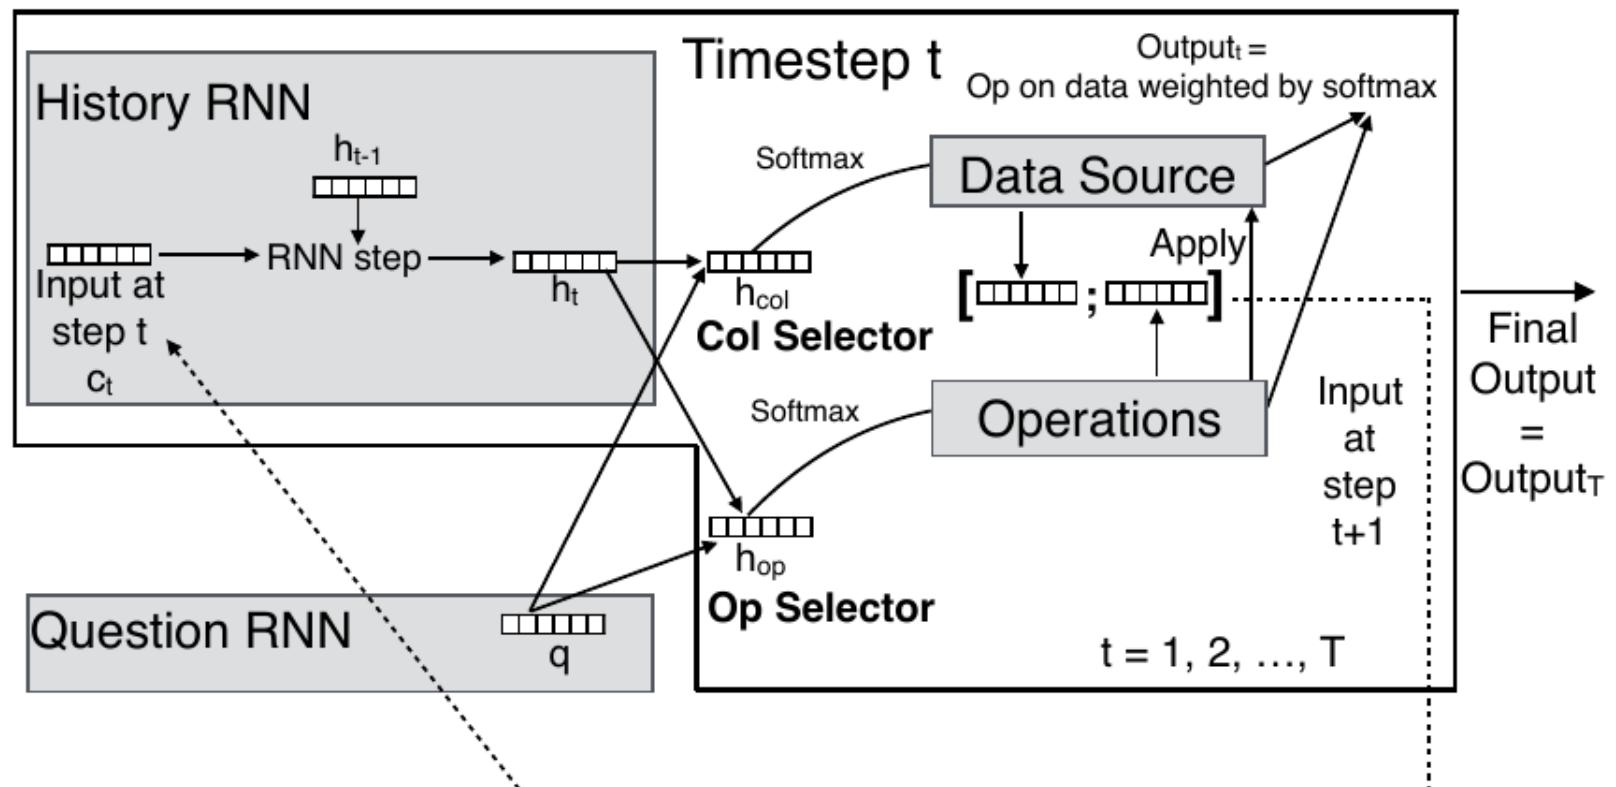


Idea

- Step-by-step fusion of different operators (and arguments)
- Trained by MSE
- Fully differentiable
- Drawbacks
 - Only numerical results (?)
 - Exponential # of combinatorial states

Arthitecture

- Question RNN
- Selector (Controller)
- Operators
- History RNN



Selector (Controller)

- Select an operator

$$\alpha_t^{op} = \text{softmax}(U \tanh(W^{op}[q; h_t]))$$

- Select a column for processing

$$\alpha_t^{col} = \text{softmax}(P \tanh(W^{col}[q; h_t]))$$

Operators

Type	Operation	Definition
Aggregate	Sum	$sum_t[j] = \sum_{i=1}^M row_select_{t-1}[i] * table[i][j], \forall j = 1, 2, \dots, C$
	Count	$count_t = \sum_{i=1}^M row_select_{t-1}[i]$
Arithmetic	Difference	$diff_t = scalar_output_{t-3} - scalar_output_{t-1}$
Comparison	Greater	$g_t[i][j] = table[i][j] > pivot_g, \forall (i, j), i = 1, \dots, M, j = 1, \dots, C$
	Lesser	$l_t[i][j] = table[i][j] < pivot_l, \forall (i, j), i = 1, \dots, M, j = 1, \dots, C$
Logic	And	$and_t[i] = min(row_select_{t-1}[i], row_select_{t-2}[i]), \forall i = 1, 2, \dots, M$
	Or	$or_t[i] = max(row_select_{t-1}[i], row_select_{t-2}[i]), \forall i = 1, 2, \dots, M$
Assign Lookup	assign	$assign_t[i][j] = row_select_{t-1}[i], \forall (i, j) i = 1, 2, \dots, M, j = 1, 2, \dots, C$
Reset	Reset	$reset_t[i] = 1, \forall i = 1, 2, \dots, M$

$$\beta_{op} = softmax(ZU(op))$$

$$pivot_{op} = \sum_{i=1}^N \beta_{op}(i) qn_i$$

Type	Operation	Definition
Aggregate	Sum	$sum_t[j] = \sum_{i=1}^M row_select_{t-1}[i] * table[i][j], \forall j = 1, 2, \dots, C$
	Count	$count_t = \sum_{i=1}^M row_select_{t-1}[i]$
Arithmetic	Difference	$diff_t = scalar_output_{t-3} - scalar_output_{t-1}$
Comparison	Greater	$g_t[i][j] = table[i][j] > pivot_g, \forall (i, j), i = 1, \dots, M, j = 1, \dots, C$
	Lesser	$l_t[i][j] = table[i][j] < pivot_l, \forall (i, j), i = 1, \dots, M, j = 1, \dots, C$
Logic	And	$and_t[i] = min(row_select_{t-1}[i], row_select_{t-2}[i]), \forall i = 1, 2, \dots, M$
	Or	$or_t[i] = max(row_select_{t-1}[i], row_select_{t-2}[i]), \forall i = 1, 2, \dots, M$
Assign Lookup	assign	$assign_t[i][j] = row_select_{t-1}[i], \forall (i, j) i = 1, 2, \dots, M, j = 1, 2, \dots, C$
Reset	Reset	$reset_t[i] = 1, \forall i = 1, 2, \dots, M$

More formally, the output variables are given by:

Execution

Output

$$scalar_answer_t = \alpha_t^{op}(\text{count}) count_t + \alpha_t^{op}(\text{difference}) diff_t + \sum_{j=1}^C \alpha_t^{col}(j) \alpha_t^{op}(\text{sum}) sum_t[j],$$

$$lookup_answer_t[i][j] = \alpha_t^{col}(j) \alpha_t^{op}(\text{assign}) assign_t[i][j], \forall (i, j) i = 1, 2, \dots, M, j = 1, 2, \dots, C$$

The row selector variable is given by:

$$row_select_t[i] = \alpha_t^{op}(\text{and}) and_t[i] + \alpha_t^{op}(\text{or}) or_t[i] + \alpha_t^{op}(\text{reset}) reset_t[i] +$$

$$\sum_{j=1}^C \alpha_t^{col}(j) (\alpha_t^{op}(\text{greater}) g_t[i][j] + \alpha_t^{op}(\text{lesser}) l_t[i][j]), \forall i = 1, \dots, M$$

Text Matching

- Example: what is the sum of elements in column B whose field in column C is **word:1** and field in column A is **word:7**?
- Text matching (no textual operation or output)

$$\begin{aligned} \text{row_select}_t[i] = & \alpha_t^{op}(\text{and})\text{and}_t[i] + \alpha_t^{op}(\text{or})\text{or}_t[i] + \alpha_t^{op}(\text{reset})\text{reset}_t[i] + \\ & \sum_{j=K+1}^C \alpha_t^{col}(j)(\alpha_t^{op}(\text{greater})g_t[i][j] + \alpha_t^{op}(\text{lesser})l_t[i][j]) + \\ & \sum_{j=1}^K \alpha_t^{col}(j)(\alpha_t^{op}(\text{text_match})\text{text_match}_t[i][j], \forall i = 1, \dots, M) \end{aligned}$$

Training Objective

- Scalar answer

$$L_{scalar}(scalar_answer_T, y) = \begin{cases} \frac{1}{2}a^2, & \text{if } a \leq \delta \\ \delta a - \frac{1}{2}\delta^2, & \text{otherwise} \end{cases}$$

- List answer

$$L_{lookup}(lookup_answer_T, y) = -\frac{1}{MC} \sum_{i=1}^M \sum_{j=1}^C \left(y[i, j] \log(lookup_answer_T[i, j]) + (1 - y[i, j]) \log(1 - lookup_answer_T[i, j]) \right)$$

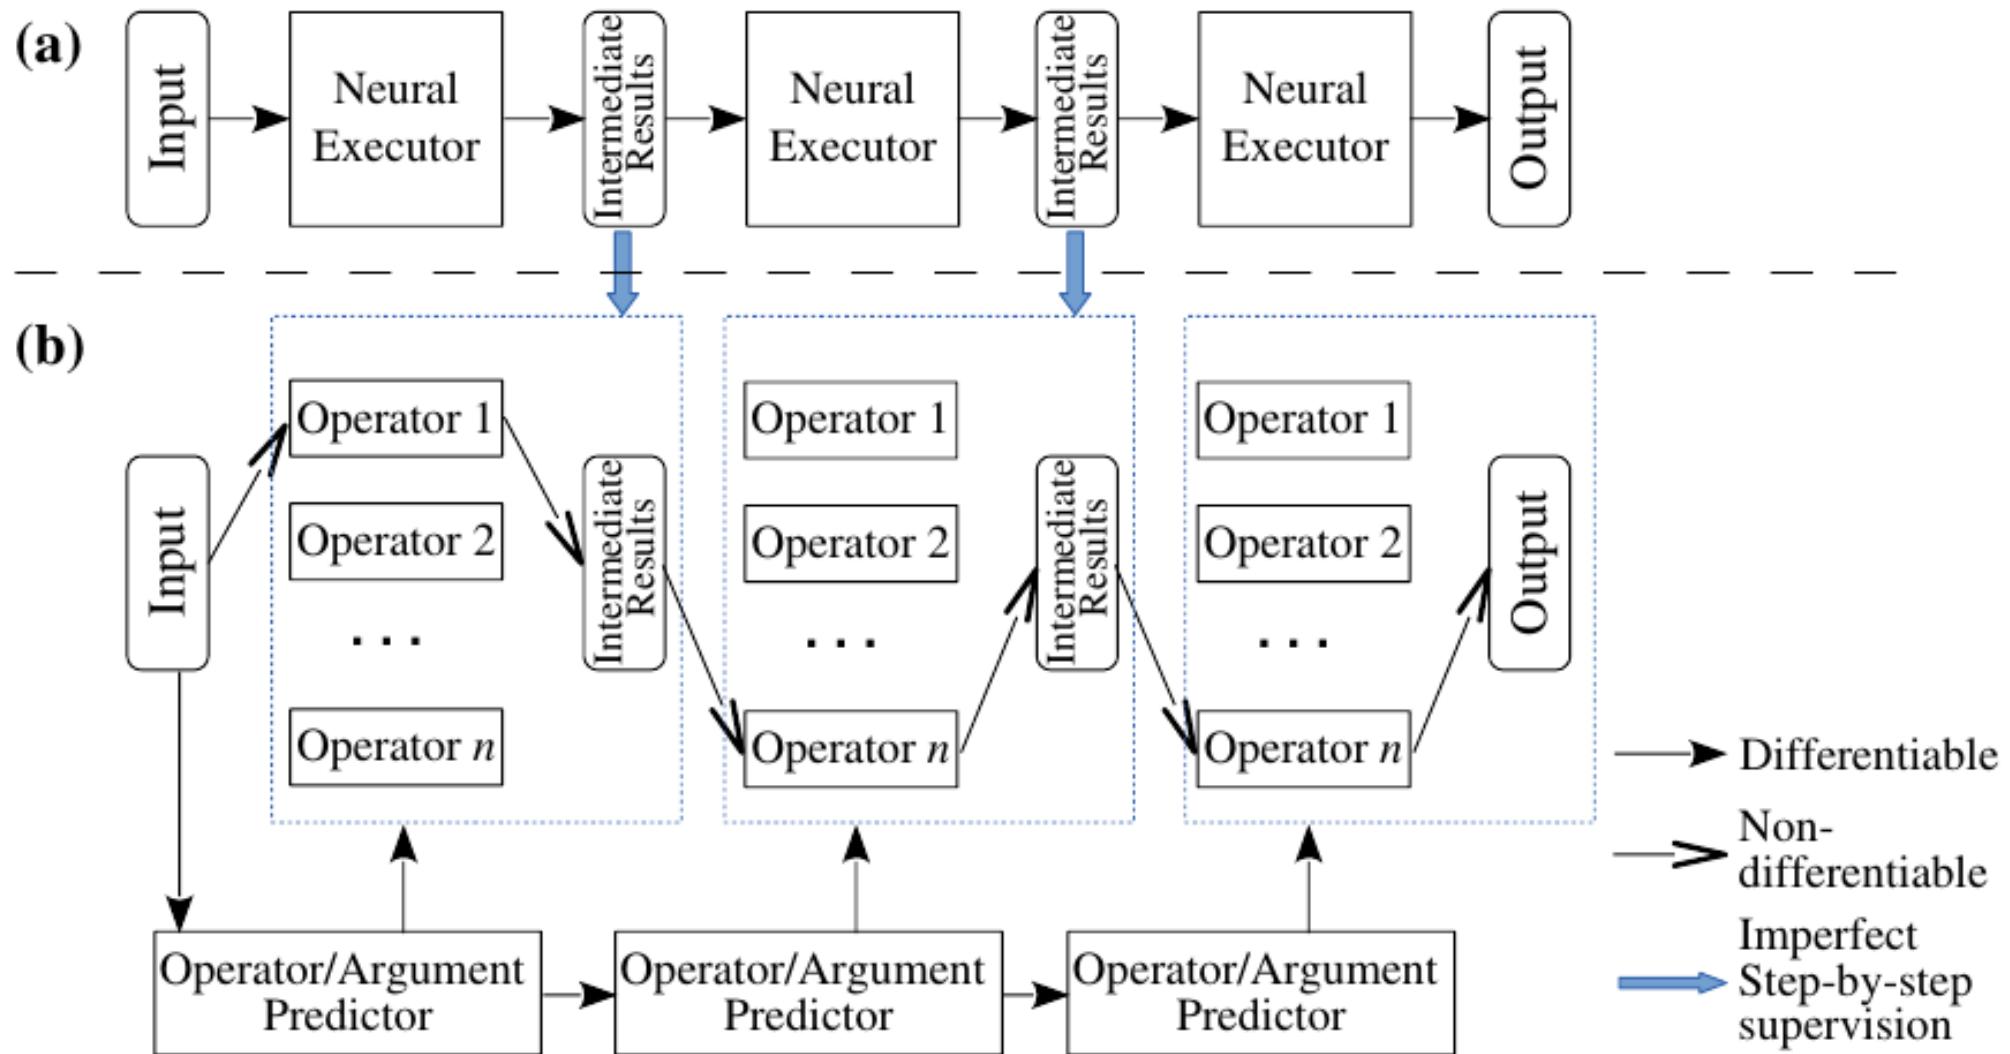
- Overall

$$L = \frac{1}{N} \sum_{k=1}^N \left([n_k == True] L_{scalar}^{(k)} + [n_k == False] \lambda L_{lookup}^{(k)} \right)$$

Outline

- Introduction
- Symbolic execution
 - (Fully supervised) Neural programmer interpreter
 - (Weakly supervised) Neural symbolic machine
 - “Spurious programs” and inductive programming
 - Learning semantic parsers from denotations
 - DeepCoder
 - More thoughts on “spurious programs”
- Distributed execution
 - Learning to execute
 - Neural enquirer
 - Differentiable neural computer
- Hybrid execution
 - Neural programmer
 - Coupling approach

Overview



Idea

- Fully neuralized model exhibits (imperfect) symbolic interpretation
- Use neural networks' intermediate results to learn initial policy of the symbolic executor
- Improve policy by REINFORCE or MML variants

Thank you for listening

Q & A?

References

1. Scott Reed, Nando de Freitas. "Neural programmer-interpreters." In *ICLR*, 2016.
2. Jonathon Cai, Richard Shin, Dawn Song. "Making neural programming architectures generalize via recursion." In *ICLR*, 2017.
3. Chen Liang, Jonathan Berant, Quoc Le, Kenneth D. Forbus, Ni Lao. "Neural symbolic machines: Learning semantic parsers on freebase with weak supervision." *arXiv preprint arXiv:1611.00020*, 2016.
4. Arun Chaganty, Percy Liang. "Inferring logical forms from denotations." In *ACL*, 2016.
5. Matej Balog, Alexander L. Gaunt, Marc Brockschmidt, Sebastian Nowozin, Daniel Tarlow. "DeepCoder: Learning to write programs." In *ICLR*, 2017.
6. Wojciech Zaremba, Ilya Sutskever. "Learning to execute." *arXiv preprint arXiv:1410.4615*, 2014.
7. Pengcheng Yin, Zhengdong Lu, Hang Li, Ben Kao. "Neural enquirer: Learning to query tables with natural language." In *IJCAI*, 2016.
8. Alex Graves, et al. "Hybrid computing using a neural network with dynamic external memory." *Nature*, 2016.
9. Arvind Neelakantan, Quoc Le, Ilya Sutskever. "Neural programmer: Inducing latent programs with gradient descent." In *ICLR*, 2016.
10. Lili Mou, Zhengdong Lu, Hang Li, Zhi Jin. "Coupling distributed and symbolic execution for natural language queries." *arXiv preprint arXiv:1612.02741*, 2016.