

hmm

Lili Mou
l mou@ualberta.ca
lili-mou.github.io



UNIVERSITY OF
ALBERTA

Drawbacks of LR/Softmax

- Classification is non-linear
 - May not even represented as fixed-dimensional features
- Do not consider the relationship of labels within one data sample



The lecture is really boring
 determiner ? verb adverb adjective

Three professors lecture IntroNLP

CardinalNumber Noun ? ProperNoun

<https://www.merriam-webster.com/dictionary/lecture>

lecture noun

lecture | \ 'lek-chər \, -shər \

Definition of *lecture* (Entry 1 of 2)

1 : a discourse given before an audience or c

2 : a formal reproof

lecture verb

lectured; lecturing \ 'lek-chə-rin \, 'lek-shrin \

Definition of *lecture* (Entry 2 of 2)

intransitive verb



Motivation

- One data sample may have different labels, e.g.,
 - POS tagging
 - Parsing
 - Sentence generation
 - etc.

Markov Model

- Finite states $S = \{s_1, s_2, \dots, s_n\}$
- You start from a state following the distribution
 $\pi = [\pi_1, \pi_2, \dots, \pi_n]$
- Transition only depends on the current state
 $\mathbb{P}[S^{(t)} = s_i | S^{(t-1)} = s_j]$



Example

- Weather
- N-gram model



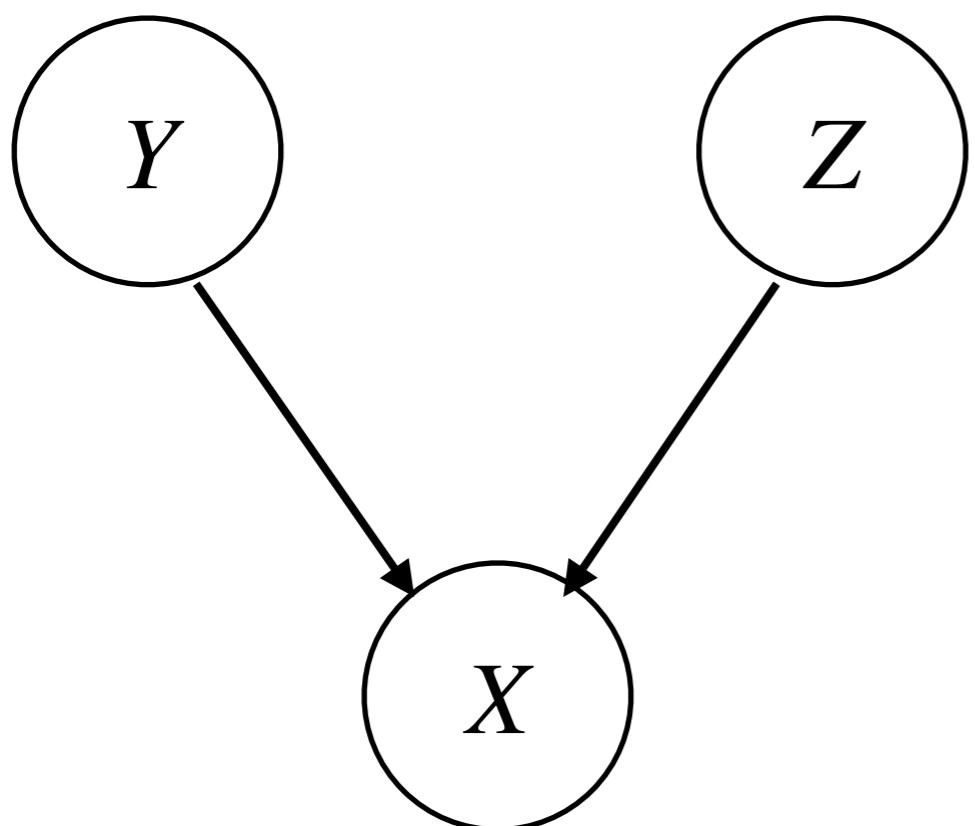
Bayesian Network in General

- Directed Acyclic Graph $G = \langle V, E \rangle$
 - Each node is a random variable
 - Each edge $a \rightarrow b$ represents that a is a direct “cause” of b
 - The joint probability can be represented as

$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i | \text{Par}(x_i))$$

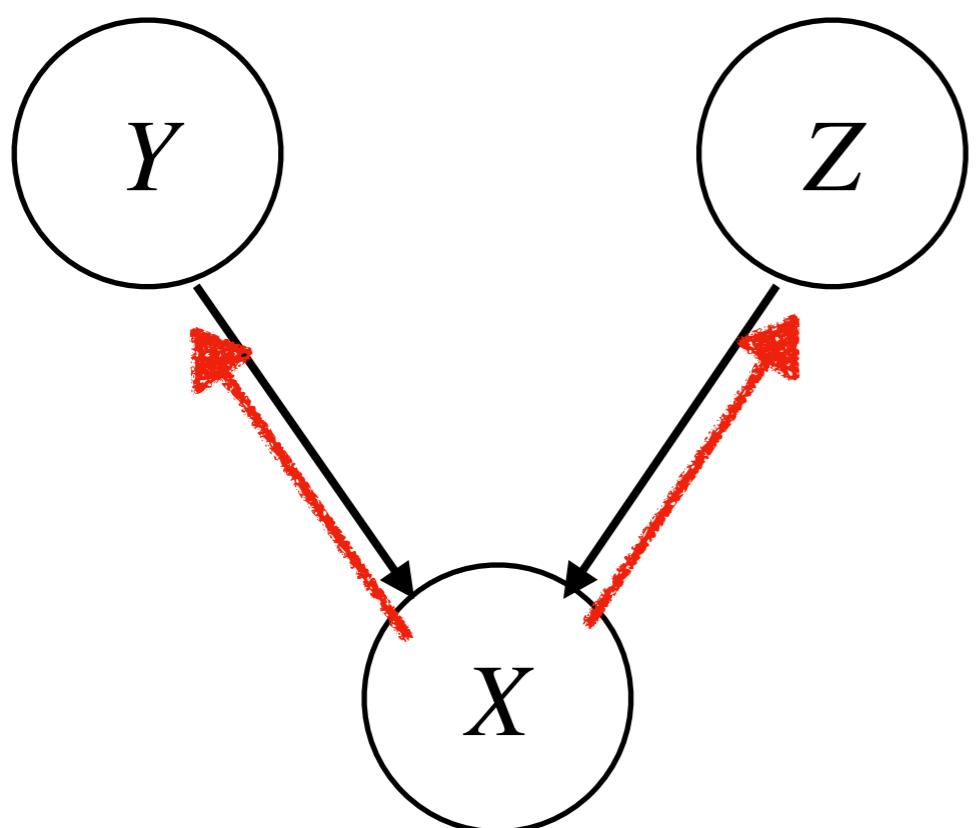
All parents

Is it correct?

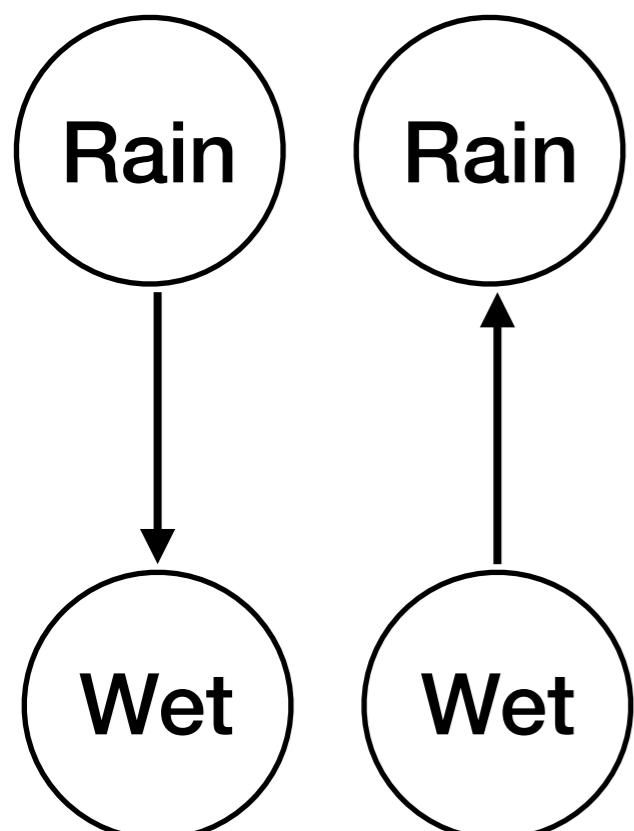


$$p(X, Y, Z) = p(X | Y)p(X | Z)$$

Can we reverse cause & effect?



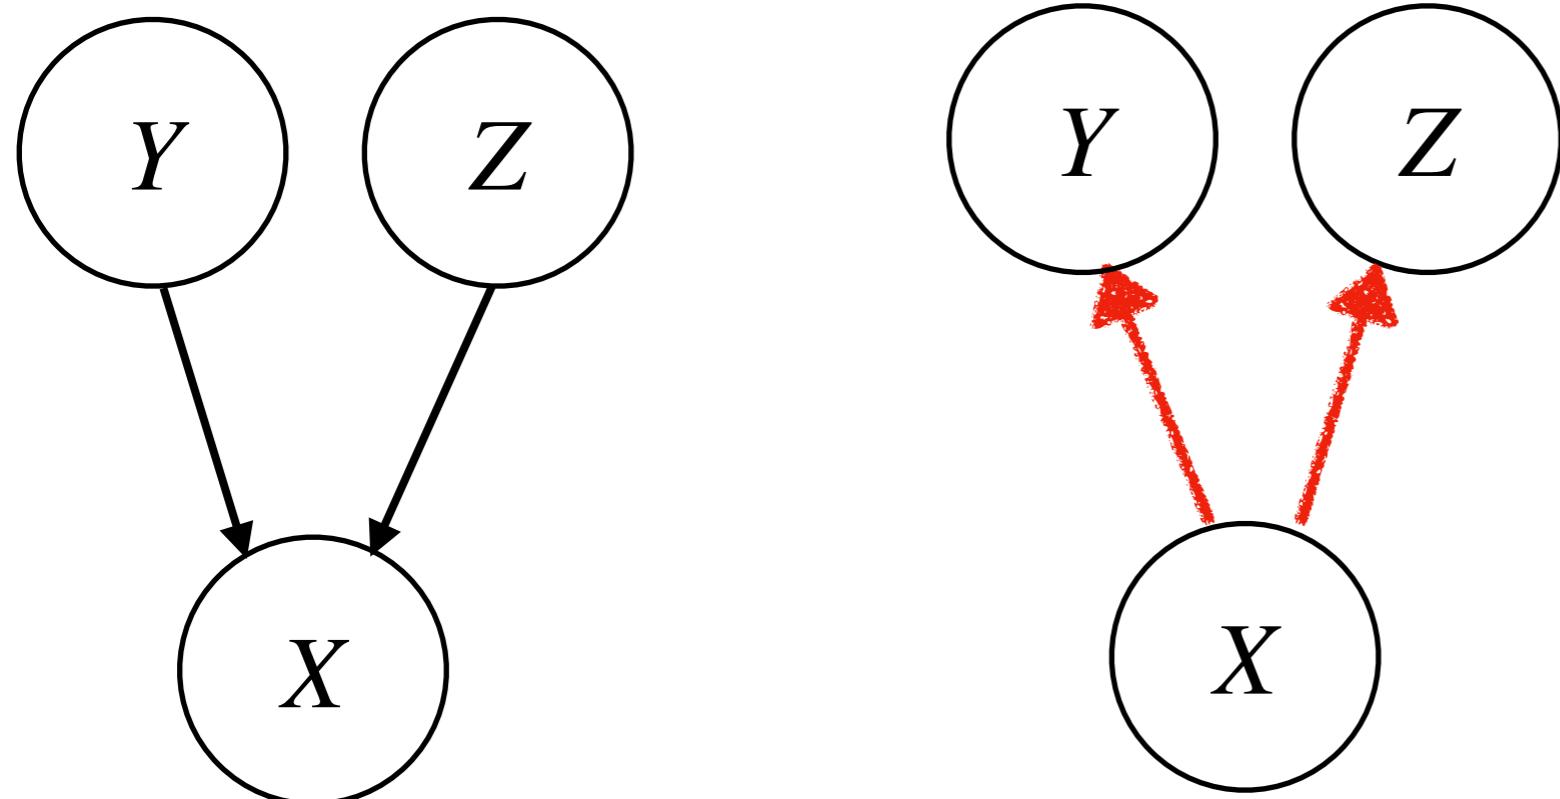
Can we reverse cause & effect?



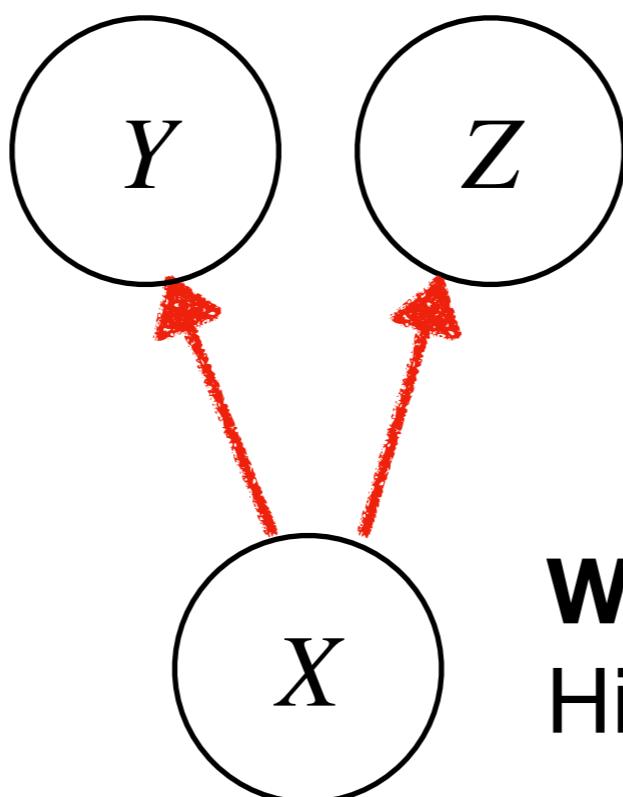
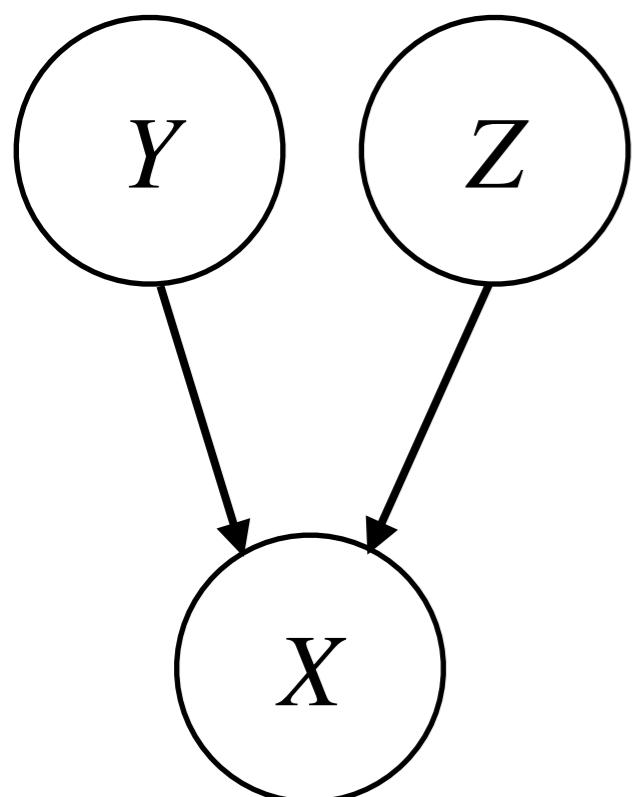
$$p(R, W) = p(R)p(W|R)$$

$$p(R, W) = p(W)p(W|R)$$

Can we reverse cause & effect?



Can we reverse cause & effect?



$$Y \perp Z | X$$

Written assignment: Prove.
Hint: By definition.

$$Y \perp Z | X$$

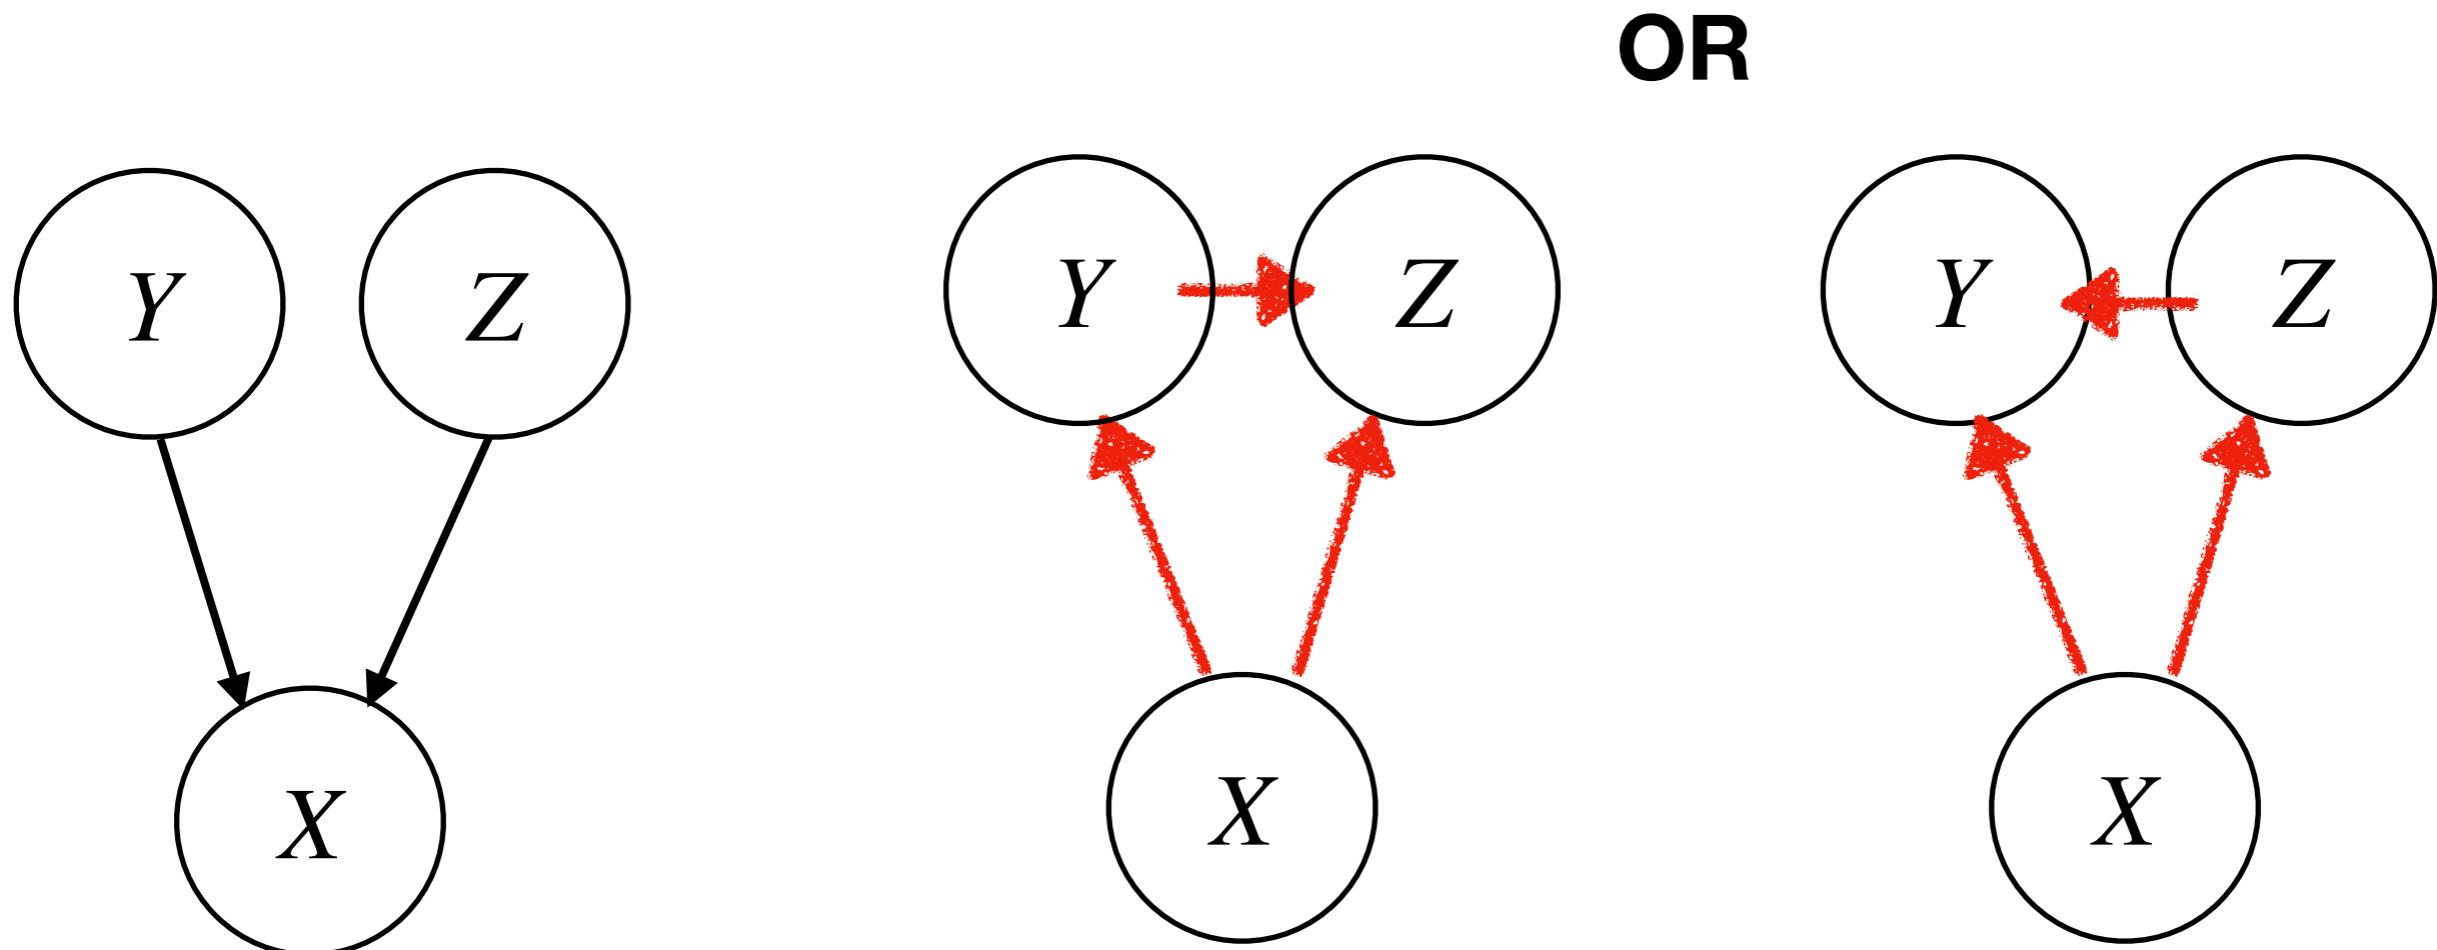
By the property of BNs,
 $Y \perp Z | X \Rightarrow 0$ mark

does not hold in general



UNIVERSITY OF
ALBERTA

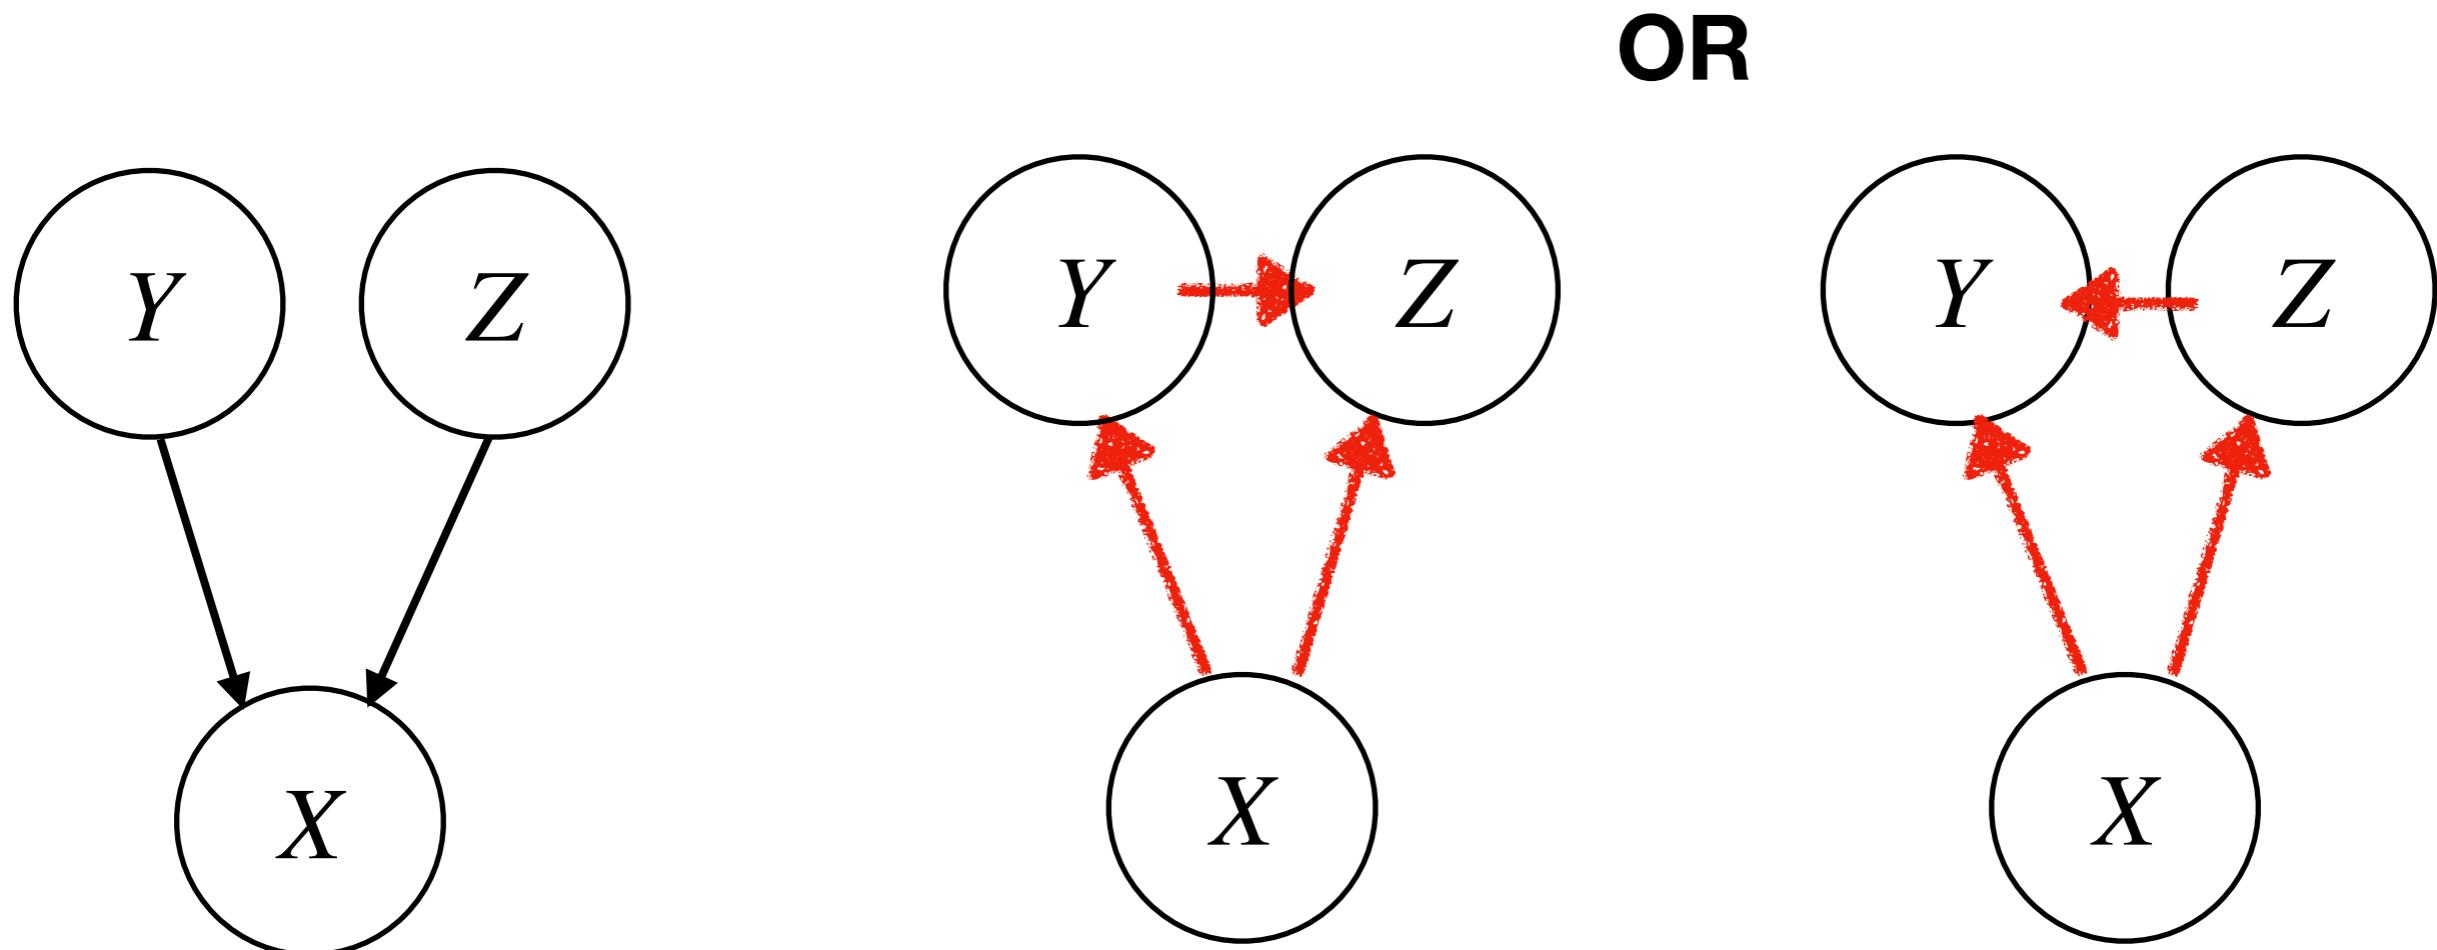
Can we reverse cause & effect?



Cause and effect cannot be formally defined.

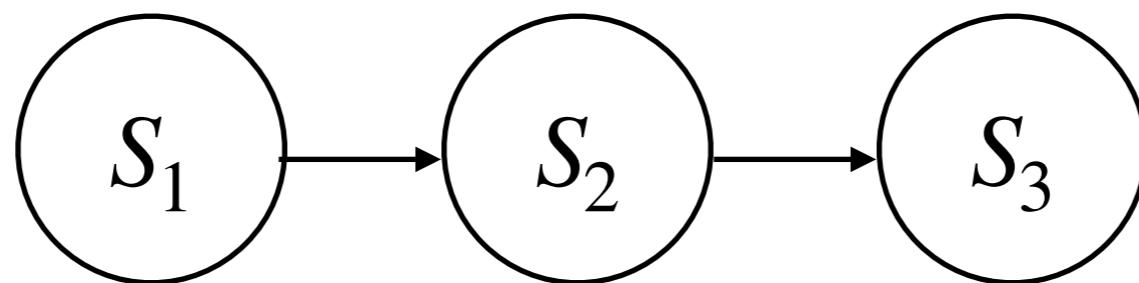
- In BN, “ \rightarrow ” refers to conditional probability
- In logics, “ \rightarrow ” refers to entailment

Can we reverse cause & effect?

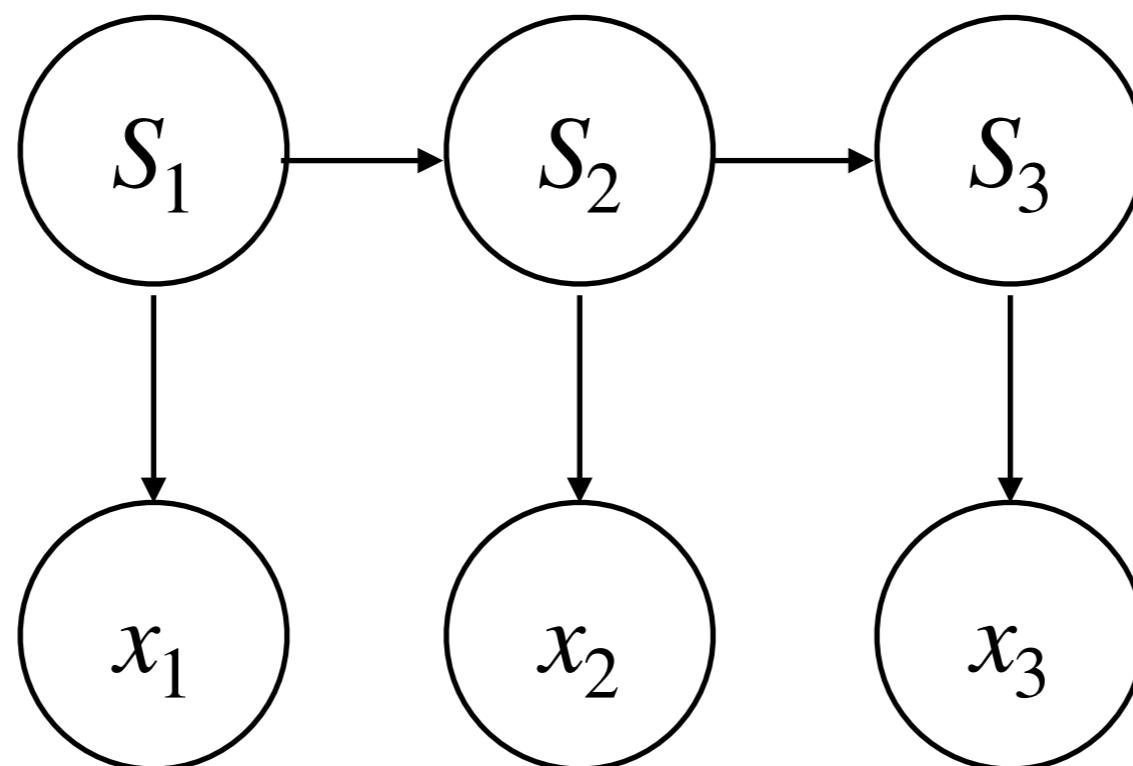


Cause and effect cannot be formally defined.
But with our intuition of cause and effect, we
can simplify our model.

Markov Model

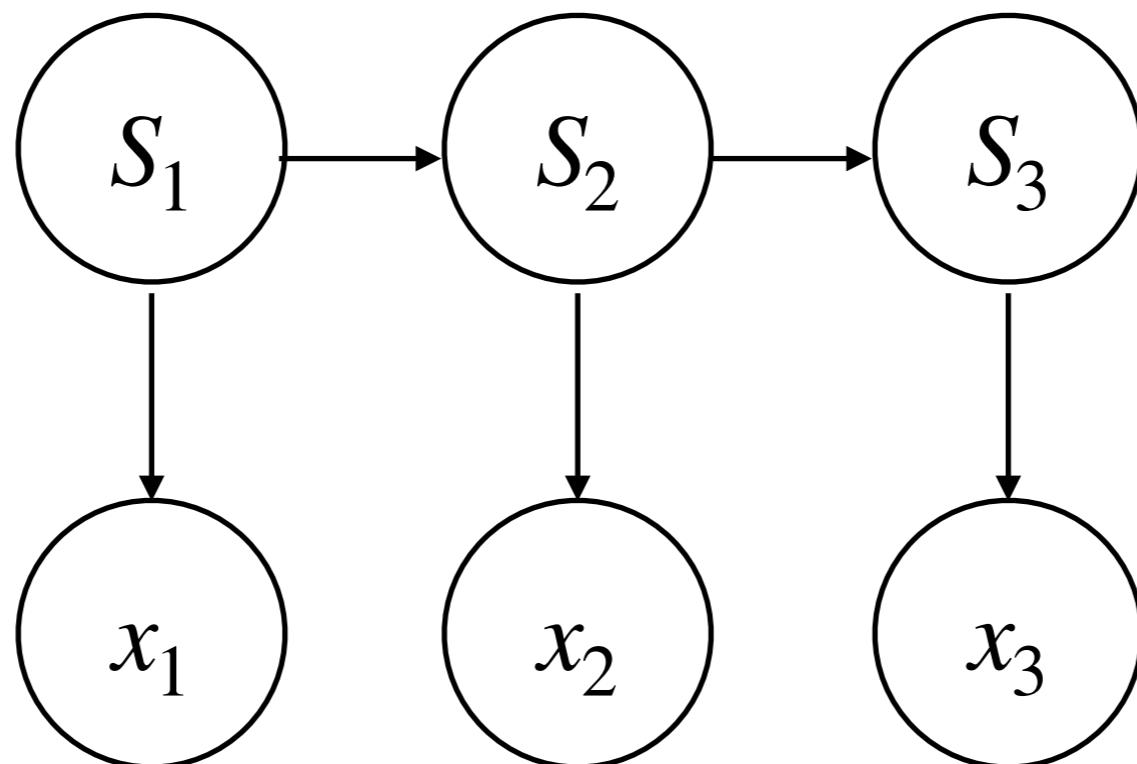


Hidden Markov Model





Hidden Markov Model

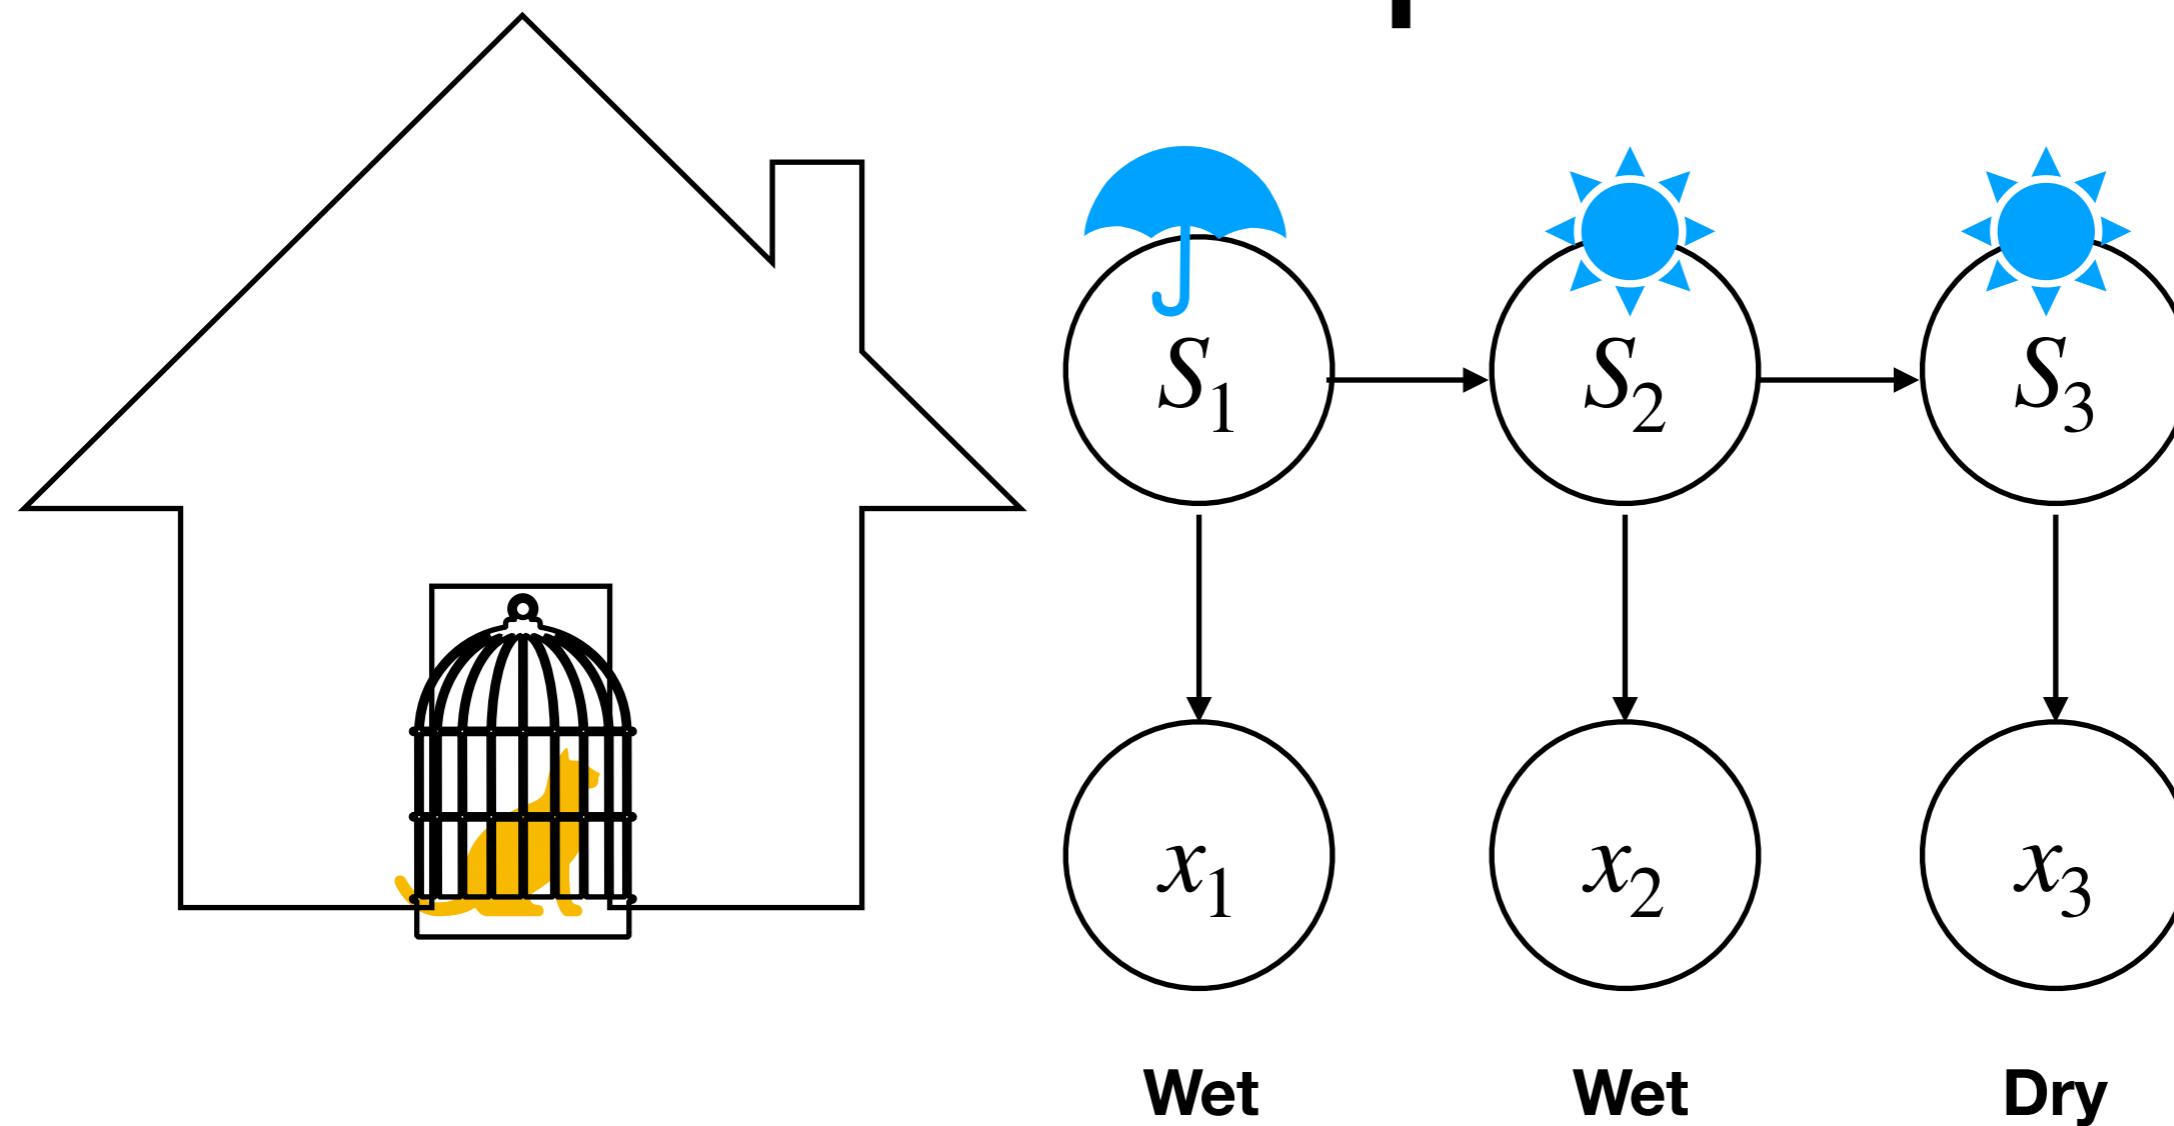


$$p(s_1, \dots, s_T, x_1, \dots, x_T) = p(s_1) \prod_{t=2}^T p(s_t | s_{t-1}) \prod_{t=1}^T p(x_t | s_t)$$

Initial State Prob. **Transition Prob.** **Emission Prob.**

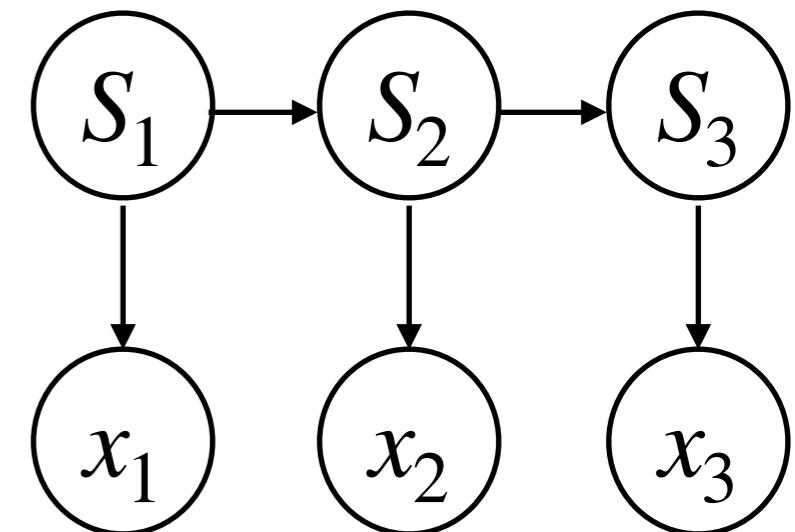
n n^2 $v \cdot n$

Example



Maximum Likelihood Estimation

- Training if fully observable
 - E.g., annotated by experts



$$p(s_1, \dots, s_T, x_1, \dots, x_T) = p(s_1) \prod_{t=2}^T p(s_t | s_{t-1}) \prod_{t=1}^T p(x_t | s_t)$$

$$\log p(\cdot) = \boxed{\log p(s_1)} + \boxed{\sum_{t=2}^T \log p(s_t | s_{t-1})} + \boxed{\sum_{t=1}^T \log p(x_t | s_t)}$$

Parameters factorize



UNIVERSITY OF
ALBERTA

MLE for Multinomial Distribution

- Counting
 - With one constraint $\pi_1 + \dots + \pi_n = 1$
 - You need to explicitly represent $\pi_n = 1 - \pi_1 - \dots - \pi_{n-1}$
 - Or, you apply the Lagrangian multiplier method

$$\log p(\cdot) = \boxed{\log p(s_1)} + \sum_{t=2}^T \log p(s_t | s_{t-1}) + \sum_{t=1}^T \log p(x_t | s_t)$$

$$\pi_i = \frac{\sum_{i=1}^M \mathbb{I}\{S_1 = i\}}{M} = \frac{\text{\# of samples that start with state } i}{\text{\# of all samples}}$$



MLE for Multinomial Distribution

- Counting
 - With one constraint $\pi_1 + \dots + \pi_n = 1$
 - You need to explicitly represent $\pi_n = 1 - \pi_1 - \dots - \pi_{n-1}$
 - Or, you apply the Lagrangian multiplier method

$$\log p(\cdot) = \log p(s_1) + \boxed{\sum_{t=2}^T \log p(s_t | s_{t-1})} + \boxed{\sum_{t=1}^T \log p(x_t | s_t)}$$

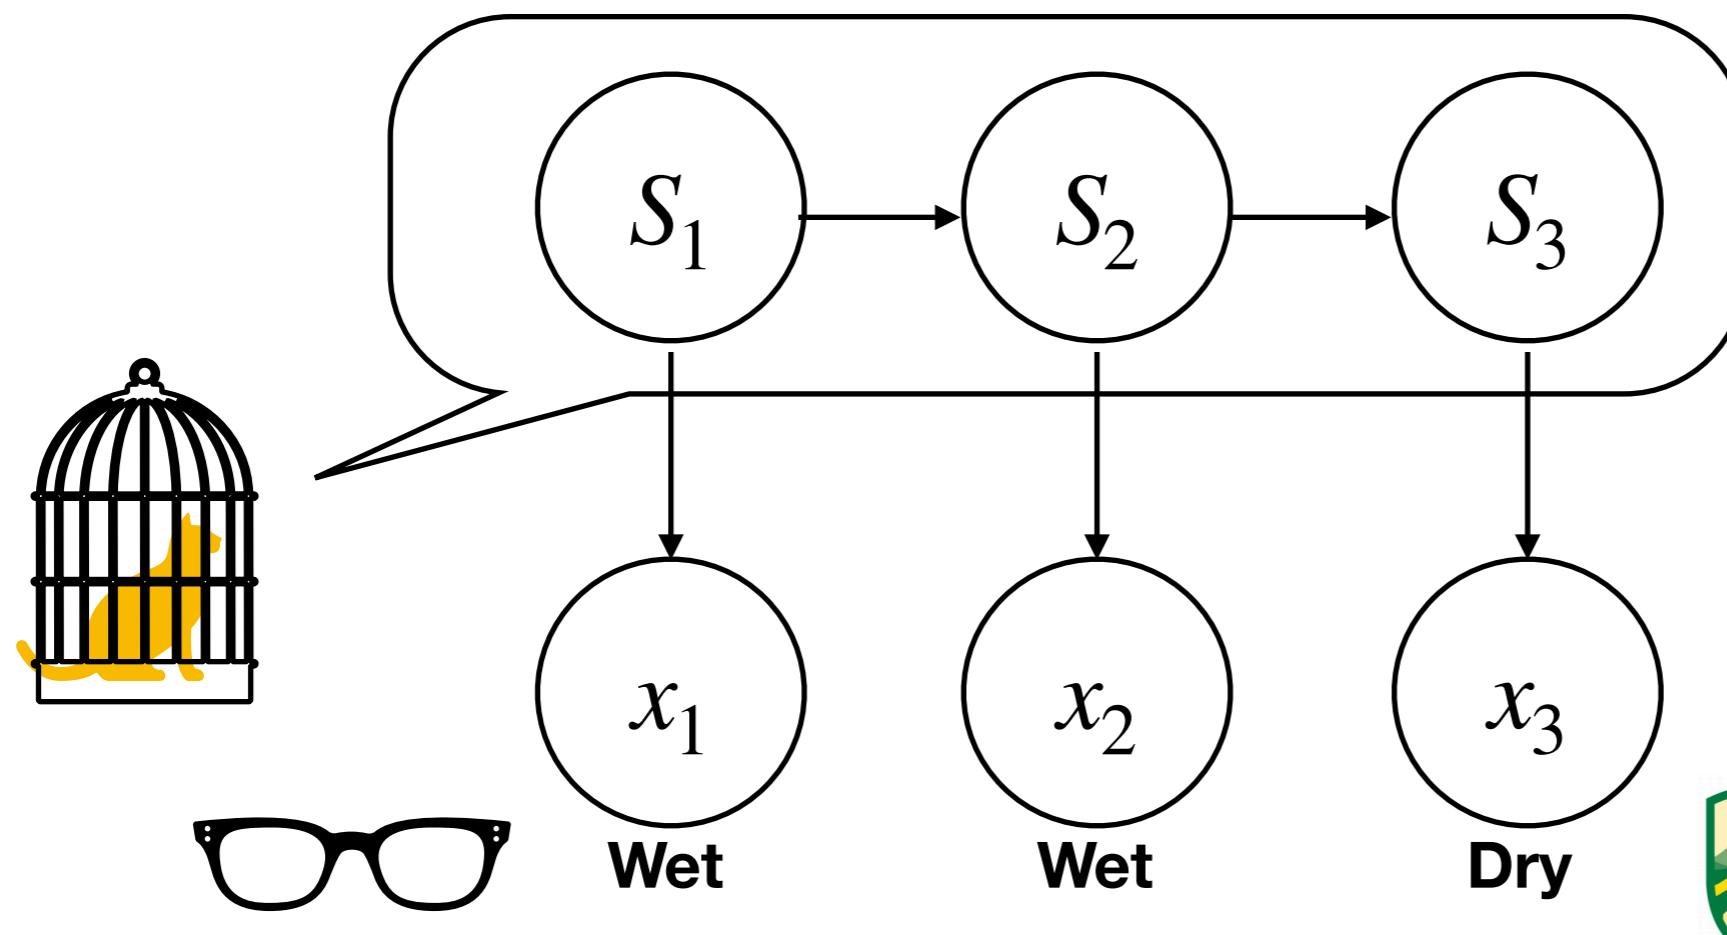
Written assignment



UNIVERSITY OF
ALBERTA

Inference

- Suppose the model is full trained
- During prediction, we observe x_1, \dots, x_T
 - How can we know the states s_1, \dots, s_T that best explain x_1, \dots, x_T ?

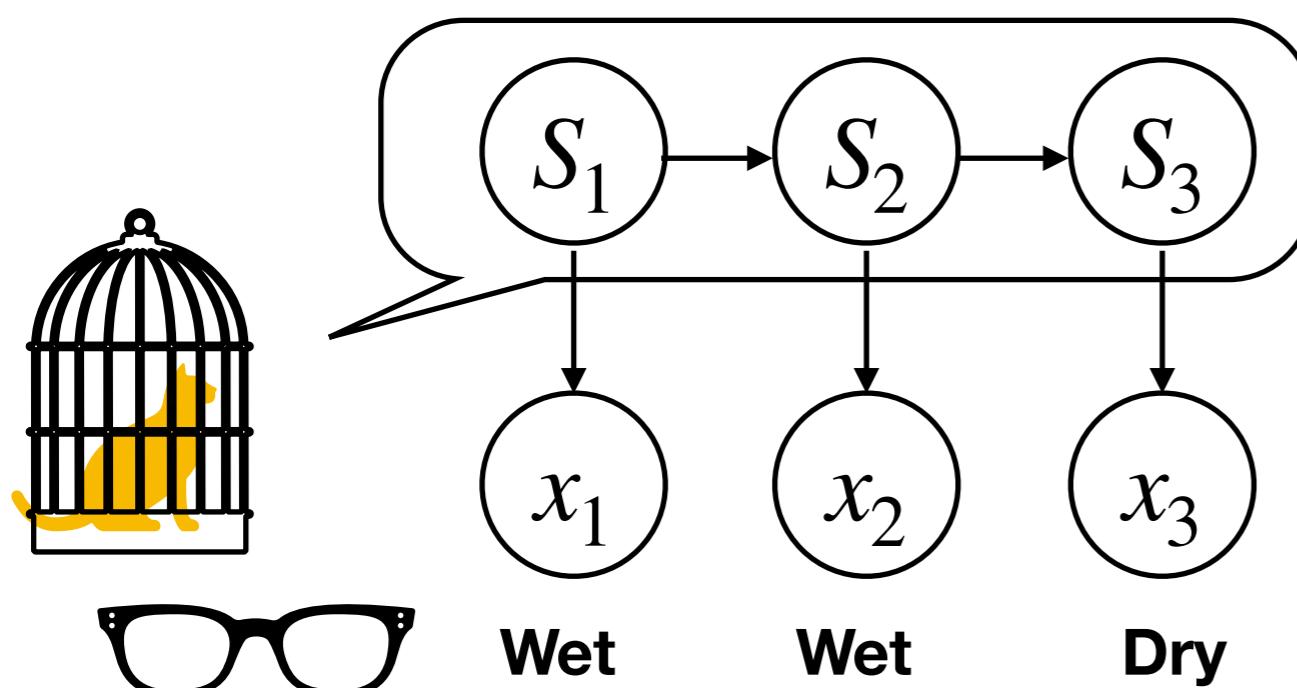


Inference Criteria

- We would like to predict the best (most probable) states
- Max *a posteriori* inference

$$s_1, \dots, s_T = \underset{s_1, \dots, s_T}{\operatorname{argmax}} p(s_1, \dots, s_T | x_1, \dots, x_T)$$

$$= \underset{s_1, \dots, s_T}{\operatorname{argmax}} p(s_1, \dots, s_T, x_1, \dots, x_T)$$



Simplified notation may be used:

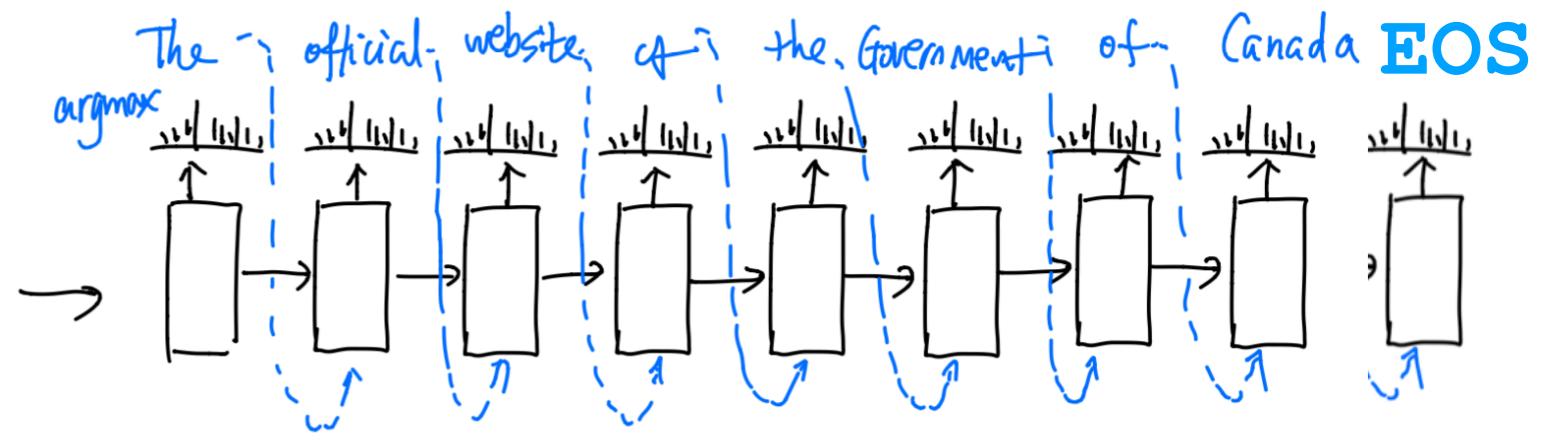
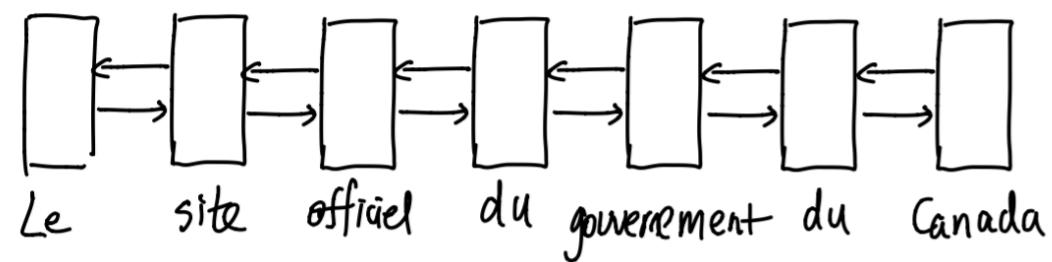
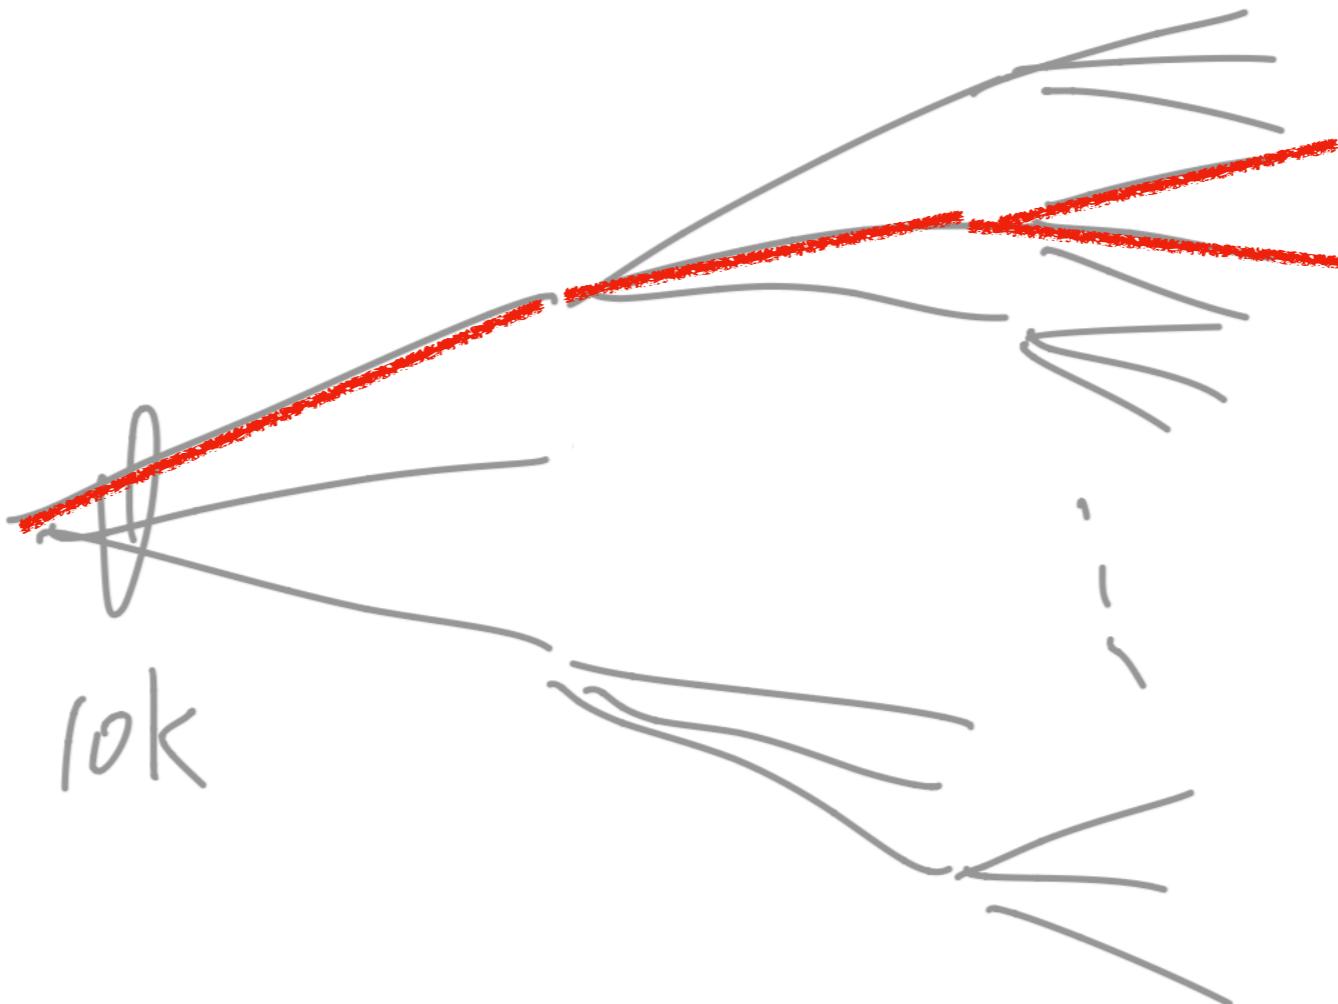
$$x_{1:t}, \quad x_1^t$$





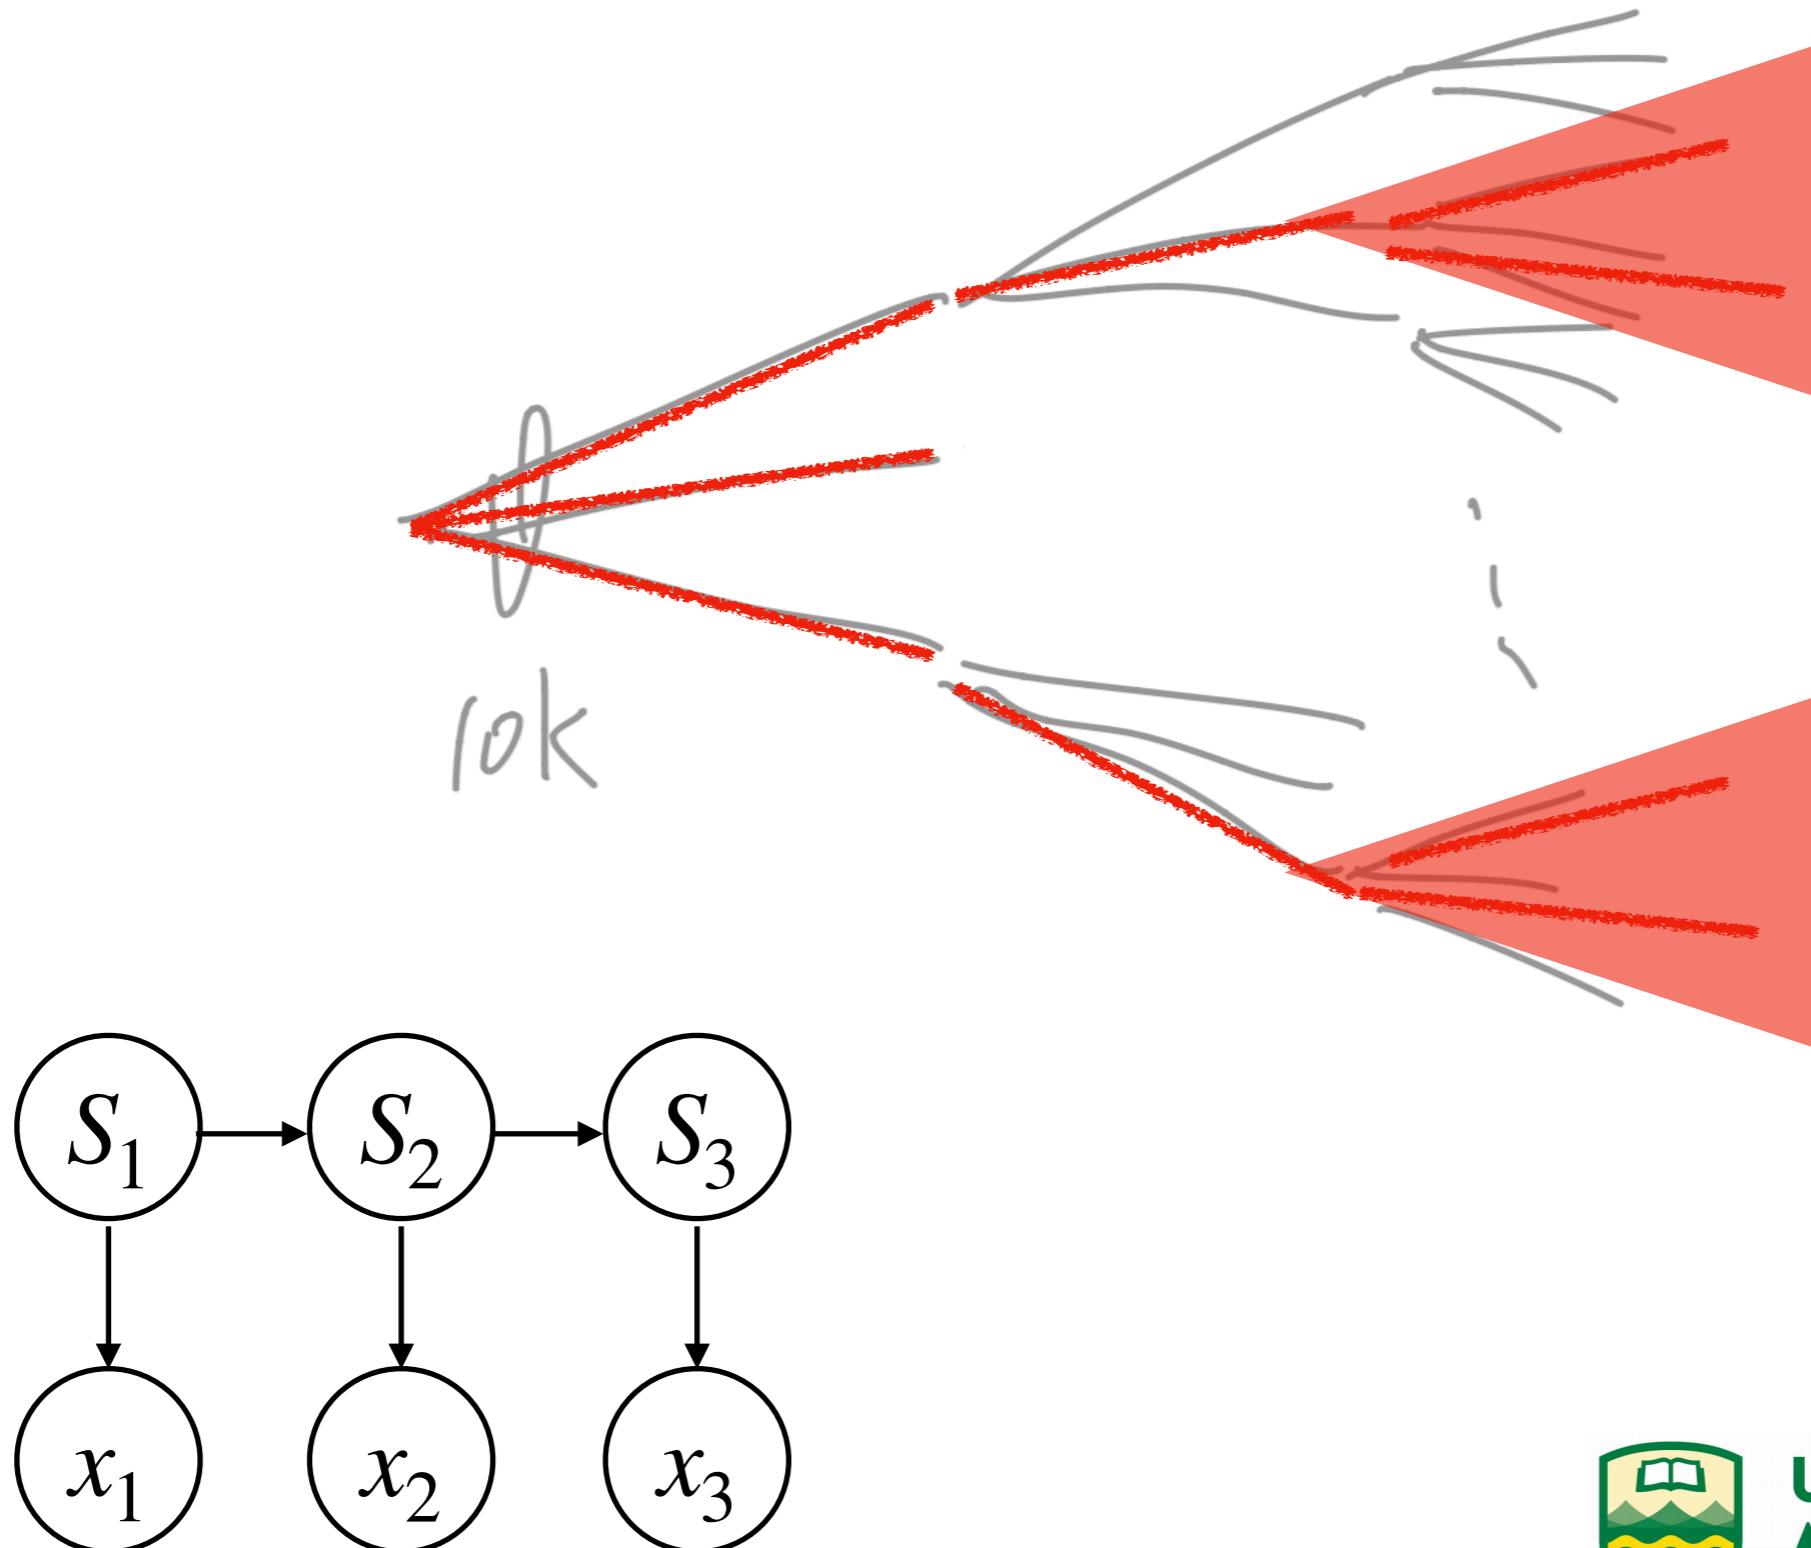
Recall Beam Search

B=2



Search in HMM

Some sub-structures are shared in different paths



Markov Blanket

$$p(s_{1:T}, x_{1:T}) = \prod_{i=1}^n [p(s_i | s_{i-1}) p(x_i | s_i)]$$

For simplicity, the first state's probability is denoted as

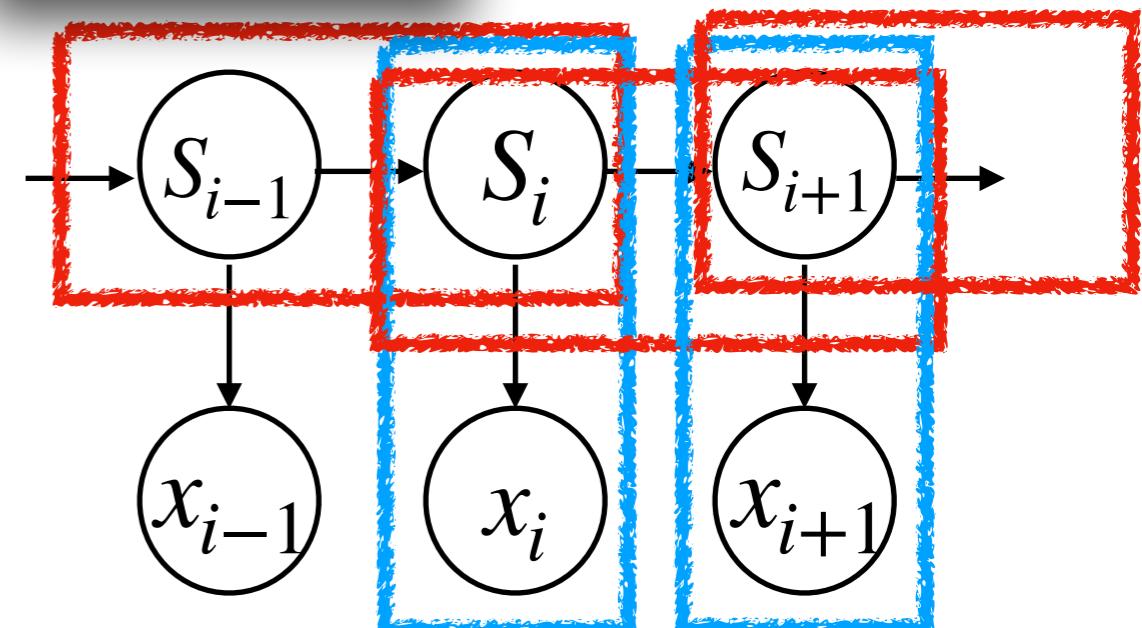
$$\mathbb{P}[s_1] \triangleq p(s_1 | s_0)$$



Key observation:

Factorized probability is local.

- $s_{i:T}, x_{i:T}$ only depends on s_{i-1}
- but not $s_{\leq i-2}, x_{\leq i-1}$



Recursion Variable

$$s_1, \dots, s_T = \operatorname{argmax}_{s_1, \dots, s_T} p(s_1, \dots, s_T, x_1, \dots, x_T)$$

$$p(s_{1:T}, x_{1:T}) = \prod_{i=1}^n \left[p(s_i | s_{i-1}) p(x_i | s_i) \right]$$

- Attempt#1: $\operatorname{argmax}_{s_{1:t}} p(x_1, \dots, x_t, s_t)$ with the best s_t
 - But best choice for every step \neq best choice globally

Recursion Variable

$$s_1, \dots, s_T = \operatorname{argmax}_{s_1, \dots, s_T} p(s_1, \dots, s_T, x_1, \dots, x_T)$$

$$p(s_{1:T}, x_{1:T}) = \prod_{i=1}^n \left[p(s_i | s_{i-1}) p(x_i | s_i) \right]$$

- Attempt#1: $\max_{s_{1:t}} p(x_1, \dots, x_t, s_t)$ with the best s_t
 - But best choice for every step \neq best choice globally
- Attempt#2: $\max_{s_{1:t-1}} p(x_1, \dots, x_t, s_t)$, for s_t being any state

$$M[t][j] \triangleq \max_{s_{1:t-1}} p(x_{1:t}, S_t = j)$$

Dynamic Programming

$$M[t][j] \triangleq \max_{1:t-1} p(x_{1:t}, S_t = j)$$

Initialization

$$\begin{aligned} M[1][j] &= \max_{\emptyset} p(x_1, S_1 = j) && [\text{nothing to choose for "max"]} \\ &= p(x_1, S_1 = j) \\ &= p(S_1 = j)p(x_1 | S_1 = j) \\ &= \pi_j \cdot p(x_1 | s_1 = j) && [\text{both are model parameters}] \end{aligned}$$



Dynamic Programming

$$M[t][j] \stackrel{\Delta}{=} \max_{1:t-1} p(x_{1:t}, S_t = j)$$

Recursion Step

$(\forall j)$

- Assume $M[t - 1][j] = \max_{s_{1:t-2}} p(x_{1:t-1}, S_{t-1} = j)$ known
- Goal: Figure out $M[t][j]$

$$\begin{aligned} M[t][j] &= \max_{s_{1:t-1}} p(x_1, \dots, x_t, S_t = j) \\ &= \max_{s_{1:t-1}} p(x_1, \dots, x_{t-1}, s_{t-1}) p(s_t = j | s_{t-1}) p(x_t | s_j) \\ &= \max_{s_t} \max_{s_{1:t-2}} p(x_1, \dots, x_{t-1}, s_{t-1}) p(s_t = j | s_{t-1}) p(x_t | s_j) \end{aligned}$$

Known by recursion assumption $M[t - 1][s_t]$



UNIVERSITY OF
ALBERTA

Dynamic Programming

$$M[t][j] \triangleq \max_{1:t-1} p(x_{1:t}, S_t = j)$$

Recursion Step

- Assume $M[t - 1][j] = \max_{S_{1:t-2}} p(x_{1:t-1}, S_{t-1} = j)$ known
- Goal: Figure out $M[t][j]$ $(\forall j)$

$$\begin{aligned}
 M[t][j] &= \max_{S_{1:t-1}} p(x_1, \dots, x_t, S_t = j) \\
 &= \max_{S_{1:t-1}} p(x_1, \dots, x_{t-1}, s_{t-1}) p(S_t = j | s_{t-1}) p(x_t | S_t = j) \\
 &= \max_{S_t} \max_{S_{1:t-2}} p(x_1, \dots, x_{t-1}, s_{t-1}) p(S_t = j | s_{t-1}) p(x_t | S_t = j)
 \end{aligned}$$

Known by recursion assumption $M[t - 1][s_t]$

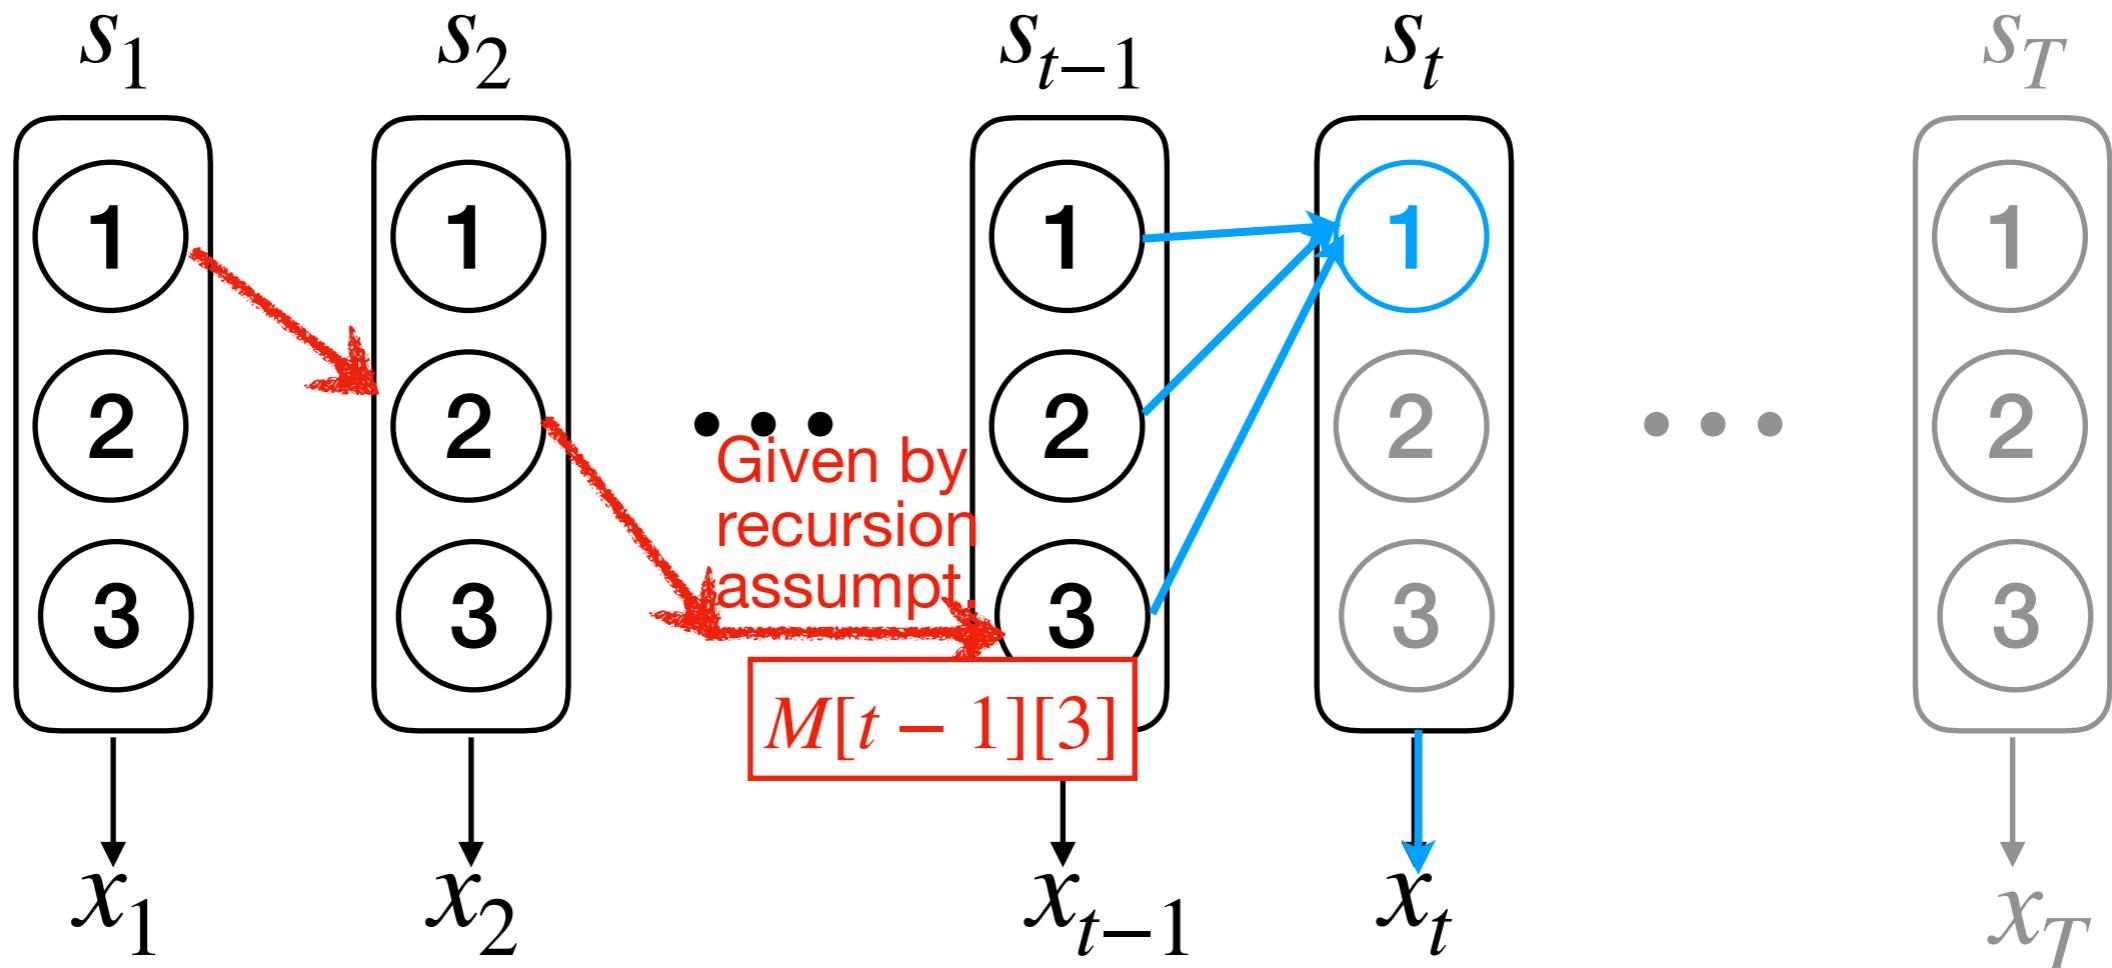


Illustration

$$M[t][j] \triangleq \max_{1:t-1} p(x_{1:t}, S_t = j)$$

Recursion Step

- Assume $M[t - 1][j] = \max_{S_{1:t-2}} p(x_{1:t-1}, S_{t-1} = j)$ known
- Goal: Figure out $M[t][j]$ $(\forall j)$

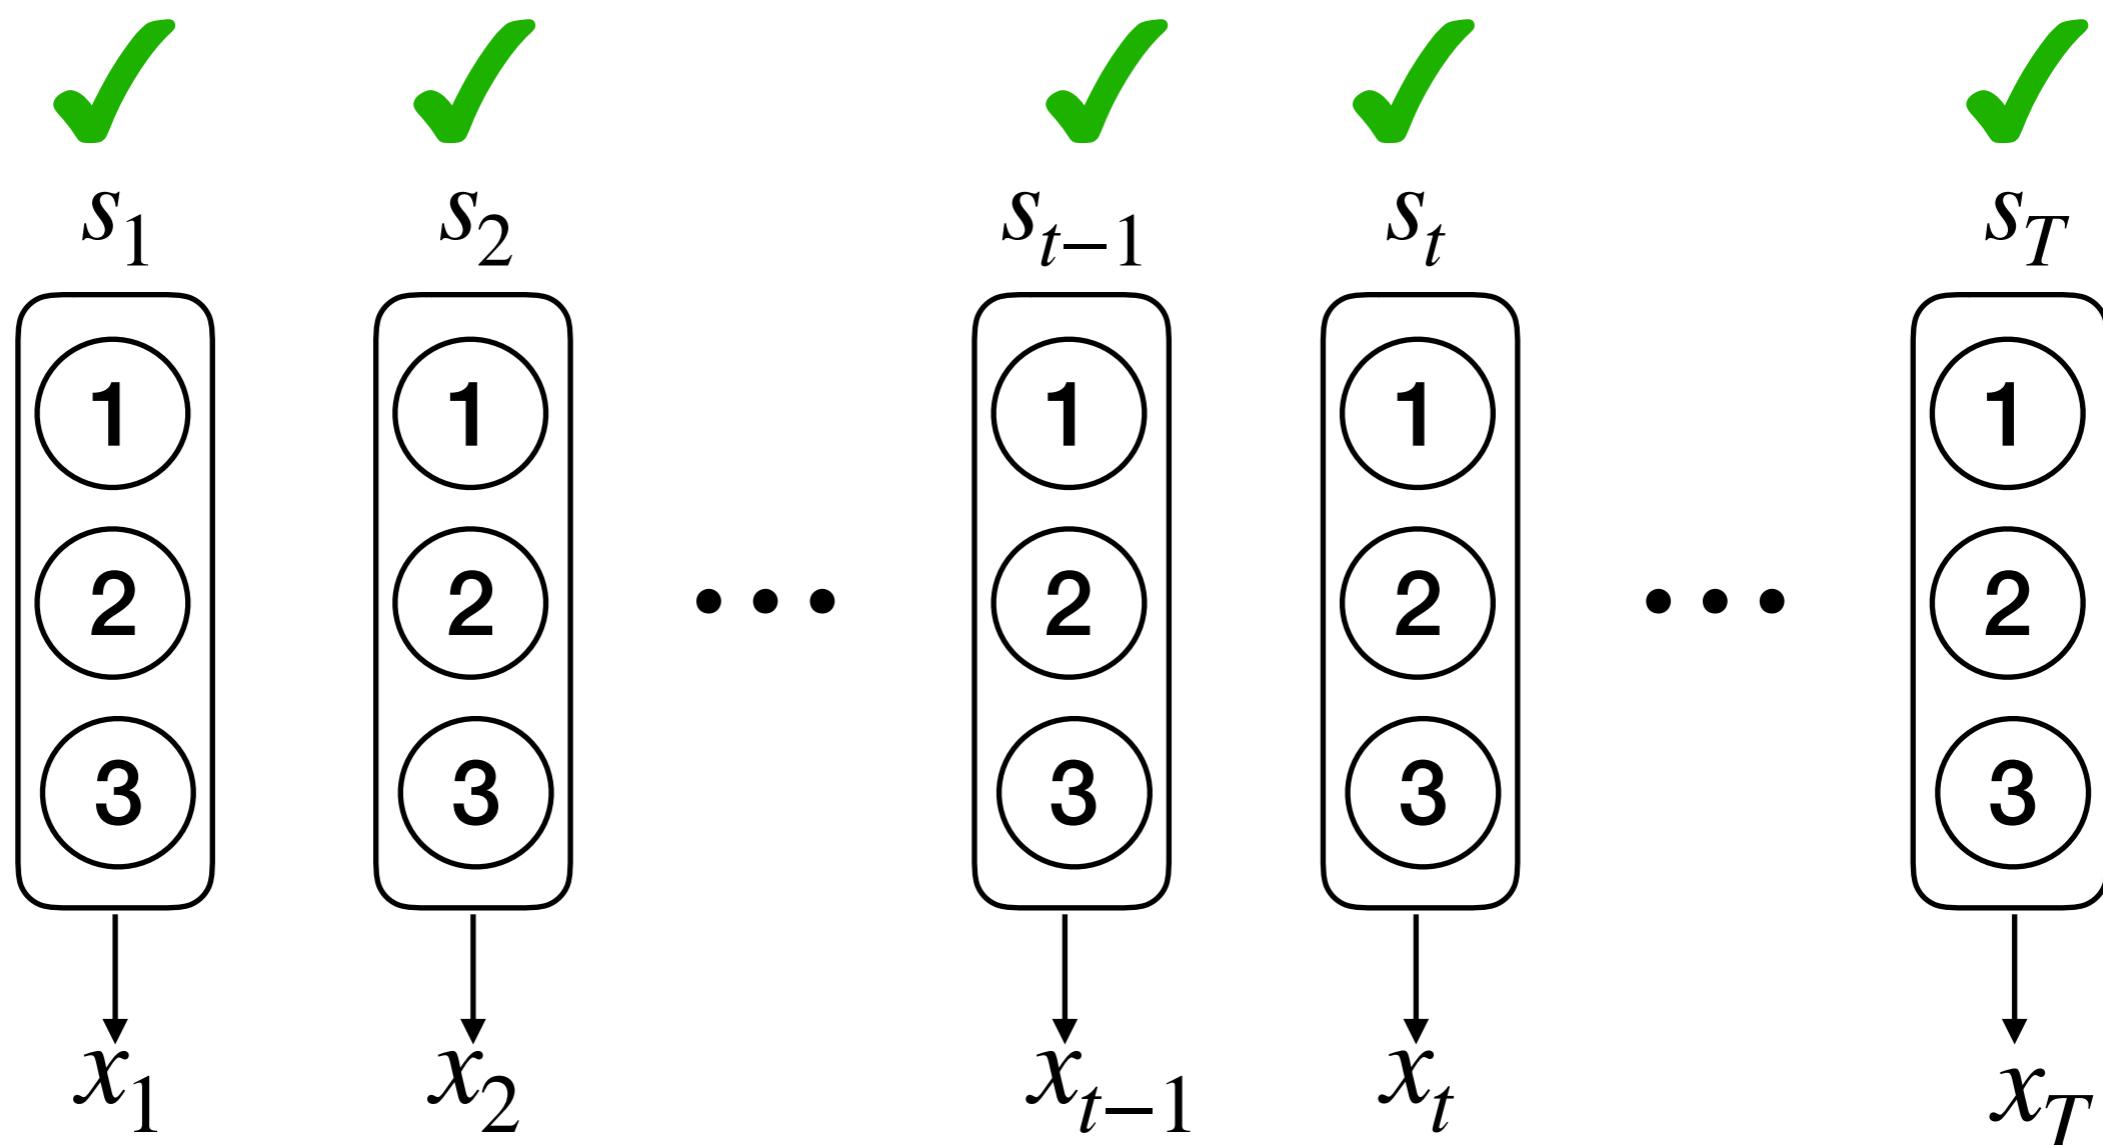


$$M[t][j] = \max_{S_{t-1}} \{ \rightarrow \rightarrow \rightarrow \nearrow \downarrow \}$$



Dynamic Programming

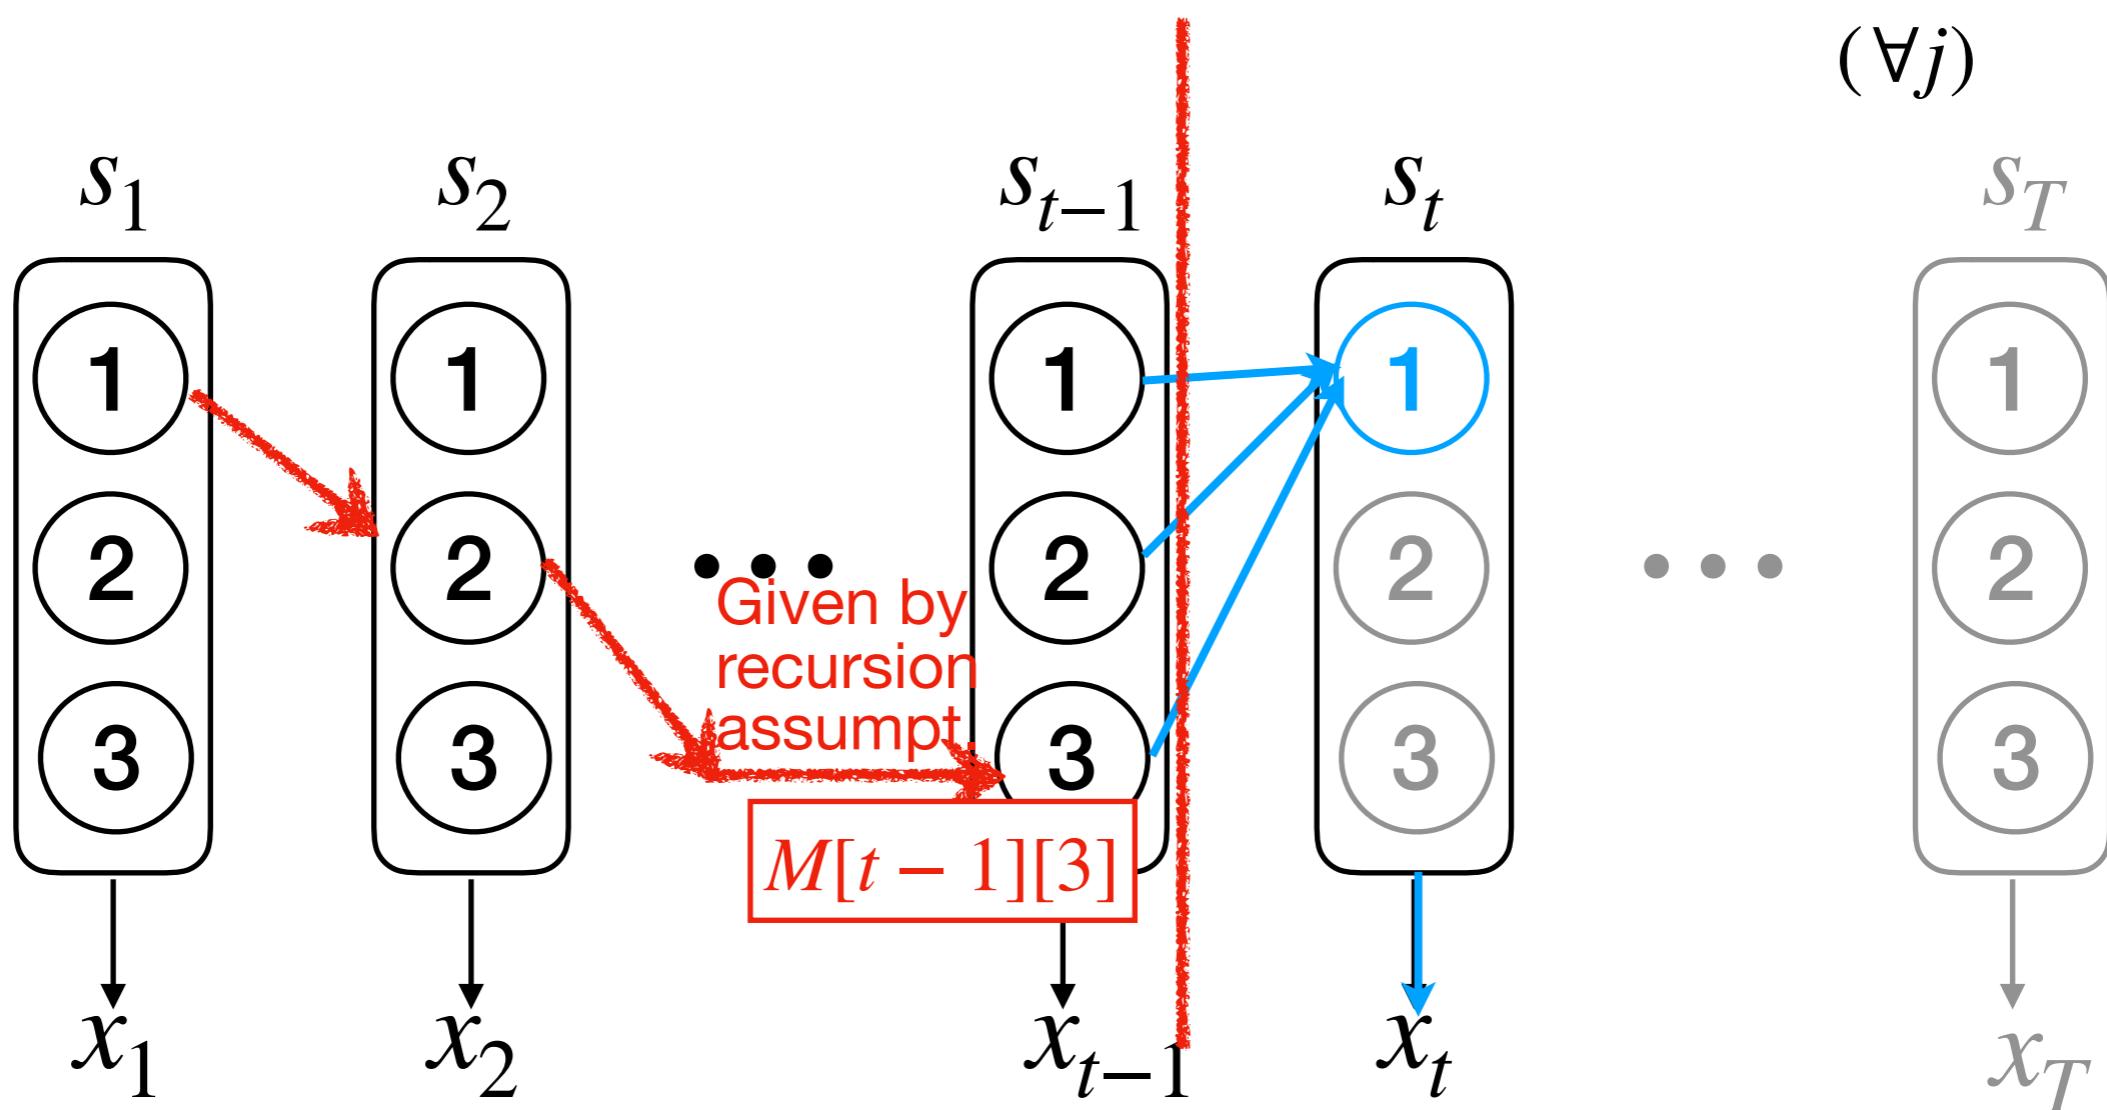
Termination: $M[T][j]$ is done ($\forall j$)



Backtracking the States

$$M[t][j] \triangleq \max_{1:t-1} p(x_{1:t}, S_t = j)$$

$$B[t][j] = \operatorname{argmax}_i \{ M[t-1][i] \cdot P(S_t = j | S_{t-1} = i) \cdot P(x_t | S_j) \}$$



$$M[t][j] = \max_{S_{t-1}} \{ \rightarrow \rightarrow \rightarrow \nearrow \downarrow \}$$



Written Assignment

- Suppose an HMM is given
 - States $S = \{1, \dots, n\}$
 - Parameters $\pi_j, P(S_{t-1} = j | S_t = i), P(x_t | S_t = j)$ known
- Suppose an HMM is given
 - We would like to find the state and output sequences of length T that have the highest jointly probability

$$s_{1:T}, x_{1:T} = \underset{s_{1:T}, x_{1:T}}{\operatorname{argmax}} p(s_{1:T}, x_{1:T})$$

- Think of the problem $x_{1:T} = \operatorname{argmax}_{x_{1:T}} p(x_{1:T})$

Written Assignment

- Suppose an HMM is given
 - States $S = \{1, \dots, n\}$
 - Parameters $\pi_j, P(S_{t-1} = j | S_t = i), P(x_t | S_t = j)$ known
- Goal
 - To find the state and output sequences of length T that have the highest jointly probability

$$s_{1:T}, x_{1:T} = \underset{s_{1:T}, x_{1:T}}{\operatorname{argmax}} p(s_{1:T}, x_{1:T})$$

- Think of the problem $x_{1:T} = \operatorname{argmax}_{x_{1:T}} p(x_{1:T})$ [optional]



Written Assignment

- Requirements
 - Design a DP algorithm, stating the initialization, recursion, and termination of the algorithm
(don't forget backpointers)
 - For any recursion variable, a clear definition is needed
 - The recursion step should be supported by derivation
 - Given pseudo code that generates $s_{1:T}, x_{1:T}$

Written Assignments

- Every week, we solve problems that have been mentioned in Monday's and Wednesday's lectures.
- Every assignment is due on next Monday
- Automatically extended to next Wednesday [**before class**]
- Further extensions require good reasons (self-approved extension may result in 0 mark).

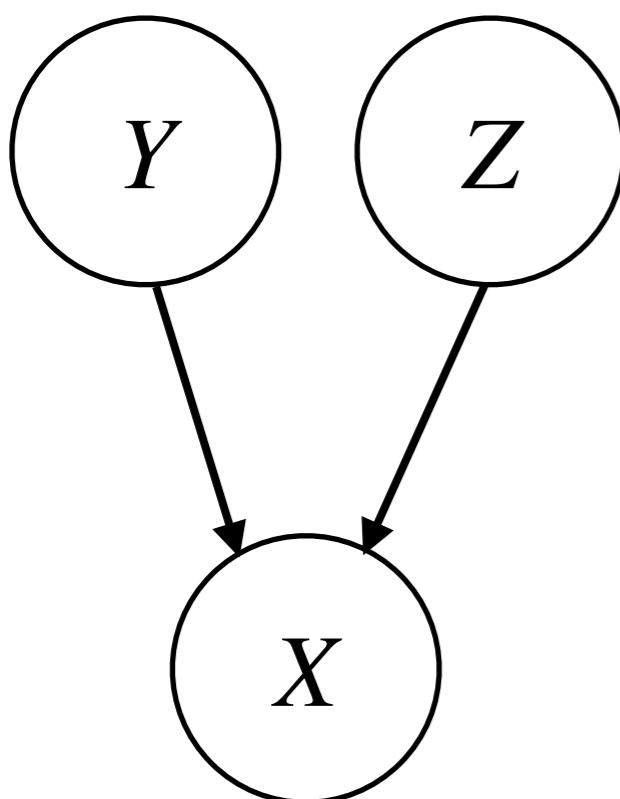


Problem 1

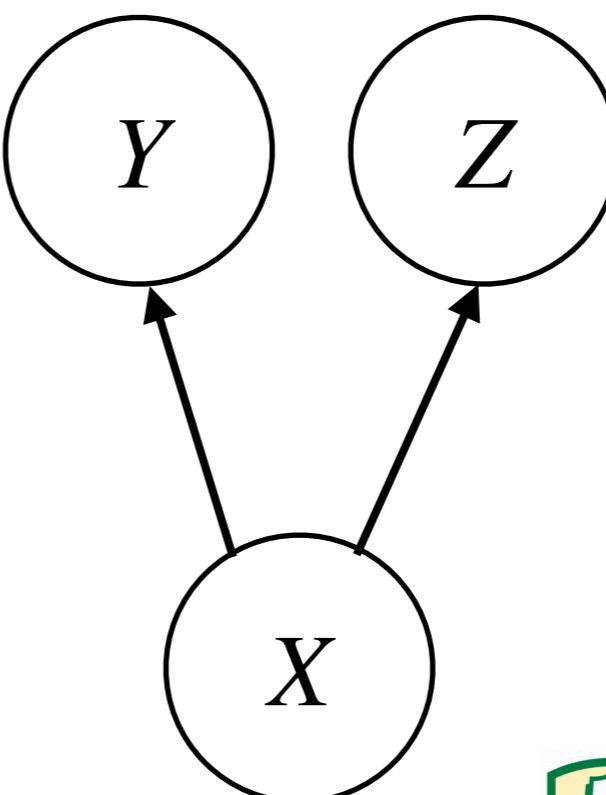
Show that $Y \perp Z | X$ does not hold in general for BN (1), but $Y \perp Z | X$ must be true for BN (2).

Note: If your solution involves showing some example, please provide your own example.

(1)



(2)



Problem 2

Give the MLE estimation for HMM transition and emission probabilities

- Figure out what are the parameters
- Give the formula to estimate these parameters (either by indicator functions or natural language expressions)

It's strongly recommended to derive MLE for multinomial distributions, but is optional for this assignment.

$$\log p(\cdot) = \log p(s_1) + \sum_{t=2}^T \log p(s_t | s_{t-1}) + \sum_{t=1}^T \log p(x_t | s_t)$$

$$\pi_i = \frac{\sum_{i=1}^M \mathbb{I}\{S_1 = i\}}{M} = \frac{\text{\# of samples that start with state } i}{\text{\# of all samples}}$$



Problem 3

- Suppose an HMM is given
 - States $S = \{1, \dots, n\}$
 - Parameters $\pi_j, P(S_{t-1} = j | S_t = i), P(x_t | S_t = j)$ known
- Goal
 - To find the state and output sequences of length T that have the highest jointly probability

$$s_{1:T}, x_{1:T} = \underset{s_{1:T}, x_{1:T}}{\operatorname{argmax}} p(s_{1:T}, x_{1:T})$$

- Think of the problem $x_{1:T} = \operatorname{argmax}_{x_{1:T}} p(x_{1:T})$ [optional]



Problem 3

- Requirements
 - Design a DP algorithm, stating the initialization, recursion, and termination of the algorithm
(don't forget back pointers)
 - For any recursion variable, a clear definition is needed
 - The recursion step should be supported by derivation
 - Given pseudo code that generates $s_{1:T}, x_{1:T}$



Thank you!

Q&A



**UNIVERSITY OF
ALBERTA**